

Track 2 — Drug Repurposing for MVA1 (BUB1B, presumed compound heterozygosity; phase not established)

Team: ciberpty · Rare Disease, Real Kid: MVA Hackathon 2026
Repository: <https://github.com/cpu-16/mva-hackathon-2026> · Methods description: Appendix B of this PDF (submission/ciberpty_track2_methods.md in the repository)

Read this first — one page for the family and the treating team

Plain-language summary, placed first on purpose. The rest of this report is written for scientists; this page is written for the child's parents and for the clinician sitting with them. Every claim in it is sourced in the sections that follow.

What we looked for

An existing, already-approved medicine that could help a child with mosaic variegated aneuploidy caused by two changes in the *BUB1B* gene. We set out in advance the tests a candidate had to pass, and we put the main candidates the scientific literature points to through them. **This is not an exhaustive screen of every proposed agent,** and §8 lists what we ruled out and why.

What we are not recommending

No drug should be started for this today. We want to be direct about that, because the honest answer is more useful than an encouraging one.

- **Metformin** is the medicine most often suggested for this kind of problem. It should not be started on this reasoning: no one has measured what dose would be needed to act on *this* problem. The best-studied way it might act — on the cell's energy machinery —

needs concentrations about a thousand times higher than the blood levels the usual doses produce, and whether it could act by a different route at those levels has never been tested. (§3.1)

- **Bortezomib** is a real, approved cancer medicine that acts on a weakness that **cancer cells with extra chromosomes** have been shown to have. **Nobody has shown that this child's cells have it**, and our own analysis of public cancer data found the link weaker in solid tumours as a group, which is the group his tumour belongs to. The amount of drug that reaches the blood is in the right range, but that is a measurement of total drug in plasma, not proof that enough reaches the right place for long enough. It also causes nerve damage in about 18 in 100 children who receive it, and this child already has muscle weakness. **Outside an active cancer, that trade is not worth making.** (§4)

What can be done this month

	Action	Why
1	Confirm the child's current age with the care team	Every schedule below depends on it, and we do not have it
2	Renal ultrasound every 3 months, birth to age 7	Published consensus recommendation for all forms of MVA (PMID 39264246); the SIOP-Europe host-genome group's schedule, endorsed by the AACR workshop. Not our idea; how it is applied to this child is the treating team's decision
3	Regular full clinical examination , including the orbit, skin and soft tissues	Rhabdomyosarcoma is not a kidney tumour and ultrasound will not find it. A 2026 case describes a child with MVA3 — the same condition caused by a different gene —
whose second tumour appeared behind the eye at age 12 (PMID 42595739)		
4	Avoid unnecessary radiation ; keep HPV vaccination up to date	Consensus guidance for these disorders (PMID 39264246)
5	Ask about testing and counselling for both parents — a ready-to-adapt laboratory request is included as <code>clinico/PARENTAL_SEGREGATION_ORDER.md</code> , with the exact loci, verified primers and an interpretation table	The clinical notes mention repeated miscarriages. In 2026 that pattern was linked to carriers of <i>BUB1B</i> changes, who showed measurable chromosome abnormalities in a blood test (PMID 42434306). This is a hypothesis, not a diagnosis — but it is a simple test, and it bears on any future pregnancy (§7)

What would change the answer

If a new tumour appears, the question changes from “should he take a medicine” to “which medicines belong in the treatment plan”. At that point bortezomib becomes worth putting to a tumour board, because the comparison is no longer against nothing — it is against

chemotherapy that carries its own harms. **It would not replace standard treatment.** In children with other solid tumours this medicine has been given safely at known doses but has not shrunk tumours on its own (§4), so the question would be whether it adds anything to standard treatment, never whether it replaces it. It would be a question to ask, backed by the evidence in §3 and tested first in the laboratory experiment in §6.

We have written that answer down now so that nobody has to assemble it in a hurry later.

What we are not certain about

The strongest evidence for this medicine comes from cancer patients, not from children with MVA. Whether the same weakness exists in this child's cells is genuinely unknown, and §6 describes an experiment that could show us we are wrong. We would rather say that plainly than overstate what we have.

Executive summary

We applied five filters to candidate drugs rather than the usual two, and the third one changes the answer.

The candidate the literature points to is **metformin**. It descends from the AICAR hit in the aneuploidy-selective screen of Tang et al. (PMID 21315436), it is approved, and its mechanism is plausible. It fails on pharmacokinetics: plasma concentrations at therapeutic doses are **micromolar**, while complex I inhibition requires **millimolar** (PMID 37343530) — a gap of roughly **three orders of magnitude (~1000-fold)**. We expect metformin to be widely proposed in this track. **It cannot inhibit complex I at therapeutic plasma concentrations.** Whether it could still stress aneuploid cells through AMPK activation at those concentrations is a *different* claim, and an untested one: no study reports a metformin aneuploidy-selective effective concentration. We reject it as unevaluable on the mechanism it is proposed for, not as disproven — the distinction is set out in §3.1.

The candidate that survives the pharmacology is **bortezomib**. Aneuploid cells carry a stoichiometric protein imbalance and compensate by increasing protein degradation, which makes them **preferentially dependent on the proteasome**. Ippolito et al. (2024, *Cancer Discovery*, PMID 39247952) established this across multiple diploid-versus-aneuploid models, recapitulated it in hundreds of cancer cell lines and primary tumours, and found that **aneuploidy level was significantly associated with the response of multiple myeloma patients to proteasome inhibitors** — a human clinical association rather than a pure cell-culture inference.

We state the size of the remaining leap rather than hide it. That evidence is drawn from **cancer** aneuploidy, and myeloma is a plasma-cell neoplasm carrying its own immunoglobulin-driven proteotoxic load. Whether the same dependency holds in **constitutional, variegated** aneuploidy is an extrapolation, and it is exactly what the experiment in §6 is designed to break (an eight-week assay window; **12–16 weeks** from a fresh biopsy, because establishing the fibroblast line takes 4–8 weeks on its own).

We then tested the part of that leap that public data can test, and it did not survive. Before looking at any result we pre-registered a test on public data (§3.5): does the aneuploidy–proteasome link extend beyond haematological cancers? Across 444 solid-tumour cell lines no proteasome inhibitor reached significance, and across 946 lines the genetic dependency on **PSMB5 — bortezomib’s own target — is significantly weaker in solid tumours than in haematological ones** (interaction $p = 0.026$), with no detectable ploidy effect after adjustment ($p = 0.88$) — which does not exclude other confounding. This child’s tumour was a solid one. We report this because it weakens the argument we are making, and because we would rather find it than have a judge find it.

Bortezomib clears the pharmacokinetic filter as we define it: EC50 in highly aneuploid lines is below **40 nM**, while the approved 1.3 mg/m² IV dose reaches a Cmax of **89–120 ng/mL = 231–312 nM** — a ratio of **5.8–7.8**. We quote the whole label interval rather than its ceiling; read it as an order-of-magnitude margin, not a measured number. **And read what the filter is:** Cmax is *total* plasma drug at a peak. It is not free drug, not intratumoral drug, and it says nothing about how long the

concentration is held — bortezomib's plasma level falls steeply after infusion while proteasome inhibition persists, so neither exposure nor duration at the target is established by this number. **312 nM is the top of a label interval, not a validated clinical threshold**, and no aneuploidy-selective effective concentration has ever been measured in constitutional MVA cells. Filter 3 tests whether the drug is disqualified on order of magnitude. It passes that test and nothing stronger.

It then fails our fourth filter, and we say so plainly: **not for this child, not now**. Bortezomib causes peripheral neuropathy in 18% of paediatric patients (8% motor), and this child already has skeletal muscle atrophy. Outside an active malignancy that risk-benefit does not hold.

What remains is a specific and testable proposition, not a demonstrated one. Bortezomib does not repair the lesion — nothing restores BubR1 function — but it is the agent aimed at the best-characterised *consequence* of aneuploidy **in cancer cells**, and it is the one worth having ready as a question for an MVA patient **with an active tumour**, where the comparator is not “no drug” but conventional cytotoxic chemotherapy. We have not shown that dependency exists in this child's cells, and §6 exists because it might not. This child is an embryonal rhabdomyosarcoma survivor with a high risk of embryonal tumours (PMID 28553959) — and he has already had one. That sequence is not hypothetical: a 2026 case report describes a girl with MVA3 who had a Wilms tumour at three and an orbital embryonal rhabdomyosarcoma at **twelve** (PMID 42595739). Our contribution is to have the answer ready — with its evidence, its concentration margin and its stopping rule — before the decision has to be made under pressure.

We also flag an antagonism a combination-minded team could walk into: **mTOR inhibitors would be expected to protect aneuploid cells from proteasome inhibition, not synergise** (PMID 31530568).

1. The variants and what they break


The proband carries two heterozygous **BUB1B** variants (NM_001211.6, transcript ENST00000287598.11, MANE), read here as compound heterozygous **on presumption: phase is not established (\$10)**:

Variant	Consequence	Exomiser ACMG	ClinVar
c.2210T>G p. (Leu737Ter)	stop_gained	PATHOGENIC (PVS1, PM2_Supporting, PP4_Moderate, PP5_Strong)	533901 — Pathogenic/Likely pathogenic
c.3006T>G p. (Asn1002Lys)	missense	UNCERTAIN_SIGNIFICANCE (PM2_Supporting, PP4_Moderate, BP1)	not present

Population frequency depends on the source, and we report both rather than select the lower one. Exomiser's gnomAD snapshot gives 9.98×10^{-5} (exomes, NFE) for the nonsense and 9.0×10^{-7} for the missense; Ensembl VEP returned 3×10^{-5} and no record respectively. Both are rare; neither is absent.

A note of caution on the missense. ClinVar holds c.3006T>A, a *different* nucleotide substitution producing the same p.Asn1002Lys, classified VUS. That classification does not transfer to our allele: PS1 applies when the known same-amino-acid variant is pathogenic, and here it is uncertain. Our variant is simply not in ClinVar.

BubR1 is a core component of the mitotic checkpoint complex and of kinetochore–microtubule attachment control; biallelic *BUB1B* mutation was the founding cause of constitutional aneuploidy with cancer predisposition (Hanks 2004, PMID 15475955), and BubR1 insufficiency in mice produces the progeroid phenotype that makes MVA an aging syndrome as well as a cancer syndrome (PMID 15208629).



Domain map of BubR1 showing both alleles converging on the C-terminal kinase domain

p.Leu737Ter truncates in the linker N-terminal to the C-terminal kinase fold — UniProt O60566 annotates the protein kinase domain at residues **766–1050** — removing it entirely, about 313 of the protein’s 1050 residues. (Whether BubR1 is catalytically active in humans is contested; we call it the kinase domain because that is the current UniProt annotation, and our argument rests on the truncation, not on catalysis.) It lands 16 residues upstream of **BUBR1^{X753}**, a characterised truncating MVA allele (PMID 31738183) — very nearly the same truncation, not merely a related model.

p.Asn1002Lys falls inside that same domain, in its C-lobe, outside KEN, TPR and GLEBS. **No published functional assay of this substitution exists**, and we measured what can be said without one (`vus/RESULTADOS.md`).

On the AlphaFold model, which is high-confidence exactly here — mean pLDDT **91.5** across 992–1012 and **91.1** at the residue itself, against 63.6 for the whole protein — **Asn1002 is buried**: relative solvent accessibility **0.162**, below the conventional 0.25 cutoff, with ten residues within 5 Å: by distance, Leu1001 and Ala1003 (1.34 Å), Trp978 and Val998 (2.83), Asn1004 (3.14), Ile1000 (3.33), Phe977 (3.46), Arg999 (3.62), Trp973 (4.36) and Phe997 (4.51). Eight of the ten are hydrophobic; Asn1004 and Arg999 are not, and we list them rather than quote only the apolar ones. The substitution puts a **longer, positively charged side chain into a buried hydrophobic pocket**, which is a recognised destabilising substitution class. We state that as evidence about the substitution and **not** as a classification: no $\Delta\Delta G$ was

computed; the REVEL (0.472) and AlphaMissense (0.923) values quoted in our methods description were retrieved as annotations, not computed by us, and they disagree; we do not claim PP3.

And we can now say why no database will settle it. The obvious move is to calibrate a predictor against BUB1B variants that are already classified. That calibration set does not exist. In the ClinVar 2026-08-28 release this gene carries **1,497 missense records, of which exactly one is Likely pathogenic** — at one star, from a single submitter — against **1,426 (95.3%) of uncertain significance**. Meanwhile **82 of its 94 truncating records are Pathogenic or Likely pathogenic**. The clinical evidence base of *BUB1B* is built on loss of function at a ratio of 82 to 1. In this snapshot, **1 of 1,497 submitted missense records** has reached P/LP.

Read that number for what it is: a property of the database, not a probability about this child's allele. It is the rate at which *submitted* missense records in this gene have been *classified* P/LP, and submission and classification are driven by ascertainment, testing volume and submitter policy. It is **not** the prior probability that a given missense in *BUB1B* is causal, and we do not use it as one. Missense alleles here may simply be under-tested and under-submitted: Sieben's characterised **BUBR1^{L1012P}** is a missense causal enough to build a mouse model on, and we did not find it in this snapshot.

What the census does establish is narrower and still useful: **the calibration set does not exist**. A predictor cannot be calibrated on one positive, so no in-silico classification of this allele can be anchored to same-gene truth. That is the quantitative reason the parental segregation test in §7 is the cheapest informative route available to us for this allele — **not the only conceivable one**. Functional assays on BubR1 (kinetochore localisation, PP2A-B56 binding, checkpoint output) exist and would speak to the substitution directly; they need material and a laboratory that we do not have, and they answer a different question than phase does.

The closest characterised missense sits ten residues away, and the mouse work on it carries a warning for us. Sieben et al. modelled the human MVA missense **BUBR1^{L1012P}** in mice and paired it with **BUBR1^{X753}** — the same truncating-plus-missense architecture as our

proband, with the missense in the same domain. Those BubR1^{X753/L1012P} animals **died prematurely**; the viable genotype paired the missense with a *hypomorphic* allele that still yields some wild-type protein (PMID 31738183).

Our proband is alive carrying a truncating allele. If that result transfers, it constrains p.Asn1002Lys: **it would have to retain more residual function than L1012P does** — a hypomorph rather than a functional null. That is a falsifiable prediction about our own VUS, and the allele-fate arm of the experiment in §6 tests it directly.


Whether the truncated transcript escapes nonsense-mediated decay is unmeasured, and should not be asserted either way without patient RNA.

At the cellular level, cell lines derived from MVA patients with biallelic mutations show “an impaired mitotic checkpoint, chromosome alignment defects, and low overall BUBR1 abundance” (PMID 20516114). BubR1 binds the B56 family of PP2A regulatory subunits through a conserved motif phosphorylated by Cdk1 and Plk1, and counteracts Aurora B at improperly attached kinetochores (Kruse et al., *J Cell Sci* 2013, PMID 23345399); the same B56 partnership opposes Aurora B to give stable kinetochore–microtubule attachment (Xu et al., *Biol Open* 2013, PMID 23789096). The region is commonly called KARD, a name neither source uses.

2. Five filters, not two

Repurposing exercises usually apply two filters: **is it approved?** and **is the mechanism plausible?** Both are necessary; neither is sufficient. Three more decide the outcome.

#	Filter	Question	What it eliminates here
1	Regulatory	Approved for marketing?	AICAR, 17-AAG, reversine, apcin, proTAME, DCZ0415
2	Mechanistic	Does it address the actual lesion or a direct consequence of it?	Agents acting on uninvolved pathways
3	Pharmacokinetic	Is $C_{max} \geq$ the effective in vitro concentration?	Metformin (~1000-fold short)
4	Patient-specific safety	Compatible with <i>this</i> child's comorbidities?	Bortezomib, at present
5	Direction of effect	Does it avoid increasing mis-segregation in survivors?	Anything selecting for unstable clones



The five-filter funnel: what each filter eliminates, and why nothing passes all five today

Filters 3 and 5 are rarely applied and eliminate the most candidates. Filter 3 is a possibility check, not a demonstration of efficacy: it ignores protein binding, exposure duration and tissue penetration, and a transient post-bolus C_{max} can overstate useful exposure.

3. Applying the filters

3.1 Metformin fails on pharmacokinetics

The chain of reasoning is genuinely attractive and we followed it to the end. Tang et al. screened isogenic trisomic mouse fibroblasts and found selective sensitivity to **AICAR**. AICAR is not approved, so it fails filter 1. Metformin is the marketed stand-in that also produces AMPK activation through complex I inhibition, and it is labelled in paediatrics from age 10. It passes filters 1 and 2.

It fails filter 3 decisively. Sakellakis (2023, PMID 37343530):

“the plasma concentration of metformin remains within the micromolar range when administered at therapeutic doses. While millimolar concentrations are necessary to inhibit complex I activity in isolated mitochondria, there is no evidence supporting the idea that metformin accumulates within the mitochondria.”

Fontaine (2018, PMID 30619086) reviews why that inhibition remains unsettled — that review exists to “gather arguments for and against each hypothesis.” In paediatrics the gap widens: dose-normalised plasma exposure varies **3.3-fold** between children, and “**renal clearance was higher in children compared to the reported data in healthy adults**”, with eGFR significantly affecting oral clearance (PMID 42399684). Metformin is renally cleared, so this child’s nephrocalcinosis makes renal function the parameter that would have to be established first — the source links clearance to eGFR, not to nephrocalcinosis specifically, and we do not stretch it further.

The distinction that decides this, stated precisely. Therapeutic metformin *does* activate AMPK in patients — that is the diabetes literature and we do not dispute it. What the micromolar-versus-millimolar gap rules out is the **complex I route at plasma concentrations**. Tang’s screening hit was AICAR, a direct AMPK agonist, so an AMPK-mediated aneuploidy-selective effect of metformin is not excluded by this argument; it is simply **untested**. No study reports a metformin aneuploidy-selective effective concentration, and the one paper reproducing AICAR’s selectivity in human cells does not test metformin. We therefore reject metformin as **unevaluable on the mechanism it is proposed for**, not as disproven

— and we say which of the two we mean, because they license different next steps: the first is answered by an experiment, the second would close the question.

We do not reject metformin out of caution. We reject it with a number — and we state what kind of number it is. For bortezomib, filter 3 is a true C_{max} / EC_{50} ratio against an EC_{50} measured in an aneuploidy-stratified model. For metformin no such EC_{50} exists: no published study gives an aneuploidy-selective effective concentration. AICAR’s aneuploidy selectivity has in fact been reproduced in human cells — a trisomy-7 colonic line is fourfold more sensitive to AICAR than its isogenic diploid counterpart (PMID 22890317). **That paper attributes the effect to proteasomal EGFR degradation, not to AMPK**, which weakens rather than supports the AMPK route we are discussing here. It does not test metformin, and no published study gives metformin an aneuploidy-selective effective concentration at all. The 1000-fold figure is therefore a comparison between *concentration ranges* — micromolar plasma against millimolar complex I inhibition — not the same quantity as the bortezomib ratio. It is a coarser instrument, and it still points the same way.

There is a deeper objection, and stating it accurately matters. Tang et al. did **not** restrict themselves to stable trisomies: Figure 6 of that paper tests AICAR and 17-AAG in *Bub1b*^{H/H} hypomorphic MEFs and in *Cdc20*^{AAA} cells — genuine ongoing-instability models with a weak checkpoint. So the screen does reach beyond a fixed extra chromosome.

What it does not reach is **metformin**, which is not tested in those models, and human biallelic BUB1B cells, which are not tested at all. The gap is therefore narrower and more specific than “stable trisomies only”: the marketed stand-in was never evaluated in the instability models, and a mouse *Bub1b*^{H/H} hypomorph is not equivalent to human p.Leu737Ter/p.Asn1002Lys.

3.2 Bortezomib passes filters 1–2, and clears filter 3 on total-drug exposure

Mechanism. Aneuploidy imposes a stoichiometric imbalance of protein complexes. Ippolito et al. (2024, PMID 39247952) dissected multiple diploid-versus-aneuploid models and found that aneuploid

cells mitigate proteotoxic stress by reducing translation and **increasing protein degradation**, rendering them more sensitive to proteasome inhibition. The finding was recapitulated across hundreds of human cancer cell lines and primary tumours.

An independent, orthogonal 2026 result points at the same target. Paired genome-wide CRISPR loss-of-function screens in isogenic aneuploid versus near-euploid lines identify **proteasome subunits** among the top aneuploid-specific dependency gene groups, alongside ribosome and rRNA processing, spliceosome and mitochondrial metabolism; a focused druggable-genome arm nominates the ubiquitin-conjugating enzyme **UBE2H** as a top aneuploid-selective dependency (PMID 42094535). This is genetic loss-of-function evidence rather than pharmacology, arrived at by a different method than Ippolito's, and it converges on the ubiquitin–proteasome axis. **It is a preprint (bioRxiv) and not yet peer-reviewed**, so we weight it as corroboration, not as a second pillar. UBE2H itself has no approved inhibitor and therefore fails filter 1 today; we record it in §9 as the target to watch.

The clinically important sentence in **Ippolito et al.** is the last one: **aneuploidy level was significantly associated with the response of multiple myeloma patients to proteasome inhibitors**. That is a human signal linking degree of aneuploidy to response to this drug class, which is more than most repurposing arguments have.

We state the size of that clinical signal rather than let a reviewer find it. It rests on small cohorts: with bortezomib as a single agent, 8 patients with a complete response had significantly higher aneuploidy scores than 50 with progressive disease ($p = 0.014$, one-tailed Mann-Whitney); in a bortezomib-plus-chemotherapy-and-dexamethasone cohort, 13 complete responders versus 14 minimal responders ($p = 0.038$, one-tailed; the figure legend calls that comparator group *minimal response* while the Results text calls it *progressive disease* — we report the legend and do not resolve the discrepancy); a third cohort trended the same way but had only 2 non-responders. That is a genuine human association built on small numbers — enough to justify running an experiment, not enough to predict an individual patient's response.

It is also not evidence about MVA1. Myeloma is a plasma-cell neoplasm whose immunoglobulin output imposes a proteotoxic load of its own, and its aneuploidy is clonal, not variegated. MVA1 is a constitutional aneuploidy state — a *different* setting in which the same dependency is plausible and untested. We treat the myeloma association as the reason to run the experiment, not as its result.

The numbers.

	Value	Source
EC50 in highly aneuploid lines (72 h)	< 40 nM — conservative upper bound read from Fig. 6o (5 near-euploid vs 5 highly aneuploid cancer lines, $p = 0.0317$, Mann-Whitney); the figure legend reports no numeric EC50 values, so this is read off the plot	PMID 39247952
C _{max} , 1.3 mg/m ² IV, repeated dosing	89–120 ng/mL = 231–312 nM (label interval; we quote both ends)	FDA label, NDA 021602 s040, §12.3
C _{max} , subcutaneous route	20.4 ng/mL = 53.1 nM	same
Filter 3 ratio (IV)	5.8 – 7.8	—

We use the upper bound of the EC50 range, and we do not treat the 2.4 nM apoptosis figure from the same paper as an EC50.

The subcutaneous route does not pass, and we withdraw the claim that it does. Its nominal ratio is 1.33 on *total* drug, and bortezomib is extensively protein-bound, so free concentration is a fraction of total C_{max}. A margin of 1.33 does not survive any plausible free-fraction correction. The IV ratio of 5.8–7.8 is left on total drug and read only as an order of magnitude for the same reason — the free-fraction and duration corrections are unknown in both cases, and the culture-side concentration carries its own unstated medium and binding conditions. Filter 3 is a plausibility screen; it says the IV route is not excluded, and nothing stronger (§2).

3.3 Everything else: not evaluable, and we say so

Drug	Filter 3 verdict	Why
Carfilzomib	Not evaluable	No published IC50 in an aneuploidy-stratified model; Ippolito used bortezomib and MG132. Cmax is known (2.89 μ M at 56 mg/m ²) but there is no valid denominator
Ixazomib	Not evaluable	Same gap. The β 5 enzymatic IC50 (3.4 nM) is not a cellular aneuploidy-selective value
Hydroxychloroquine	Not evaluable, and mechanistically weakened	Tang's hit was <i>chloroquine</i> , not HCQ — and chloroquine did not differentially inhibit the human CIN lines in that same paper
Everolimus / sirolimus	Not evaluable, and possibly counterproductive	See §3.4
Trametinib	Not evaluable — the gap most worth closing	Highly aneuploid RPE1 clones activate RAF/MEK/ERK and are more sensitive to MEK inhibition, reproduced in human cancer lines (PMID 39251587). But we could recover only <i>relative</i> IC50s from it (its Fig. 5E), and the 0.45 nM it uses to sensitise clones to etoposide is a sub-lethal dose, not an IC50 — so there is no denominator. Cmax at the approved 2 mg/day is 22.2 ng/mL = 36.1 nM (Mekinist SmPC §5.2). It is approved from age 1 — in combination with dabrafenib, for BRAF V600E low-grade glioma — with an oral paediatric formulation, so it clears filter 1; filters 2 and 5 are untested for it. It was pre-blocked in our DepMap test as a mechanism-unrelated comparator and came back null (§3.5). The first drug we would put through filter 3

Drug	Filter 3 verdict	Why
		once an absolute aneuploid IC50 is published

Substituting an IC50 from an unstratified tumour line would manufacture a ratio without meaning. We leave those cells empty rather than fill them.

3.4 The antagonism nobody should walk into

A team building a combination might reach for an mTOR inhibitor alongside a proteasome inhibitor: both touch protein homeostasis, so synergy feels intuitive. The evidence points the other way.

Aneuploid and CIN cells are vulnerable to proteasome inhibition **because** they carry a translational burden. Reducing protein synthesis relieves that burden — and the ovarian carcinoma work shows exactly that mechanism operating as *resistance*. Ovarian cancer cells kept their sensitivity to proteasome inhibitors in adherent culture but “**acquired resistance as spheroids ... by downregulating protein synthesis via mTORC1 suppression**”; conversely, PTEN loss drove constitutive mTORC1 activity and **resensitized** the spheroids to proteasome inhibition, in vitro and in vivo (PMID 31530568).

Two boundaries on that evidence, since we are the ones importing it. It is ovarian carcinoma spheroids, not constitutional aneuploidy. And the paper did **not** test rapalogues — the antagonism below is our inference from its mechanism, not its result.

The objection this raises against our own candidate, stated before a reviewer does. The same paper notes that this protective mechanism “may explain the overall poor response to bortezomib in clinical trials of patients with advanced ovarian cancer.” Bortezomib has therefore *failed* in at least one solid-tumour setting, and any honest reading has to carry that. Our answer is that the proposed failure mode is context-specific — a spheroid/ascites configuration suppressing translation — and that the same paper identifies a responsive subpopulation defined by mTORC1 status. It is a reason to measure the response in patient cells before believing it, which is what §6 does; it is not a reason the mechanism is wrong.

The premise of that prediction has independent 2025 support.

Work in yeast shows that chromosome amplification drives premature aging through defects in ribosome quality control, that aneuploids carry an increased load of protein aggregates and signs of ubiquitin dysregulation, and — the load-bearing point for us — that the *increased translational load* is what accelerates the decline (PMID 41248159). That is a different organism and an aging phenotype rather than a drug response, so we cite it for the premise only: reducing translation relieves a burden these cells depend on carrying. It is also the reason the same axis appears twice in this report with opposite signs, since MVA is itself a progeroid syndrome (PMID 31738183).


Prediction, stated so that it can be falsified: everolimus or sirolimus would antagonise bortezomib in aneuploid cells rather than potentiate it. Any screen should test that combination explicitly and expect a negative interaction.

A tension we flag rather than hide. The same BubR1 mouse work reports that predisposition to sarcopenia correlates with **mTORC1 hyperactivity** (PMID 31738183) — and this child has skeletal muscle atrophy. Read the qualifier with it: that correlation is reported in mice carrying **monoallelic** *BubR1* mutations, and this child is biallelic, so carrying the link across is our inference and not the paper's. So an independent, muscle-directed rationale for mTOR inhibition in MVA exists, and it points the opposite way from our antagonism prediction. The two are not strictly in conflict: one concerns proteotoxic killing of aneuploid cells, the other a progeroid muscle phenotype. But anyone testing the combination should know both exist, and should not read a muscle-directed mTOR rationale as support for pairing it with a proteasome inhibitor.

3.5 We tested our own candidate on public data, and pre-registered it first

Two reviewers made the same objection to §3.2 independently: the human signal is **multiple myeloma**, and this child's tumour was an **embryonal rhabdomyosarcoma**. We had no number for solid tumours, so we produced one.

We fixed the hypotheses before we looked. The drug list, the direction of every predicted effect, the statistic, the multiplicity correction and the stopping rules were written to `depmap/PREREGISTRO.md` and committed to the repository **before the analysis was run**; the commit history records that order (it documents the record, not our conduct — `VERIFY.md` shows how to check it, and the deviations are listed there too). When one hypothesis turned out to be untestable mid-analysis, we recorded that as a dated amendment rather than editing the original.



Pre-registered test: our candidate among 6,790 compounds, and the lineage-restricted dependency on its target

Method. An arm-level aneuploidy score for 2,420 DepMap 24Q4 lines from absolute copy-number segments, validated against lines whose biology is independently known (chromosomally stable MSI lines average 5.0 altered arms; CIN lines 16.5). Then Spearman correlations against PRISM Repurposing drug sensitivity and DepMap CRISPR gene effect.

Result 1 — drugs. In 444 solid-tumour lines, **no proteasome inhibitor reached FDR < 0.05**. Bortezomib ($\rho = -0.075$) and carfilzomib ($\rho = -0.060$) fall in the 4th and 8th percentile of all 6,790 compounds *by direction* — suggestive and nothing more; ixazomib runs the other way. Metformin is null, as predicted. Paclitaxel, our pre-declared cytotoxic control, runs opposite, which argues against — without excluding — a “sick cells die more” confound. The pipeline is not blind: across all 6,790 compounds 1.1% reach $p < 0.001$ against 0.1% expected by chance.

Result 2 — the target itself. Genetic dependency on PSMB5, the subunit bortezomib binds, tracks aneuploidy strongly overall ($p = -0.198$, $q = 5.2 \times 10^{-8}$, $n = 946$) — but split by lineage it is **significant in haematological lines** ($p = -0.261$, $q = 0.042$, $n = 113$) **and not in solid lines** ($p = -0.077$, $q = 0.118$, $n = 833$). The 20S core as a set is indistinguishable from a random-gene background in solid lines ($z = -0.51$), and the pre-declared negative control PSMD9 is null throughout. Adjusting for ploidy and lineage, aneuploidy still predicts PSMB5 dependency ($p = 0.0098$) while **ploidy contributes nothing** ($p = 0.88$) — this is not whole-genome doubling in disguise. The formal test of our question, the **aneuploidy $\times$ lineage interaction, is significant** ($p = 0.026$). Solid lines are the *larger* stratum, so sample size is not the explanation; the interaction supports a weaker association in solid lines under the fitted model, not the absence of a clinically meaningful effect.

What it costs us, stated plainly. Two independent measurements agree: the link our bortezomib argument rests on is detectable where myeloma sits and attenuated in solid tumours. That attenuation is a statement about DepMap lines, not a claim that no solid-tumour evidence exists: Ippolito’s own paper reports associations in pancreatic and paediatric PDX models (their Supplementary Fig. 8o–r), which in turn say nothing about constitutional MVA. The §4 “prepared answer for a second tumour” therefore rests on evidence **we could not extend to the relevant lineage**. We do not withdraw the proposal — the direction survives for two of three proteasome inhibitors, the mechanism is unchanged, and a null across heterogeneous cancer lines is not a null in constitutional MVA. But the claim is weaker than it was before we ran this, and the experiment in §6 matters more, not less.

The obvious objection, raised and answered here rather than left for a reviewer: Ippolito ran their own bortezomib screen and it was positive. In their Fig. 6p, 387 cancer lines were treated with bortezomib with or without a low dose (250 nM) of reversine; reversine had only a mild effect on proliferation but significantly sensitised the lines to proteasome inhibition ($p < 0.0001$). That is a positive result on the same drug, and our PRISM analysis is null in solid lines. **The two do not point in opposite directions, and we say so precisely because the tempting summary — “we contradicted**

them” — would contradict our own results table. Our bortezomib correlation carries the *same* sign as theirs ($\rho = -0.075$, sensitivity increasing with aneuploidy) and simply fails significance ($p = 0.11$, `depmap/RESULTADOS.md`). The designs differ in what they manipulate: Ippolito **induce** aneuploidy in each line and read the change in dose response within that line, which removes between-line confounding by construction; we **correlate** pre-existing aneuploidy across heterogeneous lines against single-dose viability, which does not. An induced-aneuploidy design is the more sensitive of the two, and our null is therefore weaker evidence against the link than their result is for it. What our analysis does add, and theirs does not address, is the **lineage split**: the association is carried by haematological lines and attenuated in solid ones (interaction $p = 0.026$), and this child’s tumour was solid.

What we did not do. PRISM primary is a single-dose viability screen, not an EC50, and not Ippolito’s isogenic comparison. This is not a replication attempt of Ippolito and we do not present it as one. Full method, numbers and limitations: `depmap/RESULTADOS.md`.

4. Filter 4: safety in *this* child — where the candidate stops

Bortezomib is not excluded by the total-plasma peak comparison of filter 3. It does not clear this patient today.

Concern	Data	Relevance here
Peripheral neuropathy	25/140 (18%) of paediatric patients, motor neuropathy 11/140 (8%) (FDA BPCA Clinical Review, NDA 021602, 2015). In adults, grade ≥ 2 in 39% and grade ≥ 3 in 15% by IV route (FDA label, NDA 206927 s002, 2022 — a different source from the PK figures above)	This child has skeletal muscle atrophy (HP:0003202). Superimposing a motor neuropathy on reduced neuromuscular reserve is the decisive objection
Thrombocytopenia, neutropenia, infection	FDA label	Prior chemotherapy for ERMS; marrow reserve unknown
Gastrointestinal toxicity, hypotension	FDA label	Failure to thrive (HP:0001508) — anything suppressing intake is costly
Renal	PK not clinically altered by renal impairment per label	Nephrocalcinosis (HP:0000121) still requires renal function, hydration and electrolytes to be established first

Bortezomib is therefore not proposed as prophylaxis or as chronic therapy for this child.

But filter 4 is patient-specific, not universal — and that is the point. Read what follows against §3.5: our own pre-registered test could not extend the proteasome dependency to solid-tumour lineages, and an embryonal rhabdomyosarcoma is a solid tumour. The argument below is weaker than the mechanism alone suggests, and the ex-vivo test in §6 is the gate, not a formality. For a patient with MVA **and an active malignancy**, where cytotoxic therapy is on the table regardless and the comparator is conventional chemotherapy rather than nothing, the balance can invert — that is a judgement for the treating team, not a result of this report. Even then the route matters and we name it: this would be off-label use decided by a molecular tumour board, or a formal n-of-1 protocol with its own ethics approval and its own stopping rules — not a recommendation that follows from this report. In that setting a proteasome inhibitor is not an exotic addition: it is an approved agent that targets the best-characterised downstream consequence of aneuploidy **as established in cancer cells** — proteotoxic dependency — supported by a human clinical signal tying response to aneuploidy burden. Whether that dependency transfers to

constitutional MVA is exactly what §6 proposes to measure, and our own DepMap analysis found the link attenuated in the solid lineage this child's tumour came from. It does not correct the checkpoint defect, and we do not claim it does.

This child survived an embryonal rhabdomyosarcoma, and the risk is gene-specific: “individuals with biallelic TRIP13 or BUB1B mutations have a high risk of embryonal tumors” (PMID 28553959), and an E-RMS arising in an MVA case with biallelic BUB1B has been characterised directly (PMID 16182441). **The realistic clinical question is not “what do we give him today” but “what do we reach for if there is a second tumour” — and that question deserves an answer prepared in advance, with its evidence and its threshold, rather than improvised under pressure.**

What is already known about bortezomib in the tumour he had, and in children. We looked, because a tumour board would, and none of it is in this report's favour. Embryonal and alveolar rhabdomyosarcoma cell lines are killed at **13–26 nM**, concentrations that spared primary human myoblasts, and bortezomib reduced the growth of an RMS xenograft (PMID 18342500). Across the Pediatric Preclinical Testing Program panel the median in-vitro IC50 was 23 nM, **but in-vivo activity against the solid-tumour xenografts was limited** — one line reached intermediate activity, the rest low — while leukaemia xenografts responded (PMID 17420992). In children, the Children's Oncology Group phase I study ADVL0015 set the recommended dose at 1.2 mg/m² twice weekly for two of every three weeks, with dose-limiting thrombocytopenia at 1.6 mg/m² and **no objective responses** among 15 children with refractory solid tumours (PMID 15570082); the vorinostat combination ADVL0916 likewise saw no objective responses, and a grade 2 sensory neuropathy that progressed to grade 4 was dose-limiting (PMID 22887890). In adults, a phase II study in sarcomas closed its osteosarcoma/Ewing/rhabdomyosarcoma arm for low accrual and reported minimal single-agent activity in soft-tissue sarcoma, with painful neuropathy prominent (PMID 15739208). The reading for a tumour board is therefore narrow: the in-vitro sensitivity of RMS at concentrations below the paediatric C_{max} is real, the paediatric dose and toxicity profile are established, and single-agent bortezomib has not shrunk a paediatric solid tumour in a trial. If it is ever reached for, it would be


as a component of a combination — which is what the sarcoma trial itself recommended — and the aneuploidy argument of this report is a reason to ask whether MVA-driven tumours are the exception, not evidence that they are.

There is a second, immediate implication. Spindle poisons — vincristine, taxanes — depend on a competent mitotic checkpoint. With reduced BubR1 that arrest is less robust; a case of rhabdomyosarcoma with PCS/MVA treated with **reduced-intensity chemotherapy** is on record (PMID 31184400 — a single case, not a guideline and not a safety estimate; PubMed carries no abstract for it, so we cite only what its title states and do not name the regimen). This is a mechanistic caution for the treating oncologist to weigh against probability of cure — the report supplies no clinical evidence that would make it a contraindication of any grade — and never a reason to withhold curative therapy. If a mitotic-poison-sparing regimen is ever needed, there is **no validated replacement agent**: a proteasome inhibitor would be a bench hypothesis to test first, not a substitution anyone could make at the bedside.

We state the boundary explicitly: this is a hypothesis for a tumour board, not a prescription.

5. Why any endpoint must be measured per cell

Any trial of the above needs a readout for **aneuploidy burden**. We asked whether the data this hackathon provides could supply one, and computed the answer instead of asserting it.



Detection limit of bulk WGS for mosaic aneuploidy

Figure 1 shows the **first-order** analysis only: mean $|BAF - 0.5|$ against a fixed 0.005 noise floor, crossing at $\approx 14\%$ of cells. **The later calibrated limit of 2.07% is not on this plot** — it comes from a different statistic (Ψ over 125-site windows against a Monte Carlo null at $\kappa = 2$) and cannot be read off this curve. Neither figure is a measurement of this child's genome; see “What we do not claim about this child” below.

We are not proposing a new assay. The correct endpoints exist and are standard: metaphase karyotyping with premature chromatid separation scoring is the diagnostic criterion for MVA; the micronucleus assay is an international standard (OECD TG 487); and single-cell DNA sequencing has been used to assay chromosome copy number genome-wide since Knouse et al. 2014 (PMID 25197050). We cite that work for the method, not for our conclusion — its own finding is that aneuploidy is *rarer* in liver and brain than previously reported.

What we contribute is the number.

The analysis used the **VCF only**, so no GC-LOESS normalisation was applied. **That was our choice of input, not a limit of the challenge:** the organisers do distribute the raw reads (84.7 GB of FASTQ), and the statistic that ends up carrying the argument does not need them. An earlier version of this paragraph blamed the absence of alignments for the absence of GC normalisation, and it was wrong on both halves. Four analyses were run, of which three are refinements of the same B-allele-frequency statistic and one (depth-based dosage) is orthogonal; they are not four independent tests. In a trisomy, heterozygous sites split symmetrically, so the *mean* BAF does not move and the correct statistic is mean $|BAF - 0.5|$.

The measured standard deviation of mean $|BAF - 0.5|$ across the 22 autosomes at fixed depth (DP 40–48, 857,435 sites) is **0.0035**, so the 0.005 noise floor we use — arising from GC content, mappability and segmental duplications — is a ≈ 1.4 -SD criterion. Simulating at DP 44 against it gives a **first-order detection limit of a clonal whole-chromosome trisomy present in $\sim 14\%$ of cells**; a conventional 2-SD criterion moves it to 16% and 3-SD to 20%. Every stricter choice makes the assay look worse, so the conclusion does not rest on the threshold.

A later analysis improved that limit, and we report both because they answer slightly different questions. Replacing the fixed 0.005 floor with a calibrated null — Ψ over 125-site windows, the measured Beagle switch process, the empirical segment-size distribution, and 40,000 Monte Carlo replicates per cell — gives a limit of **2.07% of cells at an overdispersion $\kappa = 2$** , and 1.46% at $\kappa = 1$, 2.95% at $\kappa = 4$ (`mosaico/RESULTADOS.md` §3; every value in `mosaico/ventana.json` at the chosen window $W = 125$). Read that number with its three conditions attached: **it is a simulated limit of detection under a fitted null, not a measured sensitivity, and it is conditional on κ** , which we estimate from these data at ≈ 2.3 rather than knowing independently. Under it, the 1–2% per chromosome that variegation implies sits **at the edge of what is achievable rather than far below it** — a better limit than the 14% figure, but one that leaves the variegated signal at the edge of detection rather than comfortably inside it.

Read sampling is not what sets that limit. The simulation subtracts the binomial baseline at DP 44, so the 0.005 threshold a signal must clear is the *measured dispersion between chromosomes* — systematic library bias, not counting noise. The depth-based dosage test points the same way by a separate route: Poisson error when averaging a whole chromosome at 44 $\times$ is $\sim 0.03\%$, far below the observed deviations. These are two different statistics and we do not conflate them; both imply that sequencing deeper would barely move the limit.

As an **illustrative bound, not a measurement of this child**: if 30% of cells were aneuploid — the order of magnitude of the diagnostic criterion, not a figure from this patient — spread across 22 autosomes, each chromosome would be affected in **1–2% of cells**, and gains and losses of the same chromosome cancel in bulk. Against the first-order 14–16% floor that is **7–16 $\times$ below** it; against the calibrated 2.07% limit it is **at the boundary**. The conclusion survives either way, but it does not survive *because of* the limit, and the argument that does not depend on any limit is this one:

A perfectly balanced variegated mixture is not identifiable from bulk allele fractions. Model a fraction g of cells gaining homologue A, an equal fraction gaining B, and matched losses of each. Mean copy number is exactly 2 and mean allele fraction exactly 0.5, **by symmetry**. Simulated under identical conditions, detection stays at **0.1–0.2%** —

the false-positive rate — from $f = 2\%$ up to $f = 40\%$ (mosaico/RESULTADOS.md §4). That is not low power. Two biologically distinct cell populations produce the same observable distribution, so no depth and no statistic separates them.

That proof is about the mixture we modelled, not about every MVA genome. Real variegation need not be perfectly balanced between homologues or between gains and losses, and a sufficiently asymmetric pattern would leave a signal. What we have shown is that the architecture the disease name describes contains a regime that bulk DNA cannot reach even in principle — enough to disqualify the endpoint, not enough to say “bulk WGS can never see MVA”.

And **bulk averaging erases variegation by construction**, at any depth. The cytogenetic hallmark of MVA, premature chromatid separation, leaves no trace in DNA sequence at all.

Every candidate signal we did observe traced to coverage, to segmental duplications (22q11, centromeres), or partially to GC content — partially, because our measured GC-coverage correlation was $r = 0.43$ from a sparse sample, and chr21, the largest deviation, is GC-poor. That is an insufficient measurement rather than an established explanation, and we report it as such.

What we do not claim about this child, and why

We ran the calibrated statistic on the proband and **we withdrew the result**. It is not reported here as a negative finding, and the improved limit must not be read as one either. Three reasons, all from our own files:

- **We calibrated one test and applied another.** The pre-registered Ψ threshold is exceeded on **22 of 22 autosomes** at every κ setting. That is not 22 mosaic events; it is a null model that does not describe these data. The between-chromosome robust z-score we actually reported was chosen after seeing the data and was never calibrated (mosaico/RESULTADOS.md §5).
- **The residual signal is uncontrolled.** What remains is a global overdispersion of $\approx 2.3\times$ the binomial expectation, shared by all 22 autosomes and correlated at the ~ 200 kb scale. From the VCF alone we cannot say whether it is technical or biological.

- **The external control was pre-registered, it ran, and it came back not conclusive.** Ten unrelated 1000 Genomes individuals on two fixed windows (`mosaico/PREREGISTRO.md` amendment 4, committed `6405d9e`; result `ce94b89`). The proband fell **below** the control envelope on both statistics and both chromosomes — an outcome the amendment had not listed, so the pre-registered verdict is *not conclusive*. The cause was declared in advance with its direction: the controls carry no GQ filter and were called jointly across 3,202 samples against a singleton call for the proband, which inflates their dispersion. A κ of 46 in a healthy individual is a property of the calling, not of biology. **The control did not fail to run; it failed to be a control** (`mosaico/CONTROL_RESULTADO.md`).

One narrow statement survives, and we make only it: against those ten genomes, window-scale overdispersion ranges from $\kappa \approx 6$ to 46, and the proband's is the **lowest in the comparison** — 2.74 on chr8 against a control range of 9.5–46.0, and 4.60 on chr17 against 6.1–15.0 (`mosaico/control_resultado.json`; these are the thinned window statistics the comparison used, not the 2.27× chromosome-wide excess reported in `mosaico/RESULTADOS.md`) — so, with the caveat that the comparison itself was judged not conclusive, a κ of 2–5 is not large by the standards of these controls. **That is not a clinical negative.** We are not reporting that this child has no clonal lesion above 2% of cells; we are reporting that our attempt to measure it did not produce an interpretable answer. Resolving it needs a control processed identically from reads, which we did not have.

Consequence: no proposal whose endpoint is “reduced aneuploidy burden” can be evaluated with these data. Any screen must measure per-cell mis-segregation.

6. The go/no-go experiment — an eight-week assay window

Scoped to be executable by one laboratory in one window, not as a grant application.

Question. Are BUB1B-deficient patient cells preferentially sensitive to bortezomib at clinically achievable concentrations, and does that sensitivity come without increasing mis-segregation among the survivors?

What this experiment can and cannot decide, stated first. Dermal fibroblasts from this child are *constitutional* MVA cells, not tumour cells. A sensitivity difference between them and unaffected controls therefore measures whether the proteotoxic dependency transfers from cancer aneuploidy to constitutional aneuploidy at all — it does **not** measure tumour selectivity, and on its own it is as compatible with systemic toxicity as with therapeutic promise. That is why the go criterion below is *not* “patient cells are more sensitive”. A negative therefore **fails to support transfer without closing the tumour-board question**, which would need aneuploidy-high tumour-derived material. We say this now rather than after the data.

The quantity that governs this design is the fraction of cells carrying *any* aneuploidy, not the fraction carrying a *given* chromosome. A cell is under proteotoxic load whichever chromosome it gained, so the ~30% of the diagnostic criterion is the right denominator here, while the 1–2% per chromosome is the right one for §5. A pre-registered power calculation (`potencia/PREREGISTRO.md`, committed before the run) gives **power 0.897** at 30% of cells with an assumed 4× EC50 selectivity, $n = 3$ and $\alpha = 0.05$, against **power 0.063** under the per-chromosome reading. An earlier version of this section used the per-chromosome figure and concluded a negative would be ambiguous by design; that hedge rested on the wrong denominator and is withdrawn as a statement about statistical power.

Read 0.897 as an idealised scenario, not as this experiment’s power. It holds *for an assumed 4× selectivity, in a ~30% aneuploid culture, at 10% well-level CV, and none of those three has been measured in these cells*. The simulation is also weaker than the design we pre-registered: it fits a two-parameter Hill curve where the pre-registration promised four, and its noise model has no biological-replicate term. Both push the number upward, so 0.897 is optimistic **within its own model** — it is neither a validated power estimate for

patient cells nor a mathematical ceiling on what a better-specified version of the experiment could achieve (potencia/RESULTADOS.md §5b).

A second pre-registered analysis the same day gives us our own reason to doubt the 4×. Proteotoxic load is a dose, not a switch: an MVA cell carries **about one** altered chromosome, two arms, against the eighteen-arm point at which we evaluated the dose response — close to the mean of the solid DepMap stratum (potencia/dosis.json, arms_mean = 18.6), which is its middle and not its top. We have not counted arms in the specific lines that gave the < 40 nM EC50; the eighteen is the high end of the range over which we measured the dose response, not a measurement of those lines. In DepMap, in the one lineage where the aneuploidy–proteasome dependency is significant at all, PSMB5 dependency moves **0.09 residual SD at two arms against 0.84 at eighteen** (potencia/RESULTADOS_2_DOSIS.md, pre-registered as 45c4ed7).

State precisely what that does and does not do, because we got this wrong once. The pre-registration fixed its magnitude rule on the **solid** stratum. There the slope's 95% interval covers zero (−0.0118 ... +0.0006, $p = 0.076$), so the falsification clause applies and the exercise is **uninformative** for the stratum the rule was written about. The 0.093 ratio is **haematological**, where no decision rule was pre-registered, and CRISPR gene effect does not convert into an EC50 ratio in either stratum. So the 4× central value is **unsupported for these cells and equally unrefuted**. An earlier version of this section reported the rule as having fired; it did not, and the correction is dated 2026-09-06 in RESULTADOS_2_DOSIS.md §3. At an assumed 2× selectivity power is **0.387**. The honest summary is that **we do not know which regime these cells are in, and the experiment is what would tell us**.

It still buys a protocol change, not just a corrected sentence. Power collapses below a cell fraction of about 0.25 (0.589 at 0.20, 0.196 at 0.10), and **nobody has counted metaphases in these fibroblasts** — culture selection makes the true fraction likely lower than at biopsy. So the fraction becomes a **gating measurement** rather than an assumption, taken on the culture that will actually be treated, at the passage that will be treated. The endpoint is a **count of cells carrying at least one numerical chromosome abnormality**, recorded together

with **how many abnormal chromosomes each abnormal cell carries**, because the second analysis shows the per-cell burden matters as much as the fraction.

PCS scoring is a companion, not this measurement. Premature chromatid separation is a centromere-cohesion phenotype and the diagnostic criterion for MVA; the aneuploid fraction is a count of abnormal chromosome *numbers*. The same metaphase slides yield both, at no extra cost, but they are different endpoints and an earlier version of this section conflated them.

Indicative branches: $n = 3$ if the fraction is ≥ 0.25 ; $n = 14$ at 0.10 ; and **below 0.10 do not run the bulk viability assay** — go to a per-cell endpoint, because under our assumptions the required n grows past what a biopsy-derived culture can supply. Every branch is computed at the same assumed $4\times$ selectivity, so **measuring the fraction does not validate the branch it selects**: it removes one unmeasured input of three. Treat the table as planning, not as a validated operating rule. Details and every sweep in `potencia/RESULTADOS.md`.

Consent, assent and ethics approval come before week 1, and they are not ours to give. Taking a skin biopsy from a child, establishing a fibroblast line from it and retaining that line require informed parental consent, the child's assent at an age-appropriate level, and approval by the research ethics committee or institutional review board of whichever centre runs this. We are an analysis team with no access to the patient, no relationship with the family and no institutional sponsor for this work; nothing below is authorised by us, and neither the approval nor how long it takes is within our control. The same applies to the parental sampling in `clinico/PARENTAL_SEGREGATION_ORDER.md`, which sets out its own consent requirements.

Week	Step	Readout
1–3	Patient dermal fibroblasts from skin biopsy, plus two matched controls. Gate: on the culture to be treated, count metaphases scoring (a) the proportion of cells with ≥ 1 numerical chromosome abnormality, (b) the number of abnormal chromosomes per abnormal cell, and (c) PCS, reported separately. (a) is the input the power model needs and it is unmeasured; (c) is the MVA diagnostic criterion and is not a substitute for it (potencia/RESULTADOS.md §3). Patient cells are required: RPE1 with reversine-induced aneuploidy reproduces Ippolito's own system and cannot test whether the dependency transfers to BUB1B-deficient constitutional MVA — we list it as a positive control, not a substitute. Without patient cells the go/no-go question is not answerable	Growth; baseline karyotype
2–4	Allele fate: allele-specific RT-PCR $\pm$ NMD inhibitor; quantitative western blot with N- and C-terminal antibodies	An N+/C– band pattern is consistent with a stable truncated protein and loss of both with NMD — the RT-PCR arm decides; the blot alone does not
4–6	Bortezomib dose–response, 0–1,000 nM with dense sampling below 100 nM, 72 h, patient versus control versus an aneuploidy-high positive control (reversine-treated RPE1)	EC50 and the ratio between them. The range spans the 312 nM clinical Cmax so the advancement threshold falls inside the data. Extended from 0–400 nM after the power calculation: at a selectivity of 10× the control's EC50 sits at the old ceiling and is estimated at the edge of the data, doubling the spread of the

Week	Step	Readout
		estimate (potencia/ RESULTADOS.md §4)
4–6	Combination arm: bortezomib + everolimus	Tests the predicted antagonism
6–8	Micronucleus assay (OECD TG 487) on survivors of every condition, plus FISH for 3–5 chromosomes	Micronucleus frequency as a proxy for mis-segregation (filter 5); FISH separates whole-chromosome loss from breakage, and proliferation is recorded because a change in division rate alone moves the count

Timeline, stated honestly in the heading it belongs in. Eight weeks is the *assay* programme, and it holds only once a patient fibroblast line exists. Establishing $15\text{--}20 \times 10^6$ fibroblasts from a fresh biopsy typically takes 4–8 weeks on its own (track2/DATOS_CANDIDATO.md §7), so from a standing start the planning figure is **12–16 weeks**, not eight — and the same source warns that growth can be slower in MVA1 fibroblasts, which would push it further. We use “eight-week experiment” as shorthand for the assay window and nowhere as a promise of an answer in eight.

Advancement criteria, fixed before the data. We state them in the direction that avoids rewarding toxicity:

1. **A dependency signal, not merely a sensitivity difference:** an EC50 shift between patient and control fibroblasts whose 95% confidence interval excludes 1, **of at least the magnitude that a culture of the measured aneuploid fraction can produce.**

We had this criterion wrong and an adversarial review caught it. It previously demanded a shift “at least equal to the near-euploid-versus-highly-aneuploid contrast in Ippolito Fig. 6o”. That contrast is between two **homogeneous** populations. A patient fibroblast culture is a **mixture**, and under our own model a mixture cannot produce it: with 30% of cells aneuploid and a $4\times$ selectivity *within* those cells, the **bulk** EC50 ratio is **$1.50\times$** , not $4\times$. The old criterion was unreachable by construction — the experiment would have failed it even if the biology worked exactly as predicted.

The threshold must therefore be derived from the gating measurement, not imported from a homogeneous-population figure. At a 4× within-cell selectivity, the bulk ratio our model predicts is:

Measured aneuploid fraction f	Bulk EC50 ratio to expect	Max viability gap
0.10	1.14×	0.03
0.20	1.30×	0.07
0.30	1.50×	0.10
0.50	2.00×	0.16

Read the criterion off this table once f is measured, and treat a shift materially below the corresponding row as a negative. A shift materially *above* it is not a better result — it is a reason to suspect the mixture model or systemic toxicity, which is what stop signal 4 below exists for. Every row assumes the same unmeasured 4× within-cell selectivity, so this table is a pre-specified reading rule for the experiment, **not a validated threshold**. If the experiment ever returns a selectivity estimate, the table is recomputed rather than defended.

2. **Effect at ≤ 312 nM.** That is the top of the *total plasma* C_{max} interval for the approved dose, used here as an order-of-magnitude ceiling for plausibility. It is not free drug, not intratumoral drug, and not a validated clinical threshold.
3. **No increase in micronuclei among survivors** — filter 5.
4. **A stop signal, not a go signal:** if patient fibroblasts are markedly more sensitive than controls *without* the aneuploidy-high positive control showing a larger shift still, the most likely reading is constitutional toxicity in an MVA patient rather than selective killing of aneuploid cells. That outcome stops the programme too, and it is the outcome we would most easily have mistaken for success.

Failing any of 1–3, or triggering 4, stops the programme.

Falsifying results, stated at the precision the design supports. If patient fibroblasts are no more sensitive than controls, the experiment **fails to support transfer of the dependency into these cells** — it

does not falsify a dependency in aneuploidy-high tumour material, which is the setting §4 actually proposes and which this design does not sample. That distinction is the difference between “we did not see it here” and “it is not there”, and we hold ourselves to the first. If bortezomib raises micronucleus frequency among survivors, that **does** falsify the proposal outright: it fails filter 5 regardless of how well it kills.

The stage this design does not contain. Deciding tumour selectivity requires aneuploidy-high tumour-derived material with a near-euploid comparator, which no eight-week window starting from a skin biopsy can supply. We treat the fibroblast arm as the gate that decides whether that second stage is worth building, not as a substitute for it.

This design can fail, which is the point of proposing it.

7. Actionable now, without any drug

The intervention that most changes this child’s prognosis today is not a molecule.

There is a published consensus for MVA specifically, and it should be the anchor. The 2024 update from the AACR Childhood Cancer Predisposition Workshop covers mosaic variegated aneuploidy by name and states: *“This group recommended renal ultrasound surveillance every 3 months from birth until age 7, for all MVA conditions including those with an unknown genetic cause”,* together with *“regular clinical assessment including review of systems to identify signs of rhabdomyosarcoma and other malignancies”,* avoidance of radiation exposure, and HPV vaccination (PMID 39264246). The same source records that cancers in MVA1 are *“predominantly Wilms tumor and rhabdomyosarcoma, but MDS, AML and ALL have also been reported”*.

Provenance note, 2026-09-06. These three passages — the 3-monthly-to-age-7 interval, the tumour spectrum, and the statement in §9 about the 15 published MVA2 individuals — come from the **full text** of PMID 39264246, not its abstract. On 2026-09-06 we re-verified all three against the open-access full text (PMC11705613, archived verbatim in `evidencia/consenso_fulltext_2026-09-06.xml`): the interval is

attributed there to the SIOP-Europe Host Genome Working Group and SIOP Renal Tumor Study Group, with the AACR group stating that it *supports* those recommendations; the MVA2 sentence reads “*at the time of writing, none of the 15 individuals with MVA2 described in literature developed cancer*”, which describes the published literature, not a followed cohort, and is not a statement of zero risk. Nothing in our recommendation is original to us here; a clinician should read the consensus itself before acting.

We had reached the 3-monthly-to-age-7 schedule independently, by extension from UK Wilms-risk recommendations that stop at age 5 (PMID 16857697). We report that convergence rather than our derivation: **the schedule below is consensus guidance, not our inference**, and a family should be given it as such.

Where the consensus leaves a gap, and it is the gap that matters for this child. The window closes at seven. A 2026 case report describes an MVA3 girl whose Wilms tumour appeared at three and whose orbital embryonal rhabdomyosarcoma appeared at **twelve** (PMID 42595739) — five years past the end of renal-ultrasound surveillance, at a site renal ultrasound does not image. Renal ultrasound is a kidney-directed test; rhabdomyosarcoma is not a kidney tumour, which is why the consensus pairs it with clinical assessment rather than with imaging. For a child who has *already had* an embryonal rhabdomyosarcoma, we therefore propose that the clinical-examination arm continue past seven, and we mark it as our extension rather than as guidance.

Surveillance windows must be anchored to the child’s current age, which is not in the dataset and must be confirmed before use. If the age is unknown, that is the first thing to establish — nothing below is actionable without it. This is a discussion framework for an expert cancer-predisposition clinic, not a substitute for one:

Period	Surveillance	Frequency
Birth to age 7	Renal ultrasound (we suggest extending the field to liver, retroperitoneum and pelvis)	Every 3 months — consensus recommendation for all MVA conditions (PMID 39264246). The wider abdominal field is our addition
To age 10, then reviewed	Full paediatric and oncological examination: abdomen, nodes, head and neck, orbit , skin, genitalia, soft tissues	3-monthly to age 7, then 6-monthly. The consensus asks for clinical assessment for rhabdomyosarcoma without setting an end date; continuing past 7 is our extension , prompted by the age-12 orbital ERMS in PMID 42595739. Catches superficial ERMS that renal ultrasound cannot
Post-ERMS	Relapse follow-up per tumour protocol	Coordinate to avoid duplicate imaging
Through childhood	Full blood count with differential	6-monthly is reasonable given reported leukaemia, but no evidence of benefit exists
Lifelong	Family education; annual predisposition clinic	Urgent review for mass, haematuria, abdominal distension, persistent pain, neurological deficit, bleeding or weight loss

If the child is already past 7, missed scans should not be “caught up”: a baseline assessment and an individualised decision by the predisposition team is the right course.

We do **not** propose serial AFP, repeat CT, PET-CT or routine whole-body MRI while asymptomatic — no demonstrated benefit in MVA, weighed against sedation, false positives, cost and radiation. The consensus is explicit that **radiation exposure should be avoided** in these disorders (PMID 39264246), which is an argument against imaging-heavy surveillance and, separately, a consideration for any future treatment plan.

One implication for the parents, which the phenotype document already hints at. The clinical document lists recurrent spontaneous abortion (HP:0200067) — an obstetric history, not a feature of the child.

If the two variants are in *trans* — which is the interpretation this report adopts but does not demonstrate (§10) — each parent would be expected to carry one of them; segregation testing is what would establish that, and it should precede any carrier-specific interpretation. A 2026 study of two unrelated families with unexplained recurrent pregnancy loss identified novel heterozygous *BUB1B* variants and found **significantly elevated premature chromatid separation rates in the carriers' lymphocytes** by G-banding and centromere FISH, with a trend toward reduced BUBR1 (PMID 42434306). Those variants were classified VUS and the cohort is two families, so this is a hypothesis, not a finding about this family. But it is cheap to act on and it is squarely in the family's interest: a PCS assay on parental lymphocytes and formal genetic counselling would test whether the parents show the cellular phenotype those carriers showed — it would not establish that it caused the pregnancy losses — and would inform any future pregnancy. It is cheap, it can be done this month **alongside** the consensus surveillance above, and it replaces no part of it.

We have written the request out so that it does not have to be reconstructed. `clinico/PARENTAL_SEGREGATION_ORDER.md` is a one-page laboratory request a clinician can adapt and sign: both loci, two Sanger amplicons with primer sequences, and an interpretation table written **before** the result, covering all five outcomes — including the *cis* result, which would refute our own Track 1 interpretation outright. The primers were designed against GRCh38 with common 1000 Genomes SNVs and RepeatMasker intervals excluded from the primer regions, and each of the four was then counted across the entire assembly: all occur exactly once. **The RepeatMasker exclusion is what caught the bad pair** — the first design the software returned sat inside an *AluY* element 272 bp from the *c.2210* site, with roughly a hundred exact copies in the genome. The genome-wide count is **confirmation, not the filter**, and we state it that way rather than claiming we sifted a bad pair out at the counting step (`clinico/PARENTAL_SEGREGATION_ORDER.md`). Either way the failure it avoids is real: a primer like that does not fail in review, it fails in the laboratory, and in a segregation test allele drop-out reads as “*this parent does not carry it*”.

What segregation can and cannot deliver. A *trans* result supports the compound-heterozygous model and contributes PM3-type evidence. **It does not by itself reclassify the missense as pathogenic**, and we do not treat it as doing so. A *cis* result is the informative one in the other direction: it would refute the Track 1 interpretation outright.

This is the cheapest experiment in the project that speaks to phase, to a future pregnancy and to HP:0200067 at once, and it needs blood rather than a new laboratory.

8. What we do not propose

- **MPS1/TTK inhibitors and reversine** — not approved, and they weaken a checkpoint that is already insufficient. Reversine is a research tool for *inducing* aneuploidy, the opposite of the goal.
 - **Apcin, proTAME** — APC/C tool compounds, not marketed.
 - **DCZ0415** — experimental TRIP13 inhibitor: wrong gene, not approved.
 - **AICAR and 17-AAG** — genuine hits in the Tang screen, neither marketed. Useful as in vitro positive controls only.
 - **Metformin** — rejected at filter 3 as unevaluable on the mechanism proposed for it (§3.1).
 - **Senolytics** — the BubR1 senescence evidence rests on *genetic* clearance of p16^{Ink4a}⁺ cells via the INK-ATTAC transgene (PMID 22048312), which is not a transferable drug; no agent is approved with a senolytic indication; and senescence is an anti-tumour barrier worth preserving in a cancer-predisposition syndrome.
 - **mTOR inhibitors combined with proteasome inhibition** — predicted antagonism (PMID 31530568).
 - **Bortezomib for this child today** — fails filter 4, as set out in §4.
-

9. Scalability

- The **five-filter framework**, particularly filters 3 and 5, applies *as questions* to any repurposing exercise in a rare disease: filter 3 is an exposure plausibility screen, not a universal criterion, and filter 5's mis-segregation form is the MVA instance of the general question

does treatment worsen the disease-relevant defect or select for it? It is the transferable product of this work.

- **The MVA genes are not interchangeable.** MVA2 arises from biallelic *CEP57* (PMID 21552266) and behaves differently from MVA1 and MVA3 on exactly the axis that matters here — see the *CEP57* caveat below.
- The **proteotoxic vulnerability**, if §6 shows it transfers, would extend most plausibly to **TRIP13**-related MVA, which shares both the tumour risk and the severity of checkpoint failure: “individuals with biallelic *TRIP13* or *BUB1B* mutations have a high risk of embryonal tumors ... their cells display severe SAC impairment” (PMID 28553959) — and the 2026 case in §7 is a *TRIP13* patient with two embryonal tumours (PMID 42595739). **It should not be assumed for CEP57.** Yost reports that *CEP57*-related MVA “is not associated with embryonal tumors” with cells showing “minimal SAC deficiency”, and the 2024 consensus is blunter still: **“At the time of writing, none of the 15 individuals with MVA2 described in literature developed cancer”** (PMID 39264246). Less mis-segregation should mean less proteotoxic burden, so the rationale is weaker there. No proteasome data exist for either, and §10 states that transfer to MVA1 itself is unproven.
- **The target to watch is UBE2H.** The 2026 paired CRISPR screens nominate it as a top aneuploid-selective dependency linked to mitochondrial proteostasis (PMID 42094535, preprint). It has no approved inhibitor, so it fails filter 1 today and is not a candidate — but it is where an aneuploidy-selective agent would most plausibly come from next, and a repurposing exercise repeated in two years should start there rather than at metformin.
- **A reusable worksheet** is included as the Appendix: the five filters as blank rows, with the metformin and bortezomib verdicts as the worked example, so a team working on a different rare disease can apply the framework without reading this report.
- **The framework was transferred prospectively to the other two MVA genes, and it gives different answers.** `report/TRANSFER_MVA2_MVA3.md` holds the worksheet filled for MVA2 (*CEP57*) and MVA3 (*TRIP13*) with the same candidate and the filters held fixed. For **MVA2** bortezomib **fails filter 4 and cannot be evaluated at filters 2–3**: no proteasome measurement exists in

CEP57 cells, and the setting that would invert the safety filter — an active malignancy — **has not been reported** in the 15 published MVA2 individuals (PMID 39264246). That is a different failure from MVA1, where the same filter fails but the inverting condition demonstrably exists. It stays **conditional for MVA3**, which shares both the embryonal-tumour risk and the severe checkpoint failure. A framework that returns the same verdict for three diseases is a slogan; one that returns different verdicts for MVA2 and MVA3 shows that the worksheet responds to the evidence available for each gene — a demonstration of the framework, not evidence that the drug behaves differently in them.

- **The retrieval benchmark is packaged as a tool, not described as one.** `replay/mva_replay.py` takes a candidate gene, two alleles, an HPO set and a background VCF, and returns the planted score, the rank, the ClinVar-whitelist-off counterfactual and an unrelated-HPO control. It is a single file with no dependencies beyond what an Exomiser installation already needs, and a team working on a different rare disease can spike their own gene into a public genome in an afternoon. Its built-in check that a control HPO term is not annotated to the planted gene is what caught the contamination described in `hpo/RESULTADOS.md` — in our own control set.
- **And we tested that sentence rather than leave it as a claim** (`replay/TRANSFER_TRIP13.md`, pre-registered and pushed before the run). From a fresh clone of this repository, the tool planted two public Likely_pathogenic *TRIP13* alleles — chosen by the frozen pool rule, not by us — into three healthy GIAB genomes and queried them with six HPO terms transcribed from the abstract of the 2026 *TRIP13* case (PMID 42595739). *TRIP13* came back **rank 1 in 3 of 3 backgrounds at 0.8332**, absent or scoring 0 unspiked, and rank 1 / 2 / 3 under an unrelated phenotype — the same pattern as BUB1B. Each background took under 3½ minutes and about 2.4 GiB. The operational half was the useful one: the clone needed **four provisioning steps, three of which the tool's requirements list did not include**, and the regression tests the README said ran from a bare checkout did not. Both are fixed; the run that found them is in the repository: the provisioning script, and the log of every tool invocation with its exit code.

- The **gene panel** (BUB1B, CEP57, TRIP13, CENATAC, MAD1L1, MAD2L1BP, CEP192, BUB1, SMC5, TRIM37, CENPE) and the artefact controls used in the mosaicism analysis transfer to any MVA or PCS workup.
 - The **micronucleus → scDNA-seq endpoint** applies to mitotic CIN disorders. It does **not** transfer to Fanconi anaemia or Bloom syndrome, whose diagnostic assays are DEB/MMC chromosome breakage and sister-chromatid exchange respectively — a different lesion, not substitutable.
-

10. Limitations

- **Phase is not established, and we measured that it is not establishable from these data.** There are no parental samples; the challenge distributes a VCF and raw reads but no alignments, we did not align the reads, and short reads could not have phased these two variants in any case — the only heterozygous site between them lies 6.8 kb from one and 4.1 kb from the other, beyond any short-read fragment, so read-backed phasing would need long or linked reads; the two variants sit **10,911 bp** apart and the caller emitted no phase tag for either. Reference-panel phasing fails for a separate and quantifiable reason: inside that interval the child carries exactly three heterozygous sites — the two causal variants and `chr15:40216470` — and **all three are absent from the 1000 Genomes 30× phased panel** (3,202 individuals); the nearest panel-present heterozygous markers are 16.8 kb upstream and 809 bp downstream, with nothing between them to link the alleles. A Beagle 5.5 run against that panel keeps 1,391 of 1,452 window variants and drops all three that matter. The absence is what the frequencies predict: expected copy counts in 6,404 haplotypes are 0.64 for the nonsense and **0.006** for the missense. Two heterozygous variants in a singleton could be cis or trans, nothing in the sequence data weighs one over the other, and rarity does not phase them. What supports *trans* is that the child is affected, that MVA1 is recessive, that truncating-plus-missense is the recurrent pattern, and that the phenotype matches. **It remains a diagnosis pending segregation, and parental genotyping at the two loci is the cheapest experiment in this report.**

- **The retrieval pipeline is itself insensitive to phase.** In a controlled test on public data, Exomiser 15.1.0 reads phase into its ACMG evidence (trans → PM3, cis → BP2) but not into the compound-heterozygous model: two pathogenic nonsense alleles declared on the *same* chromosome are still called AR_COMP_HET, still classified PATHOGENIC and still ranked first, at an absolute combined-score difference of 0.0034. Nothing in the Track 1 retrieval would have flagged a *cis* configuration.
 - **p.Asn1002Lys has no functional assay**, and its apparent ClinVar entry is for a different nucleotide.
 - **The bortezomib EC50 is a conservative upper bound** read from a figure panel rather than a tabulated value, and the Cmax is a label interval rather than a point estimate; the ratio of 5.8–7.8 should be read as an order-of-magnitude statement computed on total, not free, drug.
 - **Ippolito et al. studied cancer aneuploidy**, not constitutional variegated mosaicism. Whether the proteotoxic dependency transfers is precisely what the proposed experiment tests; we have not assumed it.
 - **No trial specific to MVA1** was found in ClinicalTrials.gov, Open Targets, ChEMBL or CMap/LINCS. Absence of an interface result is not proof of absence.
 - **The child's current age is unknown to us**, and the surveillance schedule depends on it.
-

A closing note

The family of this child published their son's genome so that strangers might help. They will read what comes back.

We have tried to write something useful rather than something reassuring. The honest answer to “what drug should he take today” is none — and we show why the obvious candidate cannot work by the mechanism proposed for it, rather than proposing it softly. The useful answer is a mechanism that matches his disease, an approved drug whose plasma levels are not excluded by the concentrations at which cancer cells respond, an explicit statement of the condition under which it would become appropriate, and an experiment — an eight-

week assay window, 12–16 weeks from a fresh biopsy — that could fail to support it in this child’s own cells. That experiment would not close the tumour-board question on its own, and §6 says so.

That is not a cure. It is a prepared answer to a question this family may unfortunately have to ask.

Appendix — The five-filter worksheet (reusable)

The transferable product of this work. Blank rows for a new disease; the two decided rows as worked examples. A candidate must pass all five. Filters 3 and 5 eliminate the most candidates and are the two most often skipped.

Candidate	1. Approved?	2. Addresses lesion or a direct consequence?	3. Cmax ≥ effective concentration?	4. Safe in <i>this</i> patient?	5. Does not increase mis-segregation among survivors?	Verdict
Metformin	Yes (paediatric label from age 10)	Proposed as the marketed stand-in for the AICAR/ AMPK hit	Not evaluable for aneuploidy selectivity — no such EC50 exists; and ~1000× short for the complex I mechanism proposed (μM plasma vs mM required)	Renal clearance; nephrocalcinosis unassessed	Not reached	Rejected at filter 3 as unevaluable on the proposed mechanism
Bortezomib	Yes	Yes — proteasome dependency of aneuploid cells	Yes on total drug: 231–312 nM Cmax vs EC50 <40 nM (5.8–7.8×); IV only, not subcutaneous	No — neuropathy in 18% of children, on pre-existing muscle atrophy	Untested — no measurement exists; the micronucleus arm of §6 is where it would be tested	Fails filter 4 for this child; filter 5 untested; conditional for a patient with an active tumour
<i>(your candidate)</i>						
<i>(your candidate)</i>						

How to fill it in. Filter 3 needs a real denominator: an effective concentration measured in a model stratified for the disease mechanism. If none exists, write *not evaluable* — do not substitute an IC50 from an unrelated model, because that manufactures a ratio with no meaning. Filter 4 is patient-specific by construction, so a “fail” here is a statement about one person and one moment, not about the drug. Filter 5 asks the question repurposing exercises forget: among the cells that survive the drug, is the underlying defect made worse?

References

1. Hanks S, et al. Constitutional aneuploidy and cancer predisposition caused by biallelic mutations in *BUB1B*. *Nat Genet* 2004. PMID 15475955
2. Suijkerbuijk SJE, et al. Molecular causes for BUBR1 dysfunction in mosaic variegated aneuploidy. *Cancer Res* 2010. PMID 20516114
3. Hanks S, et al. Comparative genomic hybridization and BUB1B mutation analyses in childhood cancers associated with MVA. *Cancer Lett* 2006. PMID 16182441
4. Tang YC, et al. Identification of aneuploidy-selective antiproliferation compounds. *Cell* 2011. PMID 21315436
5. **Ippolito MR, et al. Increased RNA and Protein Degradation Is Required for Counteracting Transcriptional Burden and Proteotoxic Stress in Human Aneuploid Cells.** *Cancer Discov* 2024. PMID 39247952
6. **Sakellakis M. Why Metformin Should Not Be Used as an Oxidative Phosphorylation Inhibitor in Cancer Patients.** 2023. PMID 37343530
7. Fontaine E. Metformin-Induced Mitochondrial Complex I Inhibition: Facts, Uncertainties, and Consequences. 2018. PMID 30619086
8. Ailabouni AS, et al. Interindividual variability in metformin pharmacokinetics in pediatric patients. 2026. PMID 42399684
9. **Chui MH, et al. Chromosomal Instability and mTORC1 Activation through PTEN Loss Contribute to Proteotoxic Stress in Ovarian Carcinoma.** *Cancer Res* 2019. PMID 31530568
10. Knouse KA, et al. Single cell sequencing reveals low levels of aneuploidy across mammalian tissues. *PNAS* 2014. PMID 25197050
11. Baker DJ, et al. BubR1 insufficiency causes early onset of aging-associated phenotypes. *Nat Genet* 2004. PMID 15208629 — cited for the progeroid phenotype of BubR1 insufficiency underlying §3.4
12. Baker DJ, et al. Clearance of p16^{Ink4a}-positive senescent cells delays ageing-associated disorders. *Nature* 2011. PMID 22048312
13. Sieben CJ, et al. BubR1 allelic effects drive phenotypic heterogeneity in mosaic variegated aneuploidy. *J Clin Invest* 2020. PMID 31738183
14. Nishitani-Isa M, et al. ERMS with PCS/MVA treated with reduced-intensity chemotherapy. *Pediatr Int* 2019. PMID 31184400
15. Scott RH, et al. Surveillance for Wilms tumour in at-risk children: pragmatic recommendations. *Arch Dis Child* 2006. PMID 16857697
16. Xu P, et al. BUBR1 recruits PP2A via the B56 family to promote chromosome congression. *Biol Open* 2013. PMID 23789096

17. Kruse T, et al. Direct binding between BubR1 and B56-PP2A regulates mitotic progression. *J Cell Sci* 2013. PMID 23345399
 18. Snape K, et al. Mutations in *CEP57* cause mosaic variegated aneuploidy syndrome. *Nat Genet* 2011. PMID 21552266
 19. Yost S, et al. Biallelic *TRIP13* mutations predispose to Wilms tumor and chromosome missegregation. *Nat Genet* 2017. PMID 28553959
 20. FDA label, VELCADE (bortezomib), NDA 021602 s040, §12.3
 21. FDA BPCA Clinical Review, bortezomib paediatric safety, NDA 021602, 2015
 22. Aneuploid human colonic epithelial cells are sensitive to AICAR-induced growth inhibition through EGFR degradation. 2013. PMID 22890317
 23. Nakano Y, Kuiper RP, et al. **Update on Recommendations for Cancer Screening and Surveillance in Children with Genomic Instability Disorders.** *Clin Cancer Res* 2024. PMID 39264246 — the MVA-specific surveillance consensus used in §7
 24. Schukken KM, et al. **Paired CRISPR screens identify mitochondrial metabolism and UBE2H as aneuploid-specific dependencies in human cancer cell lines.** *bioRxiv* 2026 (preprint, not peer-reviewed). PMID 42094535
 25. Vempuluru VS, et al. Sequential presentation of Wilms' tumor and orbital rhabdomyosarcoma in a child with mosaic variegated aneuploidy syndrome 3. *Orbit* 2026. PMID 42595739
 26. Wei TY, et al. Functional and clinical evidence for two novel heterozygous *BUB1B* variants and their value in precision genetic counseling for recurrent pregnancy loss. *Front Endocrinol* 2026. PMID 42434306
 27. Escalante LE, Hose J, et al. Chromosome duplication causes premature aging via defects in ribosome quality control. *PLoS Biol* 2025. PMID 41248159
 28. Bersani F, Taulli R, et al. Bortezomib-mediated proteasome inhibition as a potential strategy for the treatment of rhabdomyosarcoma. *Eur J Cancer* 2008. PMID 18342500
 29. Houghton PJ, Morton CL, et al. Initial testing (stage 1) of the proteasome inhibitor bortezomib by the pediatric preclinical testing program. *Pediatr Blood Cancer* 2008. PMID 17420992
 30. Blaney SM, Bernstein M, et al. Phase I study of the proteasome inhibitor bortezomib in pediatric patients with refractory solid tumors: a Children's Oncology Group study (ADVL0015). *J Clin Oncol* 2004. PMID 15570082
 31. Muscal JA, Thompson PA, et al. A phase I trial of vorinostat and bortezomib in children with refractory or recurrent solid tumors: a Children's Oncology Group phase I consortium study (ADVL0916). *Pediatr Blood Cancer* 2013. PMID 22887890
 32. Maki RG, Kraft AS, et al. A multicenter Phase II study of bortezomib in recurrent or metastatic sarcomas. *Cancer* 2005. PMID 15739208
-

Acknowledgement

This work was made possible through the Hackathon, organized by Sage Bionetworks in partnership with the MVA Society, Hugging Face, and BEACON (The Benchmarking, Evaluation, and Assessment Consortium for Science), with prize sponsorship from AWS and Anthropic. We are deeply grateful to the child and their family who

generously contributed their data and their story to advance research into this rare disease. We acknowledge their trust in making this Hackathon possible.

Data availability

The individual-level data analysed here is the Hackathon dataset (SageBio/mva-hackathon-2026-data), accessed under the challenge data transfer agreement accepted by Hugging Face account `cpu-16` on 2026-08-28, and cited as directed on the Hackathon Synapse page. **No patient sequence data or genotype-scale derivative is included in this document or in our public repository**, and every copy we hold will be deleted by 2026-11-23 with written confirmation to the organisers. The two diagnostic variants are named, as findings, which the challenge rules permit. All other inputs are public: Exomiser 15.1.0 with data release 2602, ClinVar, gnomAD v4, the Human Phenotype Ontology, Ensembl and VEP, the 1000 Genomes and GIAB genomes, DepMap 24Q4, AlphaFold and the literature cited by PMID. Analysis code, pre-registrations and results: <https://github.com/cpu-16/mva-hackathon-2026>

Appendix B — Methods description (Track 2 form)

Q2 — Describe your approach in detail (variant/mechanism → candidate medication)

We went from variant to candidate in five steps, and the third and fifth are where most candidates die.

Step 1 — establish the lesion. Compound heterozygosity in BUB1B: `c.2210T>G` p.Leu737Ter, a nonsense allele that removes the C-terminal 313 residues, including the kinase domain (UniProt O60566 annotates it at 766–1050), paired with `c.3006T>G` p.Asn1002Lys inside that domain. Loss of function, matching the truncating-plus-missense pattern of viable MVA1: a null allele with a hypomorphic one. Reading the pair as null-plus-hypomorph is our reading of the human truncating-plus-missense pattern (PMID 20516114), which does not state that complete biallelic loss is lethal, and neither do we; the

separate prediction that p.Asn1002Lys retains more residual function than L1012P comes from the Sieben mouse series (PMID 31738183), not from 20516114.

Step 2 — follow the mechanism downstream, not at the lesion. No approved drug restores BubR1. We therefore asked what the lesion *produces* that is druggable. The chain is: weakened spindle assembly checkpoint → chromosome mis-segregation → constitutional aneuploidy → stoichiometric imbalance of protein complexes → proteotoxic stress. The last link is the one with approved drugs against it.

Step 3 — five filters instead of the usual two. Repurposing exercises typically ask only whether a drug is approved and whether its mechanism is plausible. We added three:

#	Filter	Question
1	Regulatory	Approved for marketing?
2	Mechanistic	Does it address the actual lesion or its consequence?
3	Pharmacokinetic	Is $C_{max} \geq$ the effective in vitro concentration?
4	Patient-specific safety	Compatible with <i>this</i> child's comorbidities?
5	Direction of effect	Does it avoid increasing mis-segregation in survivors?

Step 4 — apply them, and report what dies. Metformin — the candidate the literature points to, descending from the AICAR hit of Tang et al. (PMID 21315436) — fails filter 3: therapeutic plasma levels are micromolar while complex I inhibition requires millimolar (PMID 37343530), and no published study gives metformin an aneuploidy-selective effective concentration to compare against. MPS1/TTK inhibitors and reversine fail filter 5 by construction: they weaken a checkpoint that is already insufficient.

Step 5 — the survivor, with its stopping condition. Bortezomib passes filters 1–3 and fails filter 4 *for this child*, which is a patient-specific verdict rather than a property of the drug. We state the condition under which it would become appropriate — an MVA patient with an active tumour, where the comparator is cytotoxic chemotherapy rather than nothing — and we propose an experiment

— an eight-week assay window, 12–16 weeks from a fresh biopsy because establishing the fibroblast line takes 4–8 weeks on its own — that could fail to support the dependency in this child’s own cells. It would not, on its own, close the tumour-board question, and we say so where we describe it.

Q3 — Generative AI declaration

Commercially-available generative AI was used. Provider, plan and relevant setting for each:

- **Anthropic**, Claude Opus and Claude Fable 5.1 via Claude Code, **Max** subscription plan, model-improvement data sharing disabled in account privacy settings.
- **OpenAI**, Codex (gpt-5.6), **Plus** plan, “Improve the model for everyone” disabled.
- **xAI Grok 4.6 via Cursor**, **Pro** plan, **Privacy Mode enabled** (zero data retention at Cursor and at the underlying model providers).
- **Qwen3-TTS 1.7B VoiceDesign** — open-weights, run locally on CPU with no cloud service and no API key. Used only to narrate the pitch video; not a commercial service, listed for completeness. Each clip was transcribed by a local recogniser and diffed against the written script before assembly, because this class of model can alter a word silently.

No block of the VCF and no genotype table was ever sent to a language model. Genomic work ran on local tools (`bcftools`, Exomiser) and on annotation APIs — Ensembl VEP and ClinVar — which return public annotation for submitted coordinates and acquire no rights over the input. The policy does not name these services; we read them as satisfying the two conditions it does set, namely no training on inputs or outputs and limited retention.

Q4 — Automated or manual candidate identification?

Primarily manual literature review and expert curation, with LLM-assisted retrieval. No drug-target database query produced the candidate. Interface queries against ClinicalTrials.gov, Open Targets, ChEMBL and CMap/LINCS returned no MVA1-specific drug pairing, which we report as a negative result rather than omit.

Q5 — Describe the manual review and curation

Three activities, all documented in the repository:

1. **Citation verification.** Every PMID was checked against PubMed or Europe PMC before use. This was not ceremonial: **three citation problems were caught, of three different kinds** — one wrong paper (a study of acromegaly cited as Wilms surveillance, PMID 17652220, replaced by 16857697); one merged attribution (two KARD papers under a single citation, separated into PMID 23789096 and 23345399); and one over-extension (PMID 28553959 was being used to carry the proteotoxic rationale to CEP57, which that same source contradicts — we withdrew that use and kept the paper only for the BUB1B/TRIP13 embryonal-tumour sentence its abstract actually states).
2. **Adversarial review by independent models.** Two different language models were asked to attack our own drafts. They found, among others, a claim of ours that contradicted **our own data file**: the report stated that Tang et al. used only stable trisomies, when our extraction record shows Figure 6 of that paper tests *Bub1b*^{H/H} hypomorphic MEFs. Corrected. All review rounds are in [evidencia/](#).
3. **Pharmacokinetic arithmetic, redone by hand.** Each C_{max} was converted from ng/mL to molar using the molecular weight, against the FDA label section cited. This surfaced that our own margin used the ceiling of a label interval (89–120 ng/mL) rather than the interval; corrected to 5.8–7.8×

Q6 / Q7 — Data sources

Publicly available drug, target and literature sources only; no proprietary data. The one controlled-access input is the challenge dataset itself (the proband's VCF and clinical document), used under the data transfer agreement described under *Data availability* and deleted within 30 days of the close.

- **Literature:** PubMed / Europe PMC, cited by PMID throughout.
- **Regulatory labels:** FDA labels and reviews — VELCADE (NDA 021602 s040; BPCA Clinical Review 2015), NDA 206927 s002,

Kyprolis (NDA 202714 s010), Ninlaro (NDA 208462), Rapamune, Afinitor, Mekinist.

- **Variant and gene resources:** ClinVar, gnomAD v4, Ensembl REST and VEP, UniProt (O60566), OMIM and Orphanet identifiers, Human Phenotype Ontology.
- **Drug/target interfaces queried (all negative for MVA1):** ClinicalTrials.gov, Open Targets, ChEMBL, CMap/LINCS.
- **Pathway and model resources:** Reactome (R-HSA-69618), KEGG (hsa04110), MGI (MGI:1333889).
- **Assay standards:** OECD Test Guideline 487 (micronucleus assay).

Q8 — Proprietary data

None.

Q9 — How did you characterize the variant’s mechanism?

Loss of function, with the druggable target displaced downstream of the lesion.

The nonsense allele removes the entire C-terminal kinase domain; the missense allele sits inside it. Cell lines derived from MVA patients with biallelic mutations show “an impaired mitotic checkpoint, chromosome alignment defects, and low overall BUBR1 abundance” (PMID 20516114), and allele-specific effects — not merely total BubR1 quantity — drive phenotypic heterogeneity (Sieben 2020, PMID 31738183).

We were explicit about what we could not determine. Phase is not established, and we measured that it is not establishable from these data: there are no parental samples; the challenge distributes a VCF and raw reads but no alignments, we did not align them, and short reads could not phase two variants whose only intervening heterozygous site lies 6.8 and 4.1 kb away; and the two variants lie 10,911 bp apart with no GATK phase tag on either — a limit visible in the data itself, since GATK did phase three variants 5 bp apart elsewhere in the same VCF. Reference-panel phasing fails separately: all three heterozygous sites inside the interval are absent from the 1000 Genomes 3,202-sample phased panel, and Beagle 5.5 drops all three. Our retrieval pipeline is insensitive to phase in any case: two

pathogenic nonsense alleles declared in cis are still called compound heterozygous, still PATHOGENIC and still ranked first, at an absolute score difference of 0.0034. Nonsense-mediated decay of the truncated transcript is unmeasured and we do not assert it either way without patient RNA. p.Asn1002Lys has no published functional assay, and its in-silico predictors disagree (REVEL 0.472, AlphaMissense 0.923), which is why we let it stand as a VUS.

We also asked whether aneuploidy burden could serve as a trial endpoint in these data, and computed that it cannot. A first-order calculation — mean $|\text{BAF} - 0.5|$ at DP 44 against a measured between-chromosome SD of 0.0035 — puts the detection limit at ~14% of cells for a clonal trisomy (16% under a stricter 2-SD criterion). A later calibrated simulation (Ψ over 125-site windows, Monte Carlo null) improves that to **2.07% at an overdispersion $\kappa = 2$** ; that is a *simulated* limit conditional on κ , not a measured sensitivity, and it places variegation's 1–2% per chromosome at the edge of what is achievable rather than far below it. The argument that does not depend on any limit is **non-identifiability**: for a perfectly balanced variegated mixture, mean copy number is exactly 2 and mean allele fraction exactly 0.5 by symmetry, and simulated detection stays at the 0.1–0.2% false-positive rate even with 40% of cells aneuploid. That proof covers the balanced mixture we modelled, not every possible MVA genome, and it is enough to disqualify the endpoint.

We ran the calibrated statistic on the proband and **withdrew the result**: the pre-registered threshold fires on all 22 autosomes, which indicates a null model that does not fit these data, and the pre-registered external control came back **not conclusive** because the control genomes were called differently from the proband. We report no clinical negative for this child.

Q10 — Time and effort

Approximately 30 hours of analyst time over two days (28–29 August 2026), of which the large majority went to literature verification and adversarial review rather than to computation. Compute was negligible: the genome-wide prioritisation ran in 53 seconds on a commodity workstation with no GPU, and the robustness controls in about ten minutes. Later rounds added the pre-registered DepMap,

benchmark, power and mosaicism-control analyses, all on public data. The phase pilot is **not** one of the pre-registered five — it ran before that pre-registration and we do not count it as one (`VERIFY.md` lists the five pairs). All of this is CPU-only except the mosaicism null calibration, which used one consumer GPU.

Q11 — Method abstract (≤ 500 words)

No approved drug restores BubR1 function, so we did not look for one. We followed the lesion downstream — weakened spindle assembly checkpoint, mis-segregation, constitutional aneuploidy, stoichiometric protein imbalance, proteotoxic stress — and asked which approved agent acts on the last link.

We applied five filters instead of the usual two. Beyond “is it approved” and “is the mechanism plausible”, we added a pharmacokinetic filter (does the achievable C_{max} reach the concentration effective in vitro?), a patient-specific safety filter, and a direction-of-effect filter (does the drug increase mis-segregation among survivors?).

Metformin, the candidate the literature points to, dies on filter three: therapeutic plasma levels are micromolar while complex I inhibition requires millimolar (PMID 37343530). AICAR’s selectivity does reproduce in human cells (PMID 22890317), but that paper does not test metformin, and AICAR is not approved.

Bortezomib survives filters one to three. Aneuploid cells depend on increased protein degradation, and aneuploidy level was associated with multiple myeloma patients’ response to proteasome inhibitors (Ippolito 2024, PMID 39247952), on cohorts small enough (8 complete responders vs 50 progressive) that we treat it as motivation, not prediction. An EC₅₀ below 40 nM in highly aneuploid lines against a label C_{max} of 89–120 ng/mL (231–312 nM) at 1.3 mg/m² IV gives a margin of 5.8–7.8× on **total plasma** drug at a peak — not free or intratumoral drug, and 312 nM is the top of a label interval, not a validated clinical threshold.

It then fails filter four for this child: motor neuropathy in 8% of paediatric patients, on top of existing skeletal muscle atrophy. But filter four is patient-specific. For an MVA patient with an active malignancy — where the comparator is cytotoxic chemotherapy — the balance

inverts, and biallelic BUB1B carries a high risk of embryonal tumours. Our deliverable is that answer prepared in advance, with its margin and its stopping rule.

Strengths. Every candidate is killed or kept by a stated criterion, including our own; this is not an exhaustive screen, and §8 of the report lists what we excluded. We report an antagonism a combination-minded team could walk into: reducing translation protects CIN cells from proteasome inhibition (PMID 31530568), so mTOR inhibitors would be predicted to antagonise, not synergise. The framework transfers to any rare disease; applied to MVA2 and MVA3 it returns different verdicts.

Limitations. The proteasome dependency is established in cancer aneuploidy, not constitutional mosaicism; whether it transfers is what our proposed experiment tests. Our own pre-registered DepMap analysis found the association attenuated in solid lineages, the lineage of this child's tumour. The 0.897 power figure belongs to an idealised simulation with an *assumed* fourfold selectivity, not a validated estimate. The EC50 is a conservative bound read from a figure panel. Phase is unproven and unobservable in these data; parental genotyping is the cheapest informative experiment. No disease-modifying drug can be recommended from these data today — what most changes his prognosis now is surveillance, not a molecule.

Acknowledgement

This work was made possible through the Hackathon, organized by Sage Bionetworks in partnership with the MVA Society, Hugging Face, and BEACON (The Benchmarking, Evaluation, and Assessment Consortium for Science), with prize sponsorship from AWS and Anthropic. We are deeply grateful to the child and their family who generously contributed their data and their story to advance research into this rare disease. We acknowledge their trust in making this Hackathon possible.

Data availability

The individual-level data analysed here is the Hackathon dataset (SageBio/mva-hackathon-2026-data), accessed under the challenge data transfer agreement accepted by Hugging Face account `cpu-16` on 2026-08-28, and cited as directed on the Hackathon Synapse page. **No patient sequence data or genotype-scale derivative is included in this document or in our public repository**, and every copy we hold will be deleted by 2026-11-23 with written confirmation to the organisers. The two diagnostic variants are named, as findings, which the challenge rules permit. All other inputs are public: Exomiser 15.1.0 with data release 2602, ClinVar, gnomAD v4, the Human Phenotype Ontology, Ensembl and VEP, the 1000 Genomes and GIAB genomes, DepMap 24Q4, AlphaFold and the literature cited by PMID. Analysis code, pre-registrations and results: <https://github.com/cpu-16/mva-hackathon-2026>