

Track 2 — Drug Repurposing for MVA1 (BUB1B, presumed compound heterozygosity; phase not established)

Team: ciberpty · Rare Disease, Real Kid: MVA Hackathon 2026
Repository: <https://github.com/cpu-16/mva-hackathon-2026> · **Methods description:** Appendix B of this PDF

*This document has two parts. **Part I** (this part, with its reference list) is the decision document: every conclusion, every number that carries it, and every limitation, at the length a judge can read in one sitting. **Part II** is the full evidence: the same sections at full length, with derivations, sources quoted verbatim, the correction history and every pre-registered analysis. Where Part I compresses, Part II is authoritative for detail; every number in Part I is taken from Part II or from the results files both cite.*

Read this first — for the family and the treating team

Written in plain language for the child's parents and the clinician sitting with them. Every claim is sourced in the sections that follow.

What we looked for. An existing, approved medicine that could help a child with mosaic variegated aneuploidy whose presumed cause is two changes in the *BUB1B* gene (one of them still formally uncertain, and whether they sit on opposite copies of the gene is not yet shown; §1, §10). We fixed in advance the tests a candidate had to pass and put the main candidates through them. It is not an exhaustive screen; §8 lists what we ruled out and why.

What we are not recommending. No drug should be started for this today. The honest answer is more useful than an encouraging one.

- **Metformin**, the medicine most often suggested for this kind of problem, should not be started on this reasoning: no one has measured what dose would be needed to act on *this* problem, and

the best-studied way it might act needs concentrations about a thousand times higher than the blood levels the usual doses produce, and whether it could act another way at those levels has never been tested (§3.1).

- **Bortezomib** is an approved cancer medicine that acts on a weakness that cancer cells with extra chromosomes have been shown to have. **Nobody has shown that this child's cells have it.** Our own analysis of public cancer data found the genetic link weaker in solid tumours as a group, which is the group his tumour belongs to, and no significant drug signal there; his tumour type is almost absent from that data. In one paediatric trial of bortezomib given with intensive chemotherapy, nerve injury was recorded in 25 of 140 patients; that trial did not test whether his cells have the vulnerability, and his own added risk is unknown. He already has muscle wasting (skeletal muscle atrophy). **Outside an active cancer, that trade is not worth making (§5).**

What can be done this month.

	Action	Why
1	Confirm the child's current age with the care team	Every schedule below depends on it, and we do not have it
2	Renal ultrasound every 3 months, birth to age 7	Published consensus for all forms of MVA (PMID 39264246: the SIOP-Europe host-genome group's schedule, endorsed by the AACR workshop). Not our idea; how it applies to this child is the treating team's decision
3	Regular full clinical examination , including the orbit, skin and soft tissues	Consensus already asks for regular clinical examination, with no end date; rhabdomyosarcoma is not a kidney tumour. The emphasis on the orbit is our suggestion after a single 2026 case of the related MVA3 with a second tumour behind the eye at age 12 (PMID 42595739) — not a guideline interval
4	Avoid unnecessary radiation; keep HPV vaccination up to date	Consensus guidance (PMID 39264246)
5	Ask about testing and counselling for both parents — a ready-to-adapt laboratory request is included (<code>clinico/PARENTAL_SEGREGATION_ORDER.md</code>)	Whether the two changes sit on opposite copies of the gene is unknown, so test both parents before anyone calls this compound heterozygous. The clinical notes also mention repeated miscarriages; a 2026 study of two families is a hypothesis about miscarriages in <i>BUB1B</i> carriers (PMID 42434306), not an explanation of this family's losses (§7)

What would change the answer. If a new tumour appears, the question changes from “should he take a medicine” to “which medicines belong in the treatment plan”. Standard treatment still comes first. These pages do not establish that adding bortezomib would improve the balance of benefit and harm: as a single agent in

the paediatric solid-tumour trials we found it produced no objective responses, the nerve-injury figure comes from use with intensive chemotherapy, and we have no combination response rate that would justify adding it (§5). What we have prepared is the question a tumour board could be asked — the exposure comparison, the toxicity numbers, the experiment on his own cells (§6, which measures constitutional cells, not tumour selectivity) and the reasons a board could decline — so nobody has to assemble it in a hurry.

What we are not certain about. The strongest evidence comes from cancer patients, not from children with MVA. Whether the same weakness exists in this child's cells is genuinely unknown; §6 describes the experiment that could show us we are wrong.

Executive summary

We organised the assessment around five explicit questions (§2). **Metformin**, the candidate the literature points to, is not evaluable for aneuploidy-selective activity — no effective concentration has ever been measured — and the complex I mechanism proposed for it needs millimolar concentrations against micromolar plasma levels (§3.1). **Bortezomib** is not excluded: aneuploid cancer cells depend on the proteasome (Ippolito 2024, PMID 39247952), its total-plasma peak at the approved IV dose is 5.8–7.8 times a 40 nM reference upper bound on the culture EC50 of highly aneuploid lines, and aneuploidy level tracked response in myeloma patients (§3.2). That comparison is an exposure plausibility screen on total drug against an upper bound on the EC50, not a therapeutic threshold: on free drug (83% protein-bound) the peak is $\approx$ 39–53 nM, about 1.0–1.3 times the 40 nM reference value ($\approx$ 99 nM, or 2.5 times, on the label's 223 ng/mL cohort) — a comparison to a reference value, not a measured free-drug margin, and not the six- to eightfold the total-drug ratio suggests.

The controlling conclusion, used throughout this document. No disease-modifying drug is supported for use in this child today. Bortezomib is a conditional research candidate: the stated IV total-plasma peak comparison does not exclude it, but comparable free exposure, duration at the target and tumour selectivity in MVA remain unestablished. PRISM did not detect a significant proteasome-inhibitor

association in the tested solid lines; CRISPR showed a weaker aneuploidy–PSMB5 association in solid than in haematological lines. An active malignancy would change the clinical comparator; it would not establish a favourable benefit–risk balance.

We then tested the part of that argument that public data can test, pre-registered, and the predicted drug association failed its criterion (§4): across 424–444 evaluable solid-tumour lines (depending on the drug) no proteasome inhibitor reached significance, and the genetic dependency on bortezomib’s target is significantly weaker in solid tumours than in haematological ones (interaction $\beta = -0.017$, 95% CI -0.031 to -0.002 , $p = 0.026$; removing any single solid lineage never moves that interval across zero). This child’s tumour was solid, and his tumour type is essentially absent from the drug screen.

Bortezomib then **fails the patient-specific safety filter**: peripheral neuropathy in 18% (25/140) of patients aged 1–27 years in a trial of bortezomib with intensive combination chemotherapy — his incremental risk is unknown — on pre-existing muscle atrophy. **Not for this child, not now** (§5). What remains is a prepared, evidence-bounded question for the setting of an active tumour, an experiment that could refute the premise in this child’s own cells (§6), and the actions justified today without any drug — surveillance and a parental phase test (§7).

What we contribute is eight pre-registrations that constrain the claim — DepMap and its follow-up, MVA-Replay, mosaicism, power and its four-parameter follow-up, dose–response, and the TRIP13 transfer; PRISM did not establish the predicted drug association in solid lines and CRISPR supported a weaker target association in solid than in haematological lines, the mosaicism control came back not conclusive, the power analyses showed the assay as first sized ($n = 3$) to be underpowered even under its own assumptions — about 12 cultures per arm are needed — and the dose–response analysis left the assay’s assumed 4× within-cell selectivity (separate from the exposure ratio) unsupported — a reusable five-filter worksheet, and a retrieval-benchmark tool transferred to a second MVA gene from a fresh clone of our repository (§9).

1. The variants

The proband carries two heterozygous **BUB1B** variants (NM_001211.6, MANE), read here as compound heterozygous **on presumption: phase is not established** (§10).


Variant	Consequence	Exomiser ACMG	ClinVar	gnomAD, Exomiser snapshot (exomes NFE)	gnomAD via VEP
c.2210T>G p. (Leu737Ter)	stop_gained	PATHOGENIC	533901, Pathogenic/Likely pathogenic	9.98×10 ⁻⁵	3×10 ⁻⁵
c.3006T>G p. (Asn1002Lys)	missense	UNCERTAIN_SIGNIFICANCE	not present (ClinVar holds c.3006T>A, a different nucleotide, VUS — does not transfer)	9.0×10 ⁻⁷	no record

p.Leu737Ter removes the C-terminal kinase fold entirely (UniProt annotates it at 766–1050), 16 residues from BUBR1^X753, a characterised truncating MVA allele (PMID 31738183). **p.Asn1002Lys** has no published functional assay. On the AlphaFold model (pLDDT 91.1 at the residue) Asn1002 is buried (relative solvent accessibility 0.162) in a hydrophobic pocket, and the substitution inserts a longer, charged side chain — evidence about the substitution, not a classification: no $\Delta\Delta G$ was computed, and the retrieved predictors disagree (REVEL 0.472, AlphaMissense 0.923). No calibration set exists to score it against: in ClinVar (2026-08-28 release) *BUB1B* carries 1,497 missense records of which one is Likely pathogenic — one star, single submitter — against 82 of 94 truncating records P/LP — a property of the database, not a prior about this allele. The mouse model pairing the nearby missense L1012P with a truncation died prematurely (PMID 31738183); our proband is alive; the mouse series motivates testing whether p.Asn1002Lys retains residual function, but neither his

survival nor the allele-fate assays of §6 establish its magnitude relative to L1012P — that would need a comparative checkpoint assay. Full detail: Part II §1.

2. Five filters

#	Filter	Question	What it eliminates here
1	Regulatory	Approved for marketing?	AICAR, 17-AAG, reversine, apcin, proTAME, DCZ0415
2	Mechanistic	Does it address the lesion or a direct consequence of it?	Agents on uninvolved pathways
3	Pharmacokinetic	Is the achievable exposure excluded by the effective in-vitro concentration? (“not excluded” is the best outcome this screen can return)	Metformin
4	Patient-specific safety	Compatible with <i>this</i> child’s comorbidities?	Bortezomib, at present
5	Direction of effect	Can an increase in mis-segregation among survivors beyond a prespecified margin be excluded?	Untested for the candidates that reach it; §6 is where it would be measured, against a provisional 50% margin



The five-filter funnel

Filter 3 is a possibility check on total plasma drug at a peak, not a demonstration of efficacy: it ignores protein binding, exposure duration and tissue penetration. Filter 5 asks the question repurposing exercises forget: among the cells that survive the drug, is the defect made worse?

3. Applying the filters

3.1 Metformin: not evaluable on the mechanism proposed for it

Tang et al. found aneuploidy-selective sensitivity to **AICAR** (PMID 21315436); AICAR is not approved, and metformin is the marketed stand-in that activates AMPK through complex I inhibition. Therapeutic metformin plasma concentrations are **micromolar**, while complex I inhibition requires **millimolar** (PMID 37343530) — a gap of roughly a thousandfold in concentration ranges, not a dosing calculation. Therapeutic metformin does activate AMPK in patients; what the gap rules out is the complex I route at plasma concentrations. An AMPK-mediated aneuploidy-selective effect at those concentrations is not excluded — it is **untested**: no study reports a metformin aneuploidy-selective effective concentration, the one paper reproducing AICAR's selectivity in human cells attributes it to EGFR

degradation and does not test metformin (PMID 22890317), and Tang's instability models were never given metformin. In children, dose-normalised exposure varies 3.3-fold and clearance is renal (PMID 42399684); this child has nephrocalcinosis. **Verdict: rejected at filter 3 as unevaluable, not as disproven.** Part II §3.1.

3.2 Bortezomib: passes filters 1–2; not excluded by filter 3

Aneuploid cells carry a stoichiometric protein imbalance and compensate by increasing degradation, which makes them dependent on the proteasome; Ippolito et al. showed this across diploid-versus-aneuploid models and hundreds of cancer lines, and found aneuploidy level associated with myeloma patients' response to proteasome inhibitors (PMID 39247952). That human signal rests on small cohorts and one-tailed tests (single agent: 8 complete responders vs 50 progressors, $p = 0.014$; with chemotherapy and dexamethasone: 13 complete vs 14 minimal responders per the figure legend, $p = 0.038$) — enough to justify an experiment, not to predict a patient. A 2026 paired-CRISPR preprint independently nominates proteasome subunits and UBE2H as aneuploid-specific dependencies (PMID 42094535; not peer-reviewed, weighted as corroboration).

	Value	Source
EC50, highly aneuploid lines (72 h)	< 40 nM , conservative bound read from Fig. 6o	PMID 39247952
Cmax, 1.3 mg/m ² IV, repeated dosing	89–120 ng/mL = 231–312 nM (the twice-weekly range the label reports; the same label gives 223 ng/mL in its IV-vs-SC cohort)	FDA label NDA 021602
Ratio, IV	5.8–7.8 , on total drug	—
Subcutaneous	53.1 nM; ratio 1.33 on total drug at the peak, AUC equivalent to IV — exposure unresolved , neither a pass nor a fail	same

Read the ratio as an order of magnitude on *total* plasma drug at a peak. **The exposure evidence, laid out rather than reduced to a pass:**

Quantity	Value	What it does and does not support
Effective concentration	EC50 < 40 nM at 72 h in 5 highly aneuploid cancer lines, read from Fig. 6o (PMID 39247952); no tabulated value; medium and serum binding unstated	A bound on total drug in culture, not a threshold for MVA cells
Plasma binding	83% bound to human plasma proteins over 100–1,000 ng/mL (label §12.3)	Free fraction $\approx$ 17%, applied below 100 ng/mL as an approximation
Cmax, IV 1.3 mg/m ²	112 ng/mL first dose; 89–120 ng/mL twice weekly; 223 ng/mL in the IV-vs-SC comparison cohort (label §12.3)	Total drug at a peak; free peak $\approx$ 39–53 nM (or $\approx$ 99 nM from 223 ng/mL)
Cmax, SC 1.3 mg/m ²	20.4 ng/mL, with AUC equivalent to IV (label §12.3)	Free peak $\approx$ 9 nM
Pharmacodynamics	Maximal 20S inhibition 73–83% in whole blood at 5 min (label §12.2); elimination half-life 76–108 h on multiple dosing, with plasma falling steeply after infusion	Duration at the target is not established by any number here

On total drug the IV ratio is 5.8–7.8 (231/40 and 312/40). **On free drug the peak is $\approx$ 39–53 nM, that is 0.98–1.33 times the 40 nM reference value, or $\approx$ 99 nM (2.5 times) on the 223 ng/mL cohort** — these are ratios to a reference denominator, not measured free-plasma-to-free-medium margins, and since 40 nM is an upper bound on the EC50 the true free-drug margin is unknown; the culture value is itself partly free drug in serum-containing medium, so the two sides are not on the same footing. What filter 3 supports is that the IV route is **not excluded**; it does not support comparable free exposure or duration at the target, and we do not correct one side for binding while leaving the other unexplained. The culture-side EC50 carries its own unstated medium and binding conditions. Myeloma is a plasma-cell neoplasm with its own proteotoxic load and clonal, not variegated, aneuploidy; MVA1 is a different setting in which the same dependency is plausible and **untested**. Part II §3.2.

3.3 Everything else: not evaluable

Drug	Verdict	Why
Carfilzomib, ixazomib	Not evaluable	No IC50 in an aneuploidy-stratified model; Cmax known, no valid denominator
Hydroxychloroquine	Not evaluable, mechanistically weakened	Tang's hit was chloroquine, which did <i>not</i> differentially inhibit the human CIN lines
Everolimus / sirolimus	Not evaluable, possibly counterproductive	§3.4
Trametinib	Not evaluable — the gap most worth closing	Aneuploid RPE1 clones are MEK-inhibitor-sensitive (PMID 39251587) but only relative IC50s were recoverable; Cmax 36.1 nM at 2 mg/day; approved from age 1 with dabrafenib for BRAF V600E low-grade glioma; filter 2 passes on plausibility, filter 5 untested; null in our DepMap test

We leave cells empty rather than manufacture ratios from unstratified models.

3.4 The antagonism nobody should walk into

Aneuploid cells are vulnerable to proteasome inhibition *because* of translational burden; ovarian carcinoma spheroids acquired resistance to proteasome inhibitors by suppressing mTORC1, and PTEN loss resensitised them (PMID 31530568). That paper did not test rapalogues; the antagonism is our inference from its mechanism. **Prediction, falsifiable:** everolimus or sirolimus would antagonise bortezomib in aneuploid cells; §6 tests the combination. The same paper notes bortezomib's poor clinical response in advanced ovarian cancer — a solid-tumour failure any honest reading carries. A tension we flag: mTORC1 hyperactivity correlates with sarcopenia in *monoallelic* BubR1 mice (PMID 31738183), a muscle-directed rationale that points the other way and must not be read as support for the combination.

4. We tested our own candidate on public data, and pre-registered it first

The objection to §3.2 is that the human signal is myeloma and this child's tumour was an embryonal rhabdomyosarcoma. We had no number for solid tumours, so we produced one — drug list, directions, statistic, correction and stopping rules committed to `depmap/PREREGISTRO.md` before the analysis ran (the commit history documents the record, not our conduct; deviations are listed in `VERIFY.md`, and one untestable hypothesis was recorded as a dated amendment).


Method. An arm-level aneuploidy score for 2,420 DepMap 24Q4 lines (validated: stable MSI lines average 5.0 altered arms, CIN lines 16.5), correlated with PRISM Repurposing drug sensitivity and DepMap CRISPR gene effect. The pipeline is not blind: across 6,790 compounds 1.1% reach $p < 0.001$ against 0.1% expected — an excess of small p-values that shows the screen detects associations in general, not that it is calibrated or sensitive for the proteasome-inhibitor tests in particular.

Result 1 — drugs (solid lines; $n = 444$ for bortezomib, 424 carfilzomib, 438 ixazomib, 438 metformin). No proteasome inhibitor reached $FDR < 0.05$. Bortezomib $\rho = -0.075$ (95% bootstrap CI -0.168 to $+0.017$), carfilzomib -0.060 , ixazomib $+0.028$; metformin null as predicted; paclitaxel, the cytotoxic control, ran the other way ($+0.096$), which argues against — without excluding — a “sick cells die more” confound. Haematological lines have no PRISM value for these drugs (0 of the 98 scored haematological models with a PRISM entry; the CRISPR haematological set is 113).

Result 2 — the target (946 lines). PSMB5 dependency tracks aneuploidy overall ($\rho = -0.198$, $q = 5 \times 10^{-8}$) but splits by lineage: **haematological $\rho = -0.261$ (-0.422 to -0.084); solid $\rho = -0.077$ (-0.149 to -0.009), $q = 0.12$** — the interval excludes zero narrowly while the pre-registered FDR test did not pass, and the pre-registered verdict stands. The **aneuploidy $\times$ lineage interaction is $\beta = -0.017$ (95% CI -0.031 to -0.002 , $p = 0.026$);** per ten additional altered arms PSMB5 dependency moves -0.06 in solid lines and -0.23 in haematological lines; ploidy was not significant in either model ($p = 0.82$ in the interaction model, $p = 0.875$ in the main-effects model), which does not

exclude residual confounding or whole-genome-doubling history. These model statistics were first computed ad hoc and are now reproduced by the follow-up script.

DepMap follow-up (pre-registered): the aneuploidy-proteasome association, with its uncertainty



Follow-up: effect sizes with 95% bootstrap intervals; the solid stratum by lineage

Follow-up, pre-registered after the result (depmap/PREREGISTRO_SEGUIMIENTO.md), all five predictions (P-S1 to P-S5) held. Dropping each of 26 solid lineages never flips a sign or moves the interaction interval across zero (p 0.013–0.037); all 27 within-lineage intervals overlap zero; a fraction-of-arms score gives similar estimates with unchanged signs (–0.077, –0.054, +0.032), a ploidy-residualised score shrinks the small drug correlations by more than half (–0.075 → –0.028) without changing signs. The two assays share 365 solid models and agree with each other only at $\rho = -0.070$ (–0.176 to +0.033): they are different modalities on overlapping models, not independent replications. **Rhabdomyosarcoma:** 17 scored lines, 3 with a drug value — less killed at 2.5 μM than 98%, 91% and 89% of solid lines — and 10 with a CRISPR value, unremarkable. No within-RMS association test was performed; it is a coverage limitation for exactly the tumour type this child had.

What it costs us. The link our bortezomib argument rests on is detectable where myeloma sits and attenuated in solid tumours: *the attenuation is not an artefact of any single lineage and its interval excludes zero; the solid-stratum association itself remains compatible with zero* — three results that coexist: the marginal Spearman interval

excludes zero narrowly, the pre-registered FDR test did not pass, and the adjusted regression slope's interval includes zero. Ippolito's own bortezomib screen (Fig. 6p, reversine-induced aneuploidy, $p < 0.0001$) is positive; ours carries the same sign and fails significance. The two designs estimate different quantities — a within-line perturbation versus a cross-line correlation — so their p-values do not rank them, and our null is not evidence against their result. Ippolito also report associations in pancreatic and paediatric PDX models (their Suppl. Fig. 8o–r), which say nothing about MVA. What we add is the lineage split. PRISM is a single-dose viability screen, not an EC50, and cancer lines are not constitutional MVA. Part II §3.5 and `depmap/RESULTADOS*.md`.

5. Filter 4 — safety in *this* child, where the candidate stops

Bortezomib is not excluded by the total-plasma comparison. It does not clear this patient today.

Concern	Data	Relevance here
Peripheral neuropathy	25/140 (18%) of patients in AALL07P1 — relapsed lymphoid malignancy, bortezomib with multi-agent reinduction chemotherapy including vincristine; safety population aged 1.0–26.8 years — motor 11/140 (8%) (FDA BPCA review, 2015): observed regimen-associated events, not this child's incremental risk from bortezomib alone; adults grade ≥ 2 in 39% by the IV route (FDA NDA 206927 s002, 2022)	Skeletal muscle atrophy (HP:0003202): the decisive objection
Cytopenias, infection	FDA label	Prior chemotherapy; marrow reserve unknown
GI toxicity, hypotension	FDA label	Failure to thrive (HP:0001508)
Renal	PK not clinically altered by renal impairment, per label	Nephrocalcinosis (HP:0000121): function, hydration and electrolytes first

Bortezomib is not proposed as prophylaxis or chronic therapy for this child. Filter 4 is patient-specific: for an MVA patient with an active malignancy, the comparator becomes chemotherapy rather than nothing. That changes the comparator; it does not establish that adding bortezomib improves the benefit–risk balance, which remains a judgement for the treating team and a tumour board, off label or under an n-of-1 protocol with its own ethics approval and stopping rules, informed by the constitutional-cell experiment of §6 and a separate tumour-material comparison, not a result of this report.

What is already known about bortezomib in the tumour he had, and in children — favourable preclinical observations, limited in-vivo solid-tumour activity, no demonstrated clinical benefit. RMS cell lines are killed at 13–26 nM, sparing primary myoblasts, and bortezomib reduced the growth of an RMS xenograft (PMID 18342500). But across the Pediatric Preclinical Testing Program, in-vivo activity against solid-tumour xenografts was limited while leukaemia xenografts responded (PMID 17420992). The COG phase I ADVL0015 set 1.2 mg/m² twice weekly for two of every three weeks as the paediatric dose, with dose-limiting thrombocytopenia at 1.6 mg/m² and **no objective responses** in 15 children with refractory solid tumours (PMID 15570082); ADVL0916 with vorinostat likewise saw none, with a sensory neuropathy progressing to grade 4 as a DLT (PMID 22887890); an adult sarcoma phase II found minimal single-agent activity in soft-tissue sarcoma and closed its osteosarcoma/Ewing/rhabdomyosarcoma arm for low accrual (PMID 15739208). **In the paediatric trials we found, single-agent bortezomib produced no objective responses in solid tumours** (ADVL0015: 15 enrolled, 11 assessable, no objective responses). If it is ever reached for, it would be as part of a combination, and the aneuploidy argument is a reason to ask whether MVA-driven tumours are the exception, not evidence that they are.

Spindle poisons depend on a competent checkpoint; a reduced-intensity regimen for an RMS in PCS/MVA is on record as a single case (PMID 31184400). That is a mechanistic caution to weigh against probability of cure, not a contraindication, and never a reason to withhold curative therapy.

The tumour-stage question, specified. *Research setting:* an active MVA-associated tumour. *Comparison:* tumour-derived material against constitutional normal-cell comparators under matched exposure conditions, with the tumour's standard regimen as the treatment comparator. *Question:* does adding bortezomib produce an incremental effect on the tumour material with a reproducible tumour/normal separation? *Exposure condition:* the tested schedule must be justified against clinical exposure and pharmacodynamic evidence (duration of proteasome inhibition), not against the 312 nM peak alone. *Decision (provisional planning thresholds — design choices, not findings — to be frozen before any material is obtained):* compare vehicle, bortezomib, the standard regimen, and standard regimen plus bortezomib, with viable-cell recovery at a fixed post-washout time as the primary endpoint; require at least 20 percentage points of incremental tumour killing from adding bortezomib, at least 10 points of tumour-normal separation, and no more than 10 points of incremental killing in either normal comparator, with Δ_T the incremental tumour killing (standard regimen plus bortezomib minus standard regimen, on vehicle-normalised viable-cell recovery), Δ_N the same quantity in each normal comparator and $D = \Delta_T - \Delta_N$: advance only if the simultaneous lower confidence bound of Δ_T exceeds 20 points, both lower bounds of D exceed 10 points and both upper bounds of Δ_N stay below 10 points, with culture-level resampling and replication in two independently initiated experiments; a normal-comparator lower bound above 10 points is *harm*, a tumour upper bound below 20 points is *insufficient benefit*, anything else is *inconclusive*. The operating characteristics of this rule have not been simulated; failure of the constitutional-fibroblast hypothesis (§6) does not decide this question. *Prioritisation:* there is no comparative basis today for ranking this work above trametinib or another unresolved candidate, and we say so rather than assert one.

6. Endpoints, and the experiment that could show us we are wrong

Bulk sequencing cannot serve as the aneuploidy-burden endpoint (Part II §5). From the proband's VCF we computed a first-order detection limit of ~14% of cells for a clonal trisomy, then a calibrated simulated limit of **2.07% at overdispersion $\kappa = 2$** (1.46% at $\kappa = 1$, 2.95%

at $\kappa = 4$) — a simulated limit under a fitted null, not a measured sensitivity. If ~30% of cells were aneuploid spread across 22 autosomes — an illustrative bound, not a patient measurement — each chromosome would be affected in 1–2% of cells, which sits at that edge, and the argument that needs no limit is this: **a perfectly balanced variegated mixture is not identifiable from bulk allele fractions** — detection stays at the 0.1–0.2% false-positive rate from 2% to 40% of cells. Real variegation need not be balanced, so this disqualifies the endpoint without saying bulk WGS can never see MVA. **We ran the calibrated statistic on the proband and withdrew the result:** the null did not describe the data (22 of 22 autosomes exceeded the threshold), the z-score we had reported was chosen after seeing the data and never calibrated, the residual overdispersion ($\approx 2.3\times$) is uncontrolled, and the pre-registered external control of ten 1000 Genomes individuals came back *not conclusive*: the proband fell below the control envelope, an outcome the amendment had not listed, with the controls' different calling explaining their inflated dispersion. **No clinical negative is reported.** Any screen must measure mis-segregation per cell: metaphase counts, micronuclei (OECD TG 487), or single-cell DNA.

The experiment (Part II §6). Patient dermal fibroblasts against two matched control donors (the patient-versus-control effect estimated separately for each donor, with aggregation and multiplicity fixed before sampling) and a reversine-treated RPE1 control of assay responsiveness, not a required response ranking; allele-fate assays; a bortezomib dose–response 0–1,000 nM at 72 h (spanning the 312 nM Cmax); a bortezomib + everolimus arm for the predicted antagonism; a micronucleus assay with FISH for 3–5 chromosomes on survivors of every condition, with proliferation recorded because a change in division rate alone moves the count. Eight weeks of assay after 4–8 weeks to establish the line (longer if twelve independent cultures per arm must be expanded first): a planning figure of **12–16 weeks** from biopsy (MVA1 fibroblasts may grow slower), after consent, assent and ethics approval whose timing is not ours to control.

What it decides, and what it cannot. Fibroblasts are constitutional MVA cells, not tumour cells: a sensitivity difference measures differential response in the sampled constitutional lines — it does not by itself identify a proteotoxic dependency — and is, on its own, as

compatible with systemic toxicity as with promise. A negative fails to support transfer without closing the tumour-board question, which needs aneuploidy-high tumour material this design does not sample.

Reading	Consequence
Aneuploid-cell fraction f measured first, on the culture and passage to be treated (cells with ≥ 1 numerical abnormality, plus abnormal chromosomes per abnormal cell; PCS scored separately)	Gate: the bulk EC50 ratio to expect at an assumed 4 $\times$ selectivity is 1.14 $\times$ at $f = 0.10$, 1.30 $\times$ at 0.20, 1.50 $\times$ at 0.30, 2.00 $\times$ at 0.50 — a pre-specified reading rule, not a validated threshold, fixed for this experiment; any recalibration is exploratory and requires prospective evaluation. Under the declared model (four-parameter fit, biological variability CV 0.20) the design needs about 12 independent cultures per arm at $f = 0.30$ and 23 at $f = 0.20$; below $f \approx 0.20$ use a per-cell endpoint instead of bulk viability
EC50 shift at the row for the measured f , 95% CI excluding 1, at ≤ 312 nM	Response difference in sampled constitutional cells, mechanism unresolved: motivates, but does not decide, the tumour-material stage. A shift more than twice the row's prediction is not a better result — suspect the mixture model or systemic toxicity
Patient cells markedly more sensitive than both control donors	Stop: this raises concern about constitutional-cell toxicity, flagged independently of RPE1 (which establishes assay responsiveness only); the design cannot determine the mechanism, and it is the outcome most easily mistaken for success
Micronucleus frequency among survivors (per 1,000 binucleated cells, treated vs vehicle): with Δ the treated-minus-vehicle difference and M a provisional harm margin of 50% of the vehicle rate (a design threshold set here), the upper 95% confidence bound of $\Delta - M$ must lie below zero — evidence below the margin, not proof of no increase; proliferation recorded, and centromere/kinetochore labelling plus FISH for 3–5 chromosomes used to attribute any increase to mis-segregation	An interval for $\Delta - M$ crossing zero is inconclusive (not shown to pass filter 5); an interval for $\Delta - M$ wholly above zero is harm exceeding the chosen margin in these constitutional cells, which stops the constitutional-cell proposal regardless of killing
No shift	Fails to support transfer in these cells; stops the constitutional-cell hypothesis, not tumour-directed investigation, which would need its own decision on tumour material

Power, stated under the model we pre-registered rather than the one we first ran (potencia/RESULTADOS_4P.md, pre-registered follow-up, 8 Sep): with four free Hill parameters, a biological multiplier on each replicate's EC50 (CV 0.20), a plate offset and per-well noise, **the design at $n = 3$ per arm has power 0.17** at the central scenario ($f = 0.30$, $4\times$ selectivity), not the 0.897 of the original two-parameter simulation without biological variability, which we keep only as the optimistic reference. Fitting never failed; most of the reduction is associated with changes beyond the curve fitter — biological and plate variability, the grid and the test — whose individual contributions were not isolated (a failed fit — non-convergence, EC50 at a bound, or a top below the bottom + 0.2 — counts as a non-rejection). Under that scenario the simulation first crosses 0.80 at **12 independent cultures per arm** at $f = 0.30$, 23 at $f = 0.20$ and more than 24 at $f = 0.10$ — a conditional planning figure, not a validated requirement — so the $n = 3$ and $n = 14$ branches of earlier versions are retired. n counts independent biological replicate cultures, not wells; the simulation assumes independent culture-level errors, while real cultures remain nested within one donor and share passage and plate effects, which will be recorded and modelled, with inference limited to the sampled lines. None of the inputs is measured in these cells, CV_bio = 0.20 is an assumption (0.10–0.30 moves power at $n = 3$ between 0.23 and 0.12), and our own dose–response analysis leaves the $4\times$ unsupported and unrefuted (potencia/). Twelve cultures per arm from one biopsy is heavier than an eight-week assay window implies, and we say so rather than keep the smaller number.

7. Actionable now, without any drug

The clearest clinical action supported by published guidance is surveillance, and there is an MVA-specific consensus (PMID 39264246, verified against the full text): renal ultrasound every 3 months from birth to age 7 for all MVA conditions, regular clinical assessment for rhabdomyosarcoma and other malignancies, avoidance of radiation, HPV vaccination. Cancers in MVA1 are predominantly Wilms tumour and rhabdomyosarcoma, with MDS, AML and ALL reported. **The gap that matters for this child:** the window closes at seven, and a 2026 MVA3 case had an orbital embryonal rhabdomyosarcoma at twelve (PMID 42595739), at a site renal

ultrasound does not image. The consensus asks for regular clinical assessment and sets no end date for it; what we add, for a child who has already had an RMS, is a frequency and a review point at age 10, as local suggestions rather than an extension of any stated limit.

Period	Surveillance	Frequency
Birth to 7	Renal ultrasound (we suggest widening the field to liver, retroperitoneum, pelvis)	3-monthly — consensus ; the wider field is ours
To 10, then reviewed	Full paediatric and oncological examination: abdomen, nodes, head and neck, orbit , skin, genitalia, soft tissues	3-monthly to 7, then 6-monthly — the frequency and the age-10 review point are our suggestion ; the examination itself is consensus
Post-ERMS	Relapse follow-up per tumour protocol	Avoid duplicate imaging
Through childhood	Full blood count with differential	No demonstrated benefit ; whether to do it, and how often, is an individual clinic decision
Lifelong	Family education; annual predisposition clinic	Urgent review for mass, haematuria, pain, deficit, bleeding, weight loss

The child's current age is not in the dataset and must be confirmed first. No serial AFP, CT, PET-CT or routine whole-body MRI while asymptomatic. If the child is past seven, missed scans should not be “caught up”; a baseline assessment and an individualised decision by the predisposition team is the right course. The whole table is a discussion framework for an expert predisposition clinic.

For the parents. The clinical document records recurrent spontaneous abortion (HP:0200067). If the variants are in *trans* — the interpretation adopted here, not demonstrated — each parent would be expected to carry one; segregation testing establishes that and must precede any carrier-specific interpretation. A 2026 study of two families with recurrent pregnancy loss found heterozygous *BUB1B* variants (classified VUS) with elevated premature chromatid separation in the carriers' lymphocytes (PMID 42434306): a hypothesis, cheap to act on. A PCS assay on parental lymphocytes would test whether the parents show that cellular phenotype — not that it caused

the losses — and would inform any future pregnancy. **The laboratory request is written** (`clinico/PARENTAL_SEGREGATION_ORDER.md`): both loci, two Sanger amplicons with primers verified as single-copy across GRCh38 (the first design returned sat in an *AluY* with ~100 copies; RepeatMasker exclusion caught it), and an interpretation table written before the result covering all five outcomes — including *cis*, which would refute our own Track 1 interpretation. A *trans* result contributes PM3-type evidence; **it does not by itself reclassify the missense.**

8. What we do not propose

MPS1/TTK inhibitors and reversine (not approved; they weaken a checkpoint that is already insufficient); apcin and proTAME (APC/C tool compounds, not marketed); DCZ0415 (a TRIP13 inhibitor: wrong gene, not approved); AICAR and 17-AAG (screen hits, not marketed; positive controls only); metformin (§3.1); senolytics (the BubR1 evidence rests on genetic clearance of p16⁺ cells, PMID 22048312, and senescence is an anti-tumour barrier worth keeping here); mTOR inhibitors with proteasome inhibition (§3.4); bortezomib for this child today (§5).

9. Scalability

- **The five-filter worksheet** (below) applies *as questions* to any repurposing exercise in a rare disease — filter 3 as an exposure plausibility screen, filter 5 as the MVA instance of “does treatment worsen the disease-relevant defect or select for it?” — and it is the transferable product.
- **The MVA genes are not interchangeable.** Transferred prospectively (`report/TRANSFER_MVA2_MVA3.md`) the worksheet returns different verdicts: for **MVA2** (*CEP57*), under the same standards as MVA1, filter 2 passes on plausibility (constitutional mosaic aneuploidy is the defining phenotype, PMID 21552266), filter 3 is the same generic cancer-model comparison with disease-specific transfer unestablished for all three genes, and filter 4 is not evaluable without a patient — the literature supplies no tumour-treatment setting, since an active malignancy has not been reported in the 15 published individuals (PMID 39264246; a statement about the literature, not of zero risk); for **MVA3** (*TRIP13*),

which shares the embryonal-tumour risk and severe checkpoint failure (PMID 28553959), it stays conditional. That shows the worksheet responds to the evidence available for each gene, not that the drug behaves differently in them. UBE2H is the target to watch (PMID 42094535); it has no approved inhibitor.

• **The retrieval benchmark is a tool, and we tested the claim.**

`replay/mva_replay.py` plants two alleles into a public genome and returns score, rank, a ClinVar-whitelist-off counterfactual and an unrelated-phenotype control. Pre-registered and pushed before the run, from a **fresh clone** of the repository, it planted two public Likely-pathogenic *TRIP13* alleles — chosen by the frozen pool rule, not by us — into three healthy GIAB genomes, queried with six HPO terms transcribed from the abstract of PMID 42595739:

Background	Planted, case phenotype	Unrelated phenotype	Unspiked
HG001	0.8332, rank 1	0.4240, rank 1	absent
HG002	0.8332, rank 1	0.4240, rank 2	0.0, rank 212
HG005	0.8332, rank 1	0.4240, rank 3	absent

The same pattern as BUB1B (rank 1/2/3 under an unrelated phenotype). Each background took under 3½ minutes and about 2.4 GiB. The operational half was the useful one: four provisioning steps, three of which the requirements list did not include, and regression tests the README said ran from a bare checkout did not. Both fixed; the provisioning script and every tool invocation with its exit code are committed (`replay/TRANSFER_TRIP13.md`). It measures retrieval given a perfect call, not diagnostic sensitivity. - **One command regenerates the Track 2 package.** `analysis/make_track2.sh` re-runs the DepMap follow-up (figure and tables) when the public downloads are present, renders the worksheet from `candidates.tsv`, checks a declared set of reported numbers against the results files (27 regex anchors — 10 Track 1, 15 Track 2 report, 2 methods form — covering 31 recomputed facts; not every number in the document) and builds the PDFs. Without the public downloads it rebuilds the document from committed results; with them it reruns the specified analyses, and the log says which mode ran. Run from a fresh checkout of the repository it finished in 8 s and 179 MB, naming the one missing input (the ~700 MB DepMap/PRISM downloads, which it does not bundle) and building both PDFs from the committed copies (`evidencia/`

make_track2_fresh_checkout_2026-09-08.log). - The gene panel and artefact controls transfer to any MVA or PCS workup; the micronucleus → scDNA-seq endpoint applies to mitotic CIN disorders, **not** to Fanconi anaemia or Bloom syndrome.

10. Limitations

- **Phase is not established: the called heterozygous markers do not support ordinary short-fragment bridging, and the tested reference-panel route was uninformative.** No parental samples; the challenge distributes a VCF and raw reads but no alignments, we did not align the reads, and short reads could not phase two variants whose only intervening heterozygous site lies 6.8 and 4.1 kb away; all three sites are absent from the 1000 Genomes phased panel (expected copies 0.64 and 0.006 in 6,404 haplotypes). Nothing in the data weighs *cis* over *trans*; what supports *trans* is that the child is affected, MVA1 is recessive, truncating-plus-missense is the recurrent pattern, and the phenotype matches. **A diagnosis pending segregation.** The retrieval pipeline itself is insensitive to phase: two nonsense alleles declared in *cis* are still called compound heterozygous and ranked first (score difference 0.0034).
- **p.Asn1002Lys has no functional assay**, and its apparent ClinVar entry is a different nucleotide.
- **The bortezomib EC50 is a bound read from a figure; the Cmax is a label interval; the ratio is on total drug** and says nothing about free drug, tissue or duration. No effective concentration has been measured in constitutional MVA cells.
- **Ippolito studied cancer aneuploidy.** Section 6 measures differential drug response and constitutional-cell vulnerability; mechanistic attribution and tumour selectivity require separate experiments, and we have not assumed the transfer. Our public-data test weakens the solid-tumour case and cannot speak to MVA cells.
- **The power estimate rests on unmeasured inputs:** CV_bio = 0.20 is assumed, the aneuploid fraction is uncounted, the 4× selectivity is unsupported; under the declared model the design needs about 12 cultures per arm, and 0.897 is an optimistic two-parameter reference.

- **The mosaicism control was not conclusive;** the patient-level result is withdrawn; no clinical negative.
 - **No trial specific to MVA1** was found in ClinicalTrials.gov, Open Targets, ChEMBL or CMap/LINCS; absence of an interface result is not proof of absence.
 - **The child's current age is unknown to us.**
-

A closing note

The family published their son's genome so that strangers might help. The honest answer to “what drug should he take today” is none. What we can leave them is a mechanism that matches his disease, an approved drug whose plasma levels are not excluded by the concentrations at which cancer cells respond, the condition under which it would be worth asking about — a changed comparator, not an established benefit — the experiment that could fail to support it in his own cells, and the published surveillance guidance to review with his treating and predisposition teams. That is not a cure. It is a prepared answer to a question this family may unfortunately have to ask.

The five-filter worksheet (reusable)

Rendered from `report/candidates.tsv` by `report/render_candidates.py`; the table below and the machine-readable file are the same object. Statuses only here; the model context and reason behind every cell are in the Part II worksheet and in the file. Every candidate this report evaluated is a row; add yours.

Candidate	F1 approved	F2 mechanism	F3 exposure	F4 safe in <i>this</i> patient	F5 margin excluded	Verdict
metformin	pass	pass	not evaluable	not reached	not reached	rejected at filter 3 unevaluable
bortezomib IV	pass	pass	not excluded	fail	untested	fails filter 4 for this child; conditional for active tumour
bortezomib SC	pass	pass	exposure unresolved	fail	untested	exposure unresolved; fails filter 4
carfilzomib	pass	pass	not evaluable	not reached	not reached	not evaluable
ixazomib	pass	pass	not evaluable	not reached	not reached	not evaluable
hydroxychloroquine	pass	untested	not evaluable	not reached	not reached	not evaluable mechanistically weakened
everolimus sirolimus	pass	untested	not evaluable	not reached	not reached	not evaluable possibly counterproductive
trametinib	pass	pass	not evaluable	not reached	not reached	not evaluable gap most worth closing
AICAR	fail	pass	not reached	not reached	not reached	fails filter 1
17-AAG HSP90i	fail	pass	not reached	not reached	not reached	fails filter 1
reversine MPS1 TTK inhibitors	fail	fail	not reached	not reached	untested	fails filter 1 and 2
apcin proTAME	fail	untested	not reached	not reached	not reached	fails filter 1
DCZ0415	fail	fail	not reached	not reached	not reached	fails filter 1 and 2
senolytics	not evaluable	untested	not reached	not reached	untested	not proposed
UBE2H inhibitor	fail	pass	not reached	not reached	not reached	target to watch
<i>(your candidate)</i>						

How to fill it in. Filter 3 needs a real denominator measured in a model stratified for the disease mechanism; if none exists, write *not evaluable* rather than substitute an IC50 from an unrelated model. Filter 4 is a statement about one person and one moment, not about the drug. Filter 5 asks whether an increase in mis-segregation among survivors beyond a prespecified margin can be excluded. A machine-readable version with the full status vocabulary (pass / fail / untested / not evaluable / not reached / not excluded / exposure unresolved), source and reason is `report/candidates.tsv`.

References

1. Hanks S, et al. Constitutional aneuploidy and cancer predisposition caused by biallelic mutations in *BUB1B*. *Nat Genet* 2004. PMID 15475955
2. Suijkerbuijk SJE, et al. Molecular causes for BUBR1 dysfunction in mosaic variegated aneuploidy. *Cancer Res* 2010. PMID 20516114
3. Hanks S, et al. Comparative genomic hybridization and BUB1B mutation analyses in childhood cancers associated with MVA. *Cancer Lett* 2006. PMID 16182441
4. Tang YC, et al. Identification of aneuploidy-selective antiproliferation compounds. *Cell* 2011. PMID 21315436
5. Ippolito MR, et al. Increased RNA and Protein Degradation Is Required for Counteracting Transcriptional Burden and Proteotoxic Stress in Human Aneuploid Cells. *Cancer Discov* 2024. PMID 39247952
6. Sakellakis M. Why Metformin Should Not Be Used as an Oxidative Phosphorylation Inhibitor in Cancer Patients. 2023. PMID 37343530
7. Fontaine E. Metformin-Induced Mitochondrial Complex I Inhibition: Facts, Uncertainties, and Consequences. 2018. PMID 30619086
8. Ailabouni AS, et al. Interindividual variability in metformin pharmacokinetics in pediatric patients. 2026. PMID 42399684
9. Chui MH, et al. Chromosomal Instability and mTORC1 Activation through PTEN Loss Contribute to Proteotoxic Stress in Ovarian Carcinoma. *Cancer Res* 2019. PMID 31530568
10. Knouse KA, et al. Single cell sequencing reveals low levels of aneuploidy across mammalian tissues. *PNAS* 2014. PMID 25197050
11. Baker DJ, et al. BubR1 insufficiency causes early onset of aging-associated phenotypes. *Nat Genet* 2004. PMID 15208629
12. Baker DJ, et al. Clearance of p16^{Ink4a}-positive senescent cells delays ageing-associated disorders. *Nature* 2011. PMID 22048312
13. Sieben CJ, et al. BubR1 allelic effects drive phenotypic heterogeneity in mosaic variegated aneuploidy. *J Clin Invest* 2020. PMID 31738183
14. Nishitani-Isa M, et al. ERMS with PCS/MVA treated with reduced-intensity chemotherapy. *Pediatr Int* 2019. PMID 31184400
15. Scott RH, et al. Surveillance for Wilms tumour in at-risk children: pragmatic recommendations. *Arch Dis Child* 2006. PMID 16857697

16. Xu P, et al. BUBR1 recruits PP2A via the B56 family to promote chromosome congression. *Biol Open* 2013. PMID 23789096
17. Kruse T, et al. Direct binding between BubR1 and B56-PP2A regulates mitotic progression. *J Cell Sci* 2013. PMID 23345399
18. Snape K, et al. Mutations in *CEP57* cause mosaic variegated aneuploidy syndrome. *Nat Genet* 2011. PMID 21552266
19. Yost S, et al. Biallelic *TRIP13* mutations predispose to Wilms tumor and chromosome missegregation. *Nat Genet* 2017. PMID 28553959
20. FDA label, VELCADE (bortezomib), NDA 021602 s040, §12.3
21. FDA BPCA Clinical Review, bortezomib paediatric safety, NDA 021602, 2015
22. Aneuploid human colonic epithelial cells are sensitive to AICAR-induced growth inhibition through EGFR degradation. 2013. PMID 22890317
23. Nakano Y, Kuiper RP, et al. Update on Recommendations for Cancer Screening and Surveillance in Children with Genomic Instability Disorders. *Clin Cancer Res* 2024. PMID 39264246
24. Schukken KM, et al. Paired CRISPR screens identify mitochondrial metabolism and UBE2H as aneuploid-specific dependencies in human cancer cell lines. *bioRxiv* 2026 (preprint). PMID 42094535
25. Vempuluru VS, et al. Sequential presentation of Wilms' tumor and orbital rhabdomyosarcoma in a child with mosaic variegated aneuploidy syndrome 3. *Orbit* 2026. PMID 42595739
26. Wei TY, et al. Functional and clinical evidence for two novel heterozygous *BUB1B* variants and their value in precision genetic counseling for recurrent pregnancy loss. *Front Endocrinol* 2026. PMID 42434306
27. Escalante LE, Hose J, et al. Chromosome duplication causes premature aging via defects in ribosome quality control. *PLoS Biol* 2025. PMID 41248159
28. Bersani F, Taulli R, et al. Bortezomib-mediated proteasome inhibition as a potential strategy for the treatment of rhabdomyosarcoma. *Eur J Cancer* 2008. PMID 18342500
29. Houghton PJ, Morton CL, et al. Initial testing (stage 1) of the proteasome inhibitor bortezomib by the pediatric preclinical testing program. *Pediatr Blood Cancer* 2008. PMID 17420992
30. Blaney SM, Bernstein M, et al. Phase I study of the proteasome inhibitor bortezomib in pediatric patients with refractory solid tumors: a Children's Oncology Group study (ADVL0015). *J Clin Oncol* 2004. PMID 15570082
31. Muscal JA, Thompson PA, et al. A phase I trial of vorinostat and bortezomib in children with refractory or recurrent solid tumors: a Children's Oncology Group phase I consortium study (ADVL0916). *Pediatr Blood Cancer* 2013. PMID 22887890
32. Maki RG, Kraft AS, et al. A multicenter Phase II study of bortezomib in recurrent or metastatic sarcomas. *Cancer* 2005. PMID 15739208
33. Human aneuploid cells depend on the RAF/MEK/ERK pathway for overcoming increased DNA damage. *Nat Commun* 2024. PMID 39251587

Acknowledgement

This work was made possible through the Hackathon, organized by Sage Bionetworks in partnership with the MVA Society, Hugging Face, and BEACON (The Benchmarking, Evaluation, and Assessment

Consortium for Science), with prize sponsorship from AWS and Anthropic. We are deeply grateful to the child and their family who generously contributed their data and their story to advance research into this rare disease. We acknowledge their trust in making this Hackathon possible.

Data availability

The individual-level data analysed here is the Hackathon dataset (SageBio/mva-hackathon-2026-data), accessed under the challenge data transfer agreement accepted by Hugging Face account `cpu-16` on 2026-08-28, and cited as directed on the Hackathon Synapse page. No patient sequence data or genotype-scale derivative is included in this document or in our public repository, and every copy we hold will be deleted by 2026-11-23 with written confirmation to the organisers. The two diagnostic variants are named, as findings, which the challenge rules permit. All other inputs are public: Exomiser 15.1.0 with data release 2602, ClinVar, gnomAD v4, the Human Phenotype Ontology, Ensembl and VEP, the 1000 Genomes and GIAB genomes, DepMap 24Q4, AlphaFold and the literature cited by PMID. Analysis code, pre-registrations and results: <https://github.com/cpu-16/mva-hackathon-2026>

Part II — Full evidence

The complete report: every section at full length, with derivations, sources quoted verbatim, the correction history and every pre-registered analysis. Part I is the reading version; where the two differ in wording, this part carries the detail and Part I never claims more than this part supports. Section numbers differ between the parts: Part I §4 (public-data test) is Part II §3.5, Part I §5 (safety) is Part II §4, Part I §6 (endpoints and experiment) spans Part II §5–6; §1–3 and §7–10 correspond.

Executive summary

We organised the assessment around five explicit questions, and the third and fourth change the answer.

The candidate the literature points to is **metformin**. It descends from the AICAR hit in the aneuploidy-selective screen of Tang et al. (PMID 21315436), it is approved, and its mechanism is plausible. Its aneuploidy-selective exposure cannot be evaluated, and a separate concentration mismatch weakens the proposed complex-I mechanism: plasma concentrations at therapeutic doses are **micromolar**, while complex I inhibition requires **millimolar** (PMID 37343530) — a gap of roughly **three orders of magnitude (~1000-fold)**. We expect metformin to be widely proposed in this track. **It cannot inhibit complex I at therapeutic plasma concentrations.** Whether it could still stress aneuploid cells through AMPK activation at those concentrations is a *different* claim, and an untested one: no study reports a metformin aneuploidy-selective effective concentration. We reject it as unevaluable on the mechanism it is proposed for, not as disproven — the distinction is set out in §3.1.

The candidate the pharmacology does not exclude is **bortezomib**. Aneuploid cells carry a stoichiometric protein imbalance and compensate by increasing protein degradation, which makes them **preferentially dependent on the proteasome**. Ippolito et al. (2024, *Cancer Discovery*, PMID 39247952) established this across multiple diploid-versus-aneuploid models, recapitulated it in hundreds of cancer cell lines and primary tumours, and found that **aneuploidy level was significantly associated with the response of multiple myeloma patients to proteasome inhibitors** — a human clinical association rather than a pure cell-culture inference.

We state the size of the remaining leap rather than hide it. That evidence is drawn from **cancer** aneuploidy, and myeloma is a plasma-cell neoplasm carrying its own immunoglobulin-driven proteotoxic load. Whether the same dependency holds in **constitutional, variegated** aneuploidy is an extrapolation, and it is exactly what the experiment in §6 is designed to break (an eight-week assay window; **12–16 weeks** from a fresh biopsy, because establishing the fibroblast line takes 4–8 weeks on its own — more if twelve independent cultures per arm must be expanded, §6).

We then tested the part of that leap that public data can test, and the predicted drug association failed its criterion. Before looking at any result we pre-registered a test on public data (§3.5): does the

aneuploidy–proteasome link extend beyond haematological cancers? Across 424–444 evaluable solid-tumour lines (depending on the drug) no proteasome inhibitor reached significance, and across 946 lines the genetic dependency on **PSMB5 — bortezomib’s own target — is significantly weaker in solid tumours than in haematological ones** (interaction $p = 0.026$), with ploidy not significant after adjustment ($p = 0.82$ in the interaction model, 0.875 in the main-effects model) — which does not exclude residual confounding. This child’s tumour was a solid one. We report this because it weakens the argument we are making, and because we would rather find it than have a judge find it.

Bortezomib is not excluded by the pharmacokinetic filter as we define it: EC50 in highly aneuploid lines is below **40 nM**, while the approved 1.3 mg/m^2 IV dose reaches a Cmax of **89–120 ng/mL = 231–312 nM** — a ratio of **5.8–7.8**. We quote the twice-weekly range the label reports rather than its ceiling (the same label gives 223 ng/mL in its IV-vs-SC cohort, §3.2); read it as an order-of-magnitude comparison on total drug, not a measured margin — on free drug the peak is $\approx 39\text{--}53 \text{ nM}$, about 1.0–1.3 times the 40 nM reference value ($\approx 99 \text{ nM}$, 2.5 times, on the label’s 223 ng/mL cohort), ratios to a reference denominator rather than a measured free-drug margin. **And read what the filter is:** Cmax is *total* plasma drug at a peak. It is not free drug, not intratumoral drug, and it says nothing about how long the concentration is held — bortezomib’s plasma level falls steeply after infusion while proteasome inhibition persists, so neither exposure nor duration at the target is established by this number. **312 nM is the top of a label interval, not a validated clinical threshold**, and no aneuploidy- selective effective concentration has ever been measured in constitutional MVA cells. Filter 3 tests whether the drug is disqualified on order of magnitude. It is not excluded by that test, and nothing stronger follows.

The controlling conclusion, used throughout this document. No disease-modifying drug is supported for use in this child today. Bortezomib is a conditional research candidate: the stated IV total-plasma peak comparison does not exclude it, but comparable free exposure, duration at the target and tumour selectivity in MVA remain unestablished. PRISM did not detect a significant proteasome-inhibitor association in the tested solid lines; CRISPR showed a weaker

aneuploidy-PSMB5 association in solid than in haematological lines. An active malignancy would change the clinical comparator; it would not establish a favourable benefit-risk balance.

It then fails our fourth filter, and we say so plainly: **not for this child, not now**. In the paediatric trial that measured it, peripheral neuropathy occurred in 18% of patients (25/140; 8% motor) receiving bortezomib with intensive combination chemotherapy — his incremental risk from the drug alone is unknown — and this child already has skeletal muscle atrophy. Outside an active malignancy that risk-benefit does not hold.

Eight pre-registrations were committed to git before their analyses ran — DepMap and its follow-up, MVA-Replay, mosaicism, power and its four-parameter follow-up, dose-response, and the TRIP13 transfer. PRISM did not establish the predicted drug association in solid lines and CRISPR supported a weaker target association in solid than in haematological lines, the mosaicism control came back not conclusive, the power analyses showed the assay as first sized ($n = 3$) to be underpowered even under its own assumptions — about 12 cultures per arm are needed — and the dose-response analysis left the assay's assumed 4× within-cell selectivity (separate from the exposure ratio) unsupported; all are reported.

What remains is a specific and testable proposition, not a demonstrated one. Bortezomib does not repair the lesion — nothing restores BubR1 function — but it is the agent aimed at the best-characterised *consequence* of aneuploidy **in cancer cells**, and it is one conditional research candidate worth documenting as a question (no priority over unresolved alternatives such as trametinib is established) for an MVA patient **with an active tumour**, where the comparator is not “no drug” but conventional cytotoxic chemotherapy. We have not shown that dependency exists in this child's cells, and §6 exists because it might not. This child is an embryonal rhabdomyosarcoma survivor with a high risk of embryonal tumours (PMID 28553959) — and he has already had one. That sequence is not hypothetical: a 2026 case report describes a girl with MVA3 who had a Wilms tumour at three and an orbital embryonal rhabdomyosarcoma at **twelve** (PMID

42595739). Our contribution is to have the question ready — with its evidence, its exposure comparison and its stopping rule — before the decision has to be made under pressure.

We also flag an antagonism a combination-minded team could walk into: **mTOR inhibitors would be expected to protect aneuploid cells from proteasome inhibition, not synergise** (PMID 31530568).

1. The variants and what they break


The proband carries two heterozygous **BUB1B** variants (NM_001211.6, transcript ENST00000287598.11, MANE), read here as compound heterozygous **on presumption: phase is not established (\$10)**:

Variant	Consequence	Exomiser ACMG	ClinVar
c.2210T>G p. (Leu737Ter)	stop_gained	PATHOGENIC (PVS1, PM2_Supporting, PP4_Moderate, PP5_Strong)	533901 — Pathogenic/Likely pathogenic
c.3006T>G p. (Asn1002Lys)	missense	UNCERTAIN_SIGNIFICANCE (PM2_Supporting, PP4_Moderate, BP1)	not present

Population frequency depends on the source, and we report both rather than select the lower one. Exomiser's gnomAD snapshot gives 9.98×10^{-5} (exomes, NFE) for the nonsense and 9.0×10^{-7} for the missense; Ensembl VEP returned 3×10^{-5} and no record respectively. Both are rare; neither is absent.

A note of caution on the missense. ClinVar holds c.3006T>A, a *different* nucleotide substitution producing the same p.Asn1002Lys, classified VUS. That classification does not transfer to our allele: PS1 applies when the known same-amino-acid variant is pathogenic, and here it is uncertain. Our variant is simply not in ClinVar.

BubR1 is a core component of the mitotic checkpoint complex and of kinetochore–microtubule attachment control; biallelic *BUB1B* mutation was the founding cause of constitutional aneuploidy with cancer predisposition (Hanks 2004, PMID 15475955), and BubR1 insufficiency in mice produces the progeroid phenotype that makes MVA an aging syndrome as well as a cancer syndrome (PMID 15208629).



Domain map of BubR1 showing both alleles converging on the C-terminal kinase domain

p.Leu737Ter truncates in the linker N-terminal to the C-terminal kinase fold — UniProt O60566 annotates the protein kinase domain at residues **766–1050** — removing it entirely — residues 737–1050, 314 of the protein’s 1050, if the transcript is translated at all. (Whether BubR1 is catalytically active in humans is contested; we call it the kinase domain because that is the current UniProt annotation, and our argument rests on the truncation, not on catalysis.) It lands 16 residues upstream of **BUBR1^{X753}**, a characterised truncating MVA allele (PMID 31738183) — very nearly the same truncation, not merely a related model.

p.Asn1002Lys falls inside that same domain, in its C-lobe, outside KEN, TPR and GLEBS. **No published functional assay of this substitution exists**, and we measured what can be said without one (`vus/RESULTADOS.md`).

On the AlphaFold model, which is high-confidence exactly here — mean pLDDT **91.5** across 992–1012 and **91.1** at the residue itself, against 63.6 for the whole protein — **Asn1002 is buried**: relative solvent accessibility **0.162**, below the conventional 0.25 cutoff, with ten residues within 5 Å: by distance, Leu1001 and Ala1003 (1.34 Å), Trp978 and Val998 (2.83), Asn1004 (3.14), Ile1000 (3.33), Phe977 (3.46), Arg999 (3.62), Trp973 (4.36) and Phe997 (4.51). Eight of the ten are hydrophobic; Asn1004 and Arg999 are not, and we list them rather than quote only the apolar ones. The substitution puts a **longer, positively charged side chain into a buried hydrophobic pocket**, which is a recognised destabilising substitution class. We state that as evidence about the substitution and **not** as a classification: no $\Delta\Delta G$ was

computed; the REVEL (0.472) and AlphaMissense (0.923) values quoted in our methods description were retrieved as annotations, not computed by us, and they disagree; we do not claim PP3.

And we can now say why no database will settle it. The obvious move is to calibrate a predictor against BUB1B variants that are already classified. That calibration set does not exist. In the ClinVar 2026-08-28 release this gene carries **1,497 missense records, of which exactly one is Likely pathogenic** — at one star, from a single submitter — against **1,426 (95.3%) of uncertain significance**. Meanwhile **82 of its 94 truncating records are Pathogenic or Likely pathogenic**. The clinical evidence base of *BUB1B* is built on loss of function at a ratio of 82 to 1. In this snapshot, **1 of 1,497 submitted missense records** has reached P/LP.

Read that number for what it is: a property of the database, not a probability about this child's allele. It is the rate at which *submitted* missense records in this gene have been *classified* P/LP, and submission and classification are driven by ascertainment, testing volume and submitter policy. It is **not** the prior probability that a given missense in *BUB1B* is causal, and we do not use it as one. Missense alleles here may simply be under-tested and under-submitted: Sieben's characterised **BUBR1^{L1012P}** is a missense causal enough to build a mouse model on, and we did not find it in this snapshot.

What the census does establish is narrower and still useful: **the calibration set does not exist**. A predictor cannot be calibrated on one positive, so no in-silico classification of this allele can be anchored to same-gene truth. That is the quantitative reason the parental segregation test in §7 is the cheapest informative route available to us for this allele — **not the only conceivable one**. Functional assays on BubR1 (kinetochore localisation, PP2A-B56 binding, checkpoint output) exist and would speak to the substitution directly; they need material and a laboratory that we do not have, and they answer a different question than phase does.

The closest characterised missense sits ten residues away, and the mouse work on it carries a warning for us. Sieben et al. modelled the human MVA missense **BUBR1^{L1012P}** in mice and paired it with **BUBR1^{X753}** — the same truncating-plus-missense architecture as our

proband, with the missense in the same domain. Those BubR1^{X753/L1012P} animals **died prematurely**; the viable genotype paired the missense with a *hypomorphic* allele that still yields some wild-type protein (PMID 31738183).

Our proband is alive carrying a truncating allele. The mouse allelic series motivates testing whether p.Asn1002Lys **retains residual function** — a hypomorph rather than a functional null — but survival across species and allele combinations does not establish its functional rank relative to L1012P. That is a hypothesis about our own VUS; the allele-fate assays of §6 test transcript and protein consequences, and a comparative checkpoint assay with an L1012P comparator would be needed to test residual function itself.


Whether the truncated transcript escapes nonsense-mediated decay is unmeasured, and should not be asserted either way without patient RNA.

At the cellular level, cell lines derived from MVA patients with biallelic mutations show “an impaired mitotic checkpoint, chromosome alignment defects, and low overall BUBR1 abundance” (PMID 20516114). BubR1 binds the B56 family of PP2A regulatory subunits through a conserved motif phosphorylated by Cdk1 and Plk1, and counteracts Aurora B at improperly attached kinetochores (Kruse et al., *J Cell Sci* 2013, PMID 23345399); the same B56 partnership opposes Aurora B to give stable kinetochore–microtubule attachment (Xu et al., *Biol Open* 2013, PMID 23789096). The region is commonly called KARD, a name neither source uses.

2. Five filters

We organised this assessment around five explicit questions. The first two are the ones every exercise asks: **is it approved?** and **is the mechanism plausible?** Both are necessary; neither is sufficient. Three more decide the outcome.

#	Filter	Question	What it eliminates here
1	Regulatory	Approved for marketing?	AICAR, 17-AAG, reversine, apcin, proTAME, DCZ0415
2	Mechanistic	Does it address the actual lesion or a direct consequence of it?	Agents acting on uninvolved pathways
3	Pharmacokinetic	Is the achievable exposure excluded by the effective in-vitro concentration? (“not excluded” is the best outcome this screen can return)	Metformin: no aneuploidy-selective potency; separate ~1000-fold complex-I mismatch
4	Patient-specific safety	Compatible with <i>this</i> child’s comorbidities?	Bortezomib, at present
5	Direction of effect	Can an increase in mis-segregation among survivors beyond a prespecified margin be excluded?	An increase exceeding the chosen micronucleus margin; longer-term clonal selection remains untested



The five-filter funnel: what each filter eliminates, and why nothing passes all five today

Filters 3 and 5 make exposure plausibility and disease-worsening risk explicit. Filter 3 is a possibility check, not a demonstration of efficacy: it ignores protein binding, exposure duration and tissue penetration, and a transient post-bolus C_{max} can overstate useful exposure.

3. Applying the filters

3.1 Metformin: not evaluable on the mechanism proposed for it

The chain of reasoning is genuinely attractive and we followed it to the end. Tang et al. screened isogenic trisomic mouse fibroblasts and found selective sensitivity to **AICAR**. AICAR is not approved, so it fails filter 1. Metformin is the marketed stand-in that also produces AMPK activation through complex I inhibition, and it is labelled in paediatrics from age 10. It passes filters 1 and 2.

On the complex-I mechanism it fails filter 3 by three orders of magnitude; on aneuploidy-selective activity it is not evaluable, and that is the verdict the worksheet carries. Sakellakis (2023, PMID 37343530):

“the plasma concentration of metformin remains within the micromolar range when administered at therapeutic doses. While millimolar concentrations are necessary to inhibit complex I activity in isolated mitochondria, there is no evidence supporting the idea that metformin accumulates within the mitochondria.”

Fontaine (2018, PMID 30619086) reviews why that inhibition remains unsettled — that review exists to “gather arguments for and against each hypothesis.” Paediatric exposure is variable and depends partly on renal function — it does not narrow the gap: dose-normalised plasma exposure varies **3.3-fold** between children, and “**renal clearance was higher in children compared to the reported data in healthy adults**”, with eGFR significantly affecting oral clearance (PMID 42399684). Metformin is renally cleared, so this child’s nephrocalcinosis makes renal function the parameter that would have to be established first — the source links clearance to eGFR, not to nephrocalcinosis specifically, and we do not stretch it further.

The distinction that decides this, stated precisely. Therapeutic metformin *does* activate AMPK in patients — that is the diabetes literature and we do not dispute it. What the micromolar-versus-millimolar gap rules out is the **complex I route at plasma concentrations**. Tang’s screening hit was AICAR, a direct AMPK agonist, so an AMPK-mediated aneuploidy-selective effect of metformin is not excluded by this argument; it is simply **untested**. No study reports a metformin aneuploidy-selective effective concentration, and the one paper reproducing AICAR’s selectivity in human cells does not test metformin. We therefore reject metformin as **unevaluable on the mechanism it is proposed for**, not as disproven — and we say which of the two we mean, because they license different next steps: the first is answered by an experiment, the second would close the question.

We do not reject metformin out of caution. We cannot evaluate its aneuploidy-selective exposure, and a separate concentration comparison weakens the proposed complex-I rationale — and we state what kind of number that is. For bortezomib, filter 3 compares total-plasma C_{max} with a 40 nM reference upper bound read from an aneuploidy-stratified experiment. For metformin no such EC₅₀ exists: no published study gives an aneuploidy-selective effective concentration. AICAR’s aneuploidy selectivity has in fact been

reproduced in human cells — a trisomy-7 colonic line is fourfold more sensitive to AICAR than its isogenic diploid counterpart (PMID 22890317). **That paper attributes the effect to proteasomal EGFR degradation, not to AMPK**, which weakens rather than supports the AMPK route we are discussing here. It does not test metformin, and no published study gives metformin an aneuploidy-selective effective concentration at all. The 1000-fold figure is therefore a comparison between *concentration ranges* — micromolar plasma against millimolar complex I inhibition — not the same quantity as the bortezomib ratio. It is a coarser instrument, and it still points the same way.

There is a deeper objection, and stating it accurately matters. Tang et al. did **not** restrict themselves to stable trisomies: Figure 6 of that paper tests AICAR and 17-AAG in *Bub1b*^{H/H} hypomorphic MEFs and in *Cdc20*^{AAA} cells — genuine ongoing-instability models with a weak checkpoint. So the screen does reach beyond a fixed extra chromosome.

What it does not reach is **metformin**, which is not tested in those models, and human biallelic BUB1B cells, which are not tested at all. The gap is therefore narrower and more specific than “stable trisomies only”: the marketed stand-in was never evaluated in the instability models, and a mouse *Bub1b*^{H/H} hypomorph is not equivalent to human p.Leu737Ter/p.Asn1002Lys.

3.2 Bortezomib passes filters 1–2; not excluded by filter 3 on total-drug exposure

Mechanism. Aneuploidy imposes a stoichiometric imbalance of protein complexes. Ippolito et al. (2024, PMID 39247952) dissected multiple diploid-versus-aneuploid models and found that aneuploid cells mitigate proteotoxic stress by reducing translation and **increasing protein degradation**, rendering them more sensitive to proteasome inhibition. The finding was recapitulated across hundreds of human cancer cell lines and primary tumours.

An independent, orthogonal 2026 result points at the same target. Paired genome-wide CRISPR loss-of-function screens in isogenic aneuploid versus near-euploid lines identify **proteasome subunits** among the top aneuploid-specific dependency gene groups, alongside ribosome and rRNA processing, spliceosome and mitochondrial

metabolism; a focused druggable-genome arm nominates the ubiquitin-conjugating enzyme **UBE2H** as a top aneuploid-selective dependency (PMID 42094535). This is genetic loss-of-function evidence rather than pharmacology, arrived at by a different method than Ippolito's, and it converges on the ubiquitin–proteasome axis. **It is a preprint (bioRxiv) and not yet peer-reviewed**, so we weight it as corroboration, not as a second pillar. UBE2H itself has no approved inhibitor and therefore fails filter 1 today; we record it in §9 as the target to watch.

The clinically important sentence in **Ippolito et al.** is the last one: **aneuploidy level was significantly associated with the response of multiple myeloma patients to proteasome inhibitors**. That is a human signal linking degree of aneuploidy to response to this drug class, which is more than most repurposing arguments have.

We state the size of that clinical signal rather than let a reviewer find it. It rests on small cohorts: with bortezomib as a single agent, 8 patients with a complete response had significantly higher aneuploidy scores than 50 with progressive disease ($p = 0.014$, one-tailed Mann-Whitney); in a bortezomib-plus-chemotherapy-and-dexamethasone cohort, 13 complete responders versus 14 minimal responders ($p = 0.038$, one-tailed; the figure legend calls that comparator group *minimal response* while the Results text calls it *progressive disease* — we report the legend and do not resolve the discrepancy); a third cohort trended the same way but had only 2 non-responders. That is a genuine human association built on small numbers — enough to justify running an experiment, not enough to predict an individual patient's response.

It is also not evidence about MVA1. Myeloma is a plasma-cell neoplasm whose immunoglobulin output imposes a proteotoxic load of its own, and its aneuploidy is clonal, not variegated. MVA1 is a constitutional aneuploidy state — a *different* setting in which the same dependency is plausible and untested. We treat the myeloma association as the reason to run the experiment, not as its result.

The numbers.

	Value	Source
EC50 in highly aneuploid lines (72 h)	< 40 nM — conservative upper bound read from Fig. 6o (5 near-euploid vs 5 highly aneuploid cancer lines, $p = 0.0317$, Mann-Whitney); the figure legend reports no numeric EC50 values, so this is read off the plot	PMID 39247952
Cmax, 1.3 mg/m ² IV, repeated dosing	89–120 ng/mL = 231–312 nM (the twice-weekly range; the same label gives 223 ng/mL in its IV-vs-SC cohort — see the exposure table below)	FDA label, NDA 021602 s040, §12.3
Cmax, subcutaneous route	20.4 ng/mL = 53.1 nM	same
Filter 3 ratio (IV)	5.8 – 7.8	—

We use the upper bound of the EC50 range, and we do not treat the 2.4 nM apoptosis figure from the same paper as an EC50.

The exposure evidence, laid out rather than reduced to a pass (label facts from the current DailyMed SPL, archived as `evidencia/velcade_spl_dailymed_2026-09-07.xml`):

Quantity	Value	What it does and does not support
Effective concentration	EC50 < 40 nM at 72 h in 5 highly aneuploid cancer lines, read from Fig. 6o (PMID 39247952); the legend tabulates no value; culture medium and serum binding are unstated	A bound on total drug in culture, not a threshold for constitutional MVA cells
Plasma protein binding	“averaged 83% over the concentration range of 100 to 1000 ng/mL” (§12.3)	Free fraction $\approx$ 17%, applied below 100 ng/mL as an approximation
Cmax, IV 1.3 mg/m ²	112 ng/mL after the first dose; 89–120 ng/mL twice weekly; 223 ng/mL in the repeat-dose IV-vs-SC comparison cohort (§12.3)	Total drug at a peak; free peak $\approx$ 39–53 nM from 231–312 nM, or $\approx$ 99 nM from 223 ng/mL
Cmax, SC 1.3 mg/m ²	20.4 ng/mL; AUC equivalent to IV (geometric mean ratio 0.99) (§12.3)	Free peak $\approx$ 9 nM
Pharmacodynamics	Maximal 20S proteasome inhibition 73–83% in whole blood, observed 5 min after the 1.3 mg/m ² dose (§12.2); mean elimination half-life 76–108 h on multiple dosing (§12.3) while plasma falls steeply after infusion	Neither duration of inhibition at a tumour nor recovery is established by these numbers
Not available	Free-drug EC50; intratumoral concentration; time above EC50	Marked missing rather than imputed

On total drug the IV ratio is 5.8–7.8 (231/40 and 312/40). **On free drug the peak is $\approx$ 39–53 nM, that is 0.98–1.33 times the 40 nM reference value, or $\approx$ 99 nM (2.5 times) on the 223 ng/mL cohort** — ratios to a reference denominator, not measured free-plasma-to-free-medium margins; and since 40 nM is an upper bound on the EC50, the true free-drug margin is unknown; the culture-side value is itself partly free drug in serum-containing medium, so the two sides are not on the same footing and we do not correct one for binding while leaving the other unexplained. What filter 3 supports is that the IV route is **not excluded on the stated total-plasma peak comparison**; it does not establish comparable free exposure or duration at the target.

The subcutaneous route is exposure-unresolved, and we withdraw the earlier claim that it passes. Its nominal ratio is 1.33 on *total* drug at the peak (53.1/40), with AUC equivalent to IV; bortezomib is extensively protein-bound, so the free peak is ≈ 9 nM. We do not call that a pass, and we do not call it a fail either: the free-fraction correction we decline to apply to one side of the IV comparison cannot be applied to one side of the SC comparison. The IV ratio of 5.8–7.8 is left on total drug and read only as an order of magnitude for the same reason — the free-fraction and duration corrections are unknown in both cases, and the culture-side concentration carries its own unstated medium and binding conditions. Filter 3 is a plausibility screen; it says the IV route is not excluded, and nothing stronger (§2).

3.3 Everything else: not evaluable, and we say so

Drug	Filter 3 verdict	Why
Carfilzomib	Not evaluable	No published IC50 in an aneuploidy-stratified model; Ippolito used bortezomib and MG132. Cmax is known (2.89 μ M at 56 mg/m ²) but there is no valid denominator
Ixazomib	Not evaluable	Same gap. The β 5 enzymatic IC50 (3.4 nM) is not a cellular aneuploidy-selective value
Hydroxychloroquine	Not evaluable, and mechanistically weakened	Tang's hit was <i>chloroquine</i> , not HCQ — and chloroquine did not differentially inhibit the human CIN lines in that same paper
Everolimus / sirolimus	Not evaluable, and possibly counterproductive	See §3.4
Trametinib	Not evaluable — the gap most worth closing	Highly aneuploid RPE1 clones activate RAF/MEK/ERK and are more sensitive to MEK inhibition, reproduced in human cancer lines (PMID 39251587). But we could recover only <i>relative</i> IC50s from it (its Fig. 5E), and the 0.45 nM it uses to sensitise clones to etoposide is a sub-lethal dose, not an IC50 — so there is no denominator. Cmax at the approved 2 mg/day is 22.2 ng/mL = 36.1 nM (Mekinist SmPC §5.2). It is approved from age 1 — in combination with dabrafenib, for BRAF V600E low-grade glioma — with an oral paediatric formulation, so it clears filter 1; filter 2 passes on the plausibility standard applied to every candidate (aneuploid RPE1 clones are MEK-inhibitor-sensitive); filter 5 is untested for it. It was pre-specified in our DepMap test as a mechanism-unrelated

Drug	Filter 3 verdict	Why
		comparator and came back null (\$3.5); the later literature review made that label inappropriate, so we keep the original analysis and report the revised rationale. The first drug we would put through filter 3 once an absolute aneuploid IC50 is published

Substituting an IC50 from an unstratified tumour line would manufacture a ratio without meaning. We leave those cells empty rather than fill them.

3.4 *The antagonism nobody should walk into*

A team building a combination might reach for an mTOR inhibitor alongside a proteasome inhibitor: both touch protein homeostasis, so synergy feels intuitive. The evidence points the other way.

Aneuploid and CIN cells are vulnerable to proteasome inhibition **because** they carry a translational burden. Reducing protein synthesis relieves that burden — and the ovarian carcinoma work shows exactly that mechanism operating as *resistance*. Ovarian cancer cells kept their sensitivity to proteasome inhibitors in adherent culture but “**acquired resistance as spheroids ... by downregulating protein synthesis via mTORC1 suppression**”; conversely, PTEN loss drove constitutive mTORC1 activity and **resensitized** the spheroids to proteasome inhibition, in vitro and in vivo (PMID 31530568).

Two boundaries on that evidence, since we are the ones importing it. It is ovarian carcinoma spheroids, not constitutional aneuploidy. And the paper did **not** test rapalogues — the antagonism below is our inference from its mechanism, not its result.

The objection this raises against our own candidate, stated before a reviewer does. The same paper notes that this protective mechanism “may explain the overall poor response to bortezomib in clinical trials of patients with advanced ovarian cancer.” Bortezomib has therefore *failed* in at least one solid-tumour setting, and any honest reading has to carry that. Our answer is that the proposed failure

mode is context-specific — a spheroid/ascites configuration suppressing translation — and that the same paper identifies a responsive subpopulation defined by mTORC1 status. It is a reason to measure the response in patient cells before believing it, which is what §6 does; it is not a reason the mechanism is wrong.

The premise of that prediction has independent 2025 support.

Work in yeast shows that chromosome amplification drives premature aging through defects in ribosome quality control, that aneuploids carry an increased load of protein aggregates and signs of ubiquitin dysregulation, and — the load-bearing point for us — that the *increased translational load* is what accelerates the decline (PMID 41248159). That is a different organism and an aging phenotype rather than a drug response, so we cite it for the premise only: reducing translation relieves a burden these cells depend on carrying. It is also the reason the same axis appears twice in this report with opposite signs, since MVA is itself a progeroid syndrome (PMID 31738183).


Prediction, stated so that it can be falsified: everolimus or sirolimus would antagonise bortezomib in aneuploid cells rather than potentiate it. Any screen should test that combination explicitly and expect a negative interaction.

A tension we flag rather than hide. The same BubR1 mouse work reports that predisposition to sarcopenia correlates with **mTORC1 hyperactivity** (PMID 31738183) — and this child has skeletal muscle atrophy. Read the qualifier with it: that correlation is reported in mice carrying **monoallelic** *BubR1* mutations, and this child is presumed biallelic, so carrying the link across is our inference and not the paper's. So an independent, muscle-directed rationale for mTOR inhibition in MVA exists, and it points the opposite way from our antagonism prediction. The two are not strictly in conflict: one concerns proteotoxic killing of aneuploid cells, the other a progeroid muscle phenotype. But anyone testing the combination should know both exist, and should not read a muscle-directed mTOR rationale as support for pairing it with a proteasome inhibitor.

3.5 We tested our own candidate on public data, and pre-registered it first

Two reviewers made the same objection to §3.2 independently: the human signal is **multiple myeloma**, and this child's tumour was an **embryonal rhabdomyosarcoma**. We had no number for solid tumours, so we produced one.

We fixed the hypotheses before we looked. The drug list, the direction of every predicted effect, the statistic, the multiplicity correction and the stopping rules were written to `depmap/PREREGISTRO.md` and committed to the repository **before the analysis was run**; the commit history records that order (it documents the record, not our conduct — `VERIFY.md` shows how to check it, and the deviations are listed there too). When one hypothesis turned out to be untestable mid-analysis, we recorded that as a dated amendment rather than editing the original.



Pre-registered test: our candidate among 6,790 compounds, and the lineage-restricted dependency on its target

Method. An arm-level aneuploidy score for 2,420 DepMap 24Q4 lines from absolute copy-number segments, validated against lines whose biology is independently known (chromosomally stable MSI lines average 5.0 altered arms; CIN lines 16.5). Then Spearman correlations against PRISM Repurposing drug sensitivity and DepMap CRISPR gene effect.


Result 1 — drugs. Across 424–444 evaluable solid-tumour lines, depending on the inhibitor, **no proteasome inhibitor reached FDR < 0.05**. Bortezomib ($\rho = -0.075$) and carfilzomib ($\rho = -0.060$) fall in the 4th and 8th percentile of all 6,790 compounds *by direction* — suggestive

and nothing more; ixazomib runs the other way ($\rho = +0.028$). Metformin is null, as predicted. Paclitaxel, our pre-declared cytotoxic control, runs opposite ($\rho = +0.096$), which argues against — without excluding — a “sick cells die more” confound. The pipeline is not blind: across all 6,790 compounds 1.1% reach $p < 0.001$ against 0.1% expected by chance.

Result 2 — the target itself. Genetic dependency on PSMB5, the subunit bortezomib binds, has a modest overall association with aneuploidy ($\rho = -0.198$, $q = 5.2 \times 10^{-8}$, $n = 946$) — but split by lineage it is **significant in haematological lines ($\rho = -0.261$, $q = 0.042$, $n = 113$) and not in solid lines ($\rho = -0.077$, $q = 0.118$, $n = 833$)**. The 20S core as a set is indistinguishable from a random-gene background in solid lines ($z = -0.51$), and the pre-declared negative control PSMD9 is null throughout. In a main-effects model adjusting for ploidy and lineage, aneuploidy still predicts PSMB5 dependency ($\beta = -0.0079$ per altered arm, $p = 0.0098$) while **ploidy was not significant** ($p = 0.875$; $p = 0.82$ in the interaction model) — the association persists after adjustment for measured ploidy in this model, which leaves residual confounding and whole-genome-doubling history unresolved. In the model with the interaction, the formal test of our question, the **aneuploidy $\times$ lineage term, is significant ($\beta = -0.017$, 95% CI -0.031 to -0.002 , $p = 0.026$)**, and the solid-stratum slope on its own is -0.006 (95% CI -0.012 to $+0.0002$, $p = 0.056$). Solid lines are the *larger* stratum, so a smaller sample is not the explanation, although predictor range, heterogeneity and measurement error remain possible contributors; the interaction supports a weaker association in solid lines under the fitted model, not the absence of a clinically meaningful effect. (These three numbers were first computed ad hoc; they now come from `depmap/05_seguimiento.py`, which reproduces all of them.)

Follow-up, pre-registered after seeing the result and run once (7 Sep; `depmap/PREREGISTRO_SEGUIMIENTO.md`, `depmap/RESULTADOS_SEGUIMIENTO.md`). A reader asked for the uncertainty, the structure under the pooled “solid” stratum, and the representation of this child’s tumour type. All five pre-registered predictions (P-S1 to P-S5) held:

DepMap follow-up (pre-registered): the aneuploidy-proteasome association, with its uncertainty



Follow-up: effect sizes with 95% bootstrap intervals, and the solid stratum split by lineage

- **Effect sizes.** Bortezomib $\rho = -0.075$ (95% bootstrap CI -0.168 to $+0.017$); every proteasome-inhibitor interval in the drug screen includes zero. PSMB5 dependency: haematological $\rho = -0.261$ (-0.422 to -0.084), solid $\rho = -0.077$ (-0.149 to -0.009) — that solid interval excludes zero narrowly while the pre-registered FDR test did not pass ($q = 0.12$); the pre-registered verdict stands and both are reported. Per ten additional altered arms, PSMB5 gene effect moves by -0.06 in solid lines and -0.23 in haematological lines.
- **Not driven by any single lineage.** Dropping each of 26 solid lineages in turn never flips the sign of the solid-stratum bortezomib or PSMB5 correlation, and the interaction's 95% CI never crosses zero (p from 0.013 to 0.037). Inside the solid stratum, all 27 within-lineage intervals (17 lineages with $n \geq 20$ for CRISPR, 10 for PRISM) overlap zero; no lineage is singled out.
- **The two assays barely agree with each other.** Across the 365 solid models with both measurements, bortezomib sensitivity and PSMB5 dependency correlate at $\rho = -0.070$ (-0.176 to $+0.033$). They are different modalities on overlapping models, not independent replications, and we no longer describe them as “two independent measurements”.
- **Score definition.** Under a fraction-of-arms score the estimates are similar with unchanged signs (bortezomib -0.077 , carfilzomib -0.054 , ixazomib $+0.032$); under a ploidy-residualised score the bortezomib and carfilzomib correlations shrink from $-0.075/-0.060$

to $-0.028/-0.028$ (signs unchanged), so part of the small drug signal travels with ploidy-correlated aneuploidy. The interaction is unaffected ($p = 0.024-0.029$).

- **Rhabdomyosarcoma is essentially absent from the drug screen — and what is there does not help us.** Seventeen RMS lines carry an aneuploidy score; **three** have a PRISM bortezomib value and they are *less* killed at the single $2.5\ \mu\text{M}$ dose than 98%, 91% and 89% of the 444 solid lines; **ten** have a CRISPR PSMB5 value, spread between the 24th and 84th percentile. No within-RMS association test was performed on three or ten lines; the percentiles are descriptive. This is a coverage limitation of the public data for exactly the tumour type this child had, and it is consistent with the limited in-vivo activity in solid-tumour xenografts reported by the Pediatric Preclinical Testing Program (§4).

The sentence the pre-registration commits us to, given these outcomes: **the attenuation in solid lineages is not an artefact of any single lineage and its interval excludes zero; the solid-stratum association itself remains compatible with zero.** Three results coexist behind that last clause and we keep them apart: the marginal Spearman interval (-0.149 to -0.009) excludes zero narrowly, the pre-registered FDR test did not pass ($q = 0.12$), and the adjusted regression slope's interval (-0.012 to $+0.0002$) includes zero — “compatible with zero” refers to the adjusted slope.

What it costs us, stated plainly. The CRISPR data support a stratum difference, and the PRISM data supply a separate, non-significant solid-line drug association: the link our bortezomib argument rests on is detectable where myeloma sits and attenuated in solid tumours. That attenuation is a statement about DepMap lines, not a claim that no solid-tumour evidence exists: Ippolito's own paper reports associations in pancreatic and paediatric PDX models (their Supplementary Fig. 80–r), which in turn say nothing about constitutional MVA. The §4 “prepared answer for a second tumour” therefore rests on evidence **we could not extend to the relevant lineage.** We do not withdraw the proposal — the direction survives for two of three proteasome inhibitors, the mechanism is unchanged, and a null across

heterogeneous cancer lines is not a null in constitutional MVA. But the claim is weaker than it was before we ran this, and the experiment in §6 matters more, not less.

The obvious objection, raised and answered here rather than left for a reviewer: Ippolito ran their own bortezomib screen and it was positive. In their Fig. 6p, 387 cancer lines were treated with bortezomib with or without a low dose (250 nM) of reversine; reversine had only a mild effect on proliferation but significantly sensitised the lines to proteasome inhibition ($p < 0.0001$). That is a positive result on the same drug, and our PRISM analysis is null in solid lines. **The two do not point in opposite directions, and we say so precisely because the tempting summary — “we contradicted them” — would contradict our own results table.** Our bortezomib correlation carries the *same* sign as theirs ($\rho = -0.075$, sensitivity increasing with aneuploidy) and simply fails significance ($p = 0.11$, `depmap/RESULTADOS.md`). The designs differ in what they manipulate: Ippolito **induce** aneuploidy in each line and read the change in dose response within that line, which controls stable between-line differences, while attribution specifically to induced aneuploidy still depends on the perturbation design; we **correlate** pre-existing aneuploidy across heterogeneous lines against single-dose viability, which does not. The two designs estimate different quantities — a within-line perturbation versus a cross-line correlation — so their p -values do not rank them, and our null is not evidence against their result. What our CRISPR analysis adds, and theirs does not address, is a **lineage split in PSMB5 dependency** (PRISM cannot estimate the corresponding haematological drug-response association): the association is carried by haematological lines and attenuated in solid ones (interaction $p = 0.026$), and this child’s tumour was solid.

What we did not do. PRISM primary is a single-dose viability screen, not an EC50, and not Ippolito’s isogenic comparison. This is not a replication attempt of Ippolito and we do not present it as one. Full method, numbers and limitations: `depmap/RESULTADOS.md`.

4. Filter 4: safety in *this* child — where the candidate stops

Bortezomib is not excluded by the total-plasma peak comparison of filter 3. It does not clear this patient today.

Concern	Data	Relevance here
Peripheral neuropathy	25/140 (18%) of patients in AALL07P1 — relapsed lymphoid malignancy, bortezomib with multi-agent reinduction chemotherapy including vincristine; safety population aged 1.0–26.8 years — motor 11/140 (8%) (FDA BPCA Clinical Review, NDA 021602, 2015): observed regimen-associated events, not this child’s incremental risk from bortezomib alone. In adults, grade ≥ 2 in 39% and grade ≥ 3 in 15% by IV route (FDA label, NDA 206927 s002, 2022 — a different source from the PK figures above)	This child has skeletal muscle atrophy (HP:0003202). Superimposing a motor neuropathy on reduced neuromuscular reserve is the decisive objection
Thrombocytopenia, neutropenia, infection	FDA label	Prior chemotherapy for ERMS; marrow reserve unknown
Gastrointestinal toxicity, hypotension	FDA label	Failure to thrive (HP:0001508) — anything suppressing intake is costly
Renal	PK not clinically altered by renal impairment per label	Nephrocalcinosis (HP:0000121) still requires renal function, hydration and electrolytes to be established first

Bortezomib is therefore not proposed as prophylaxis or as chronic therapy for this child.

But filter 4 is patient-specific, not universal — and that is the point.

Read what follows against §3.5: our own pre-registered test could not extend the proteasome dependency to solid-tumour lineages, and an embryonal rhabdomyosarcoma is a solid tumour. The argument below is weaker than the mechanism alone suggests; the fibroblast study of §6 examines constitutional-cell response, and tumour-directed

investigation requires a separate tumour/normal comparison. For a patient with MVA **and an active malignancy**, where cytotoxic therapy is on the table regardless and the comparator is conventional chemotherapy rather than nothing, the comparator changes — whether adding bortezomib then improves the benefit–risk balance is unestablished, and a judgement for the treating team, not a result of this report. Even then the route matters and we name it: this would be off-label use decided by a molecular tumour board, or a formal n-of-1 protocol with its own ethics approval and its own stopping rules — not a recommendation that follows from this report. In that setting a proteasome inhibitor brings documented pharmacology and toxicity, not established benefit in this combination or disease: it is an approved agent that targets the best-characterised downstream consequence of aneuploidy **as established in cancer cells** — proteotoxic dependency — supported by a human clinical signal tying response to aneuploidy burden. Section 6 measures differential drug response and constitutional-cell vulnerability — mechanistic attribution and tumour selectivity require separate experiments — and our own DepMap analysis found the link attenuated in the solid lineage this child’s tumour came from. It does not correct the checkpoint defect, and we do not claim it does.

This child survived an embryonal rhabdomyosarcoma, and the risk is gene-specific: “individuals with biallelic TRIP13 or BUB1B mutations have a high risk of embryonal tumors” (PMID 28553959), and an E-RMS arising in an MVA case with biallelic BUB1B has been characterised directly (PMID 16182441). **The realistic clinical question is not “what do we give him today” but “what do we reach for if there is a second tumour” — and that question deserves an answer prepared in advance, with its evidence and its threshold, rather than improvised under pressure.**

What is already known about bortezomib in the tumour he had, and in children. We looked, because a tumour board would: favourable preclinical observations, limited in-vivo solid-tumour activity, and no demonstrated clinical benefit in the cited trials. Embryonal and alveolar rhabdomyosarcoma cell lines are killed at **13–26 nM**, concentrations that spared primary human myoblasts, and bortezomib reduced the growth of an RMS xenograft (PMID 18342500). Across the Pediatric Preclinical Testing Program panel the median in-

vitro IC₅₀ was 23 nM, **but in-vivo activity against the solid-tumour xenografts was limited** — one line reached intermediate activity, the rest low — while leukaemia xenografts responded (PMID 17420992). In children, the Children's Oncology Group phase I study ADVL0015 set the recommended dose at 1.2 mg/m² twice weekly for two of every three weeks, with dose-limiting thrombocytopenia at 1.6 mg/m² and **no objective responses** among 15 children with refractory solid tumours (PMID 15570082); the vorinostat combination ADVL0916 likewise saw no objective responses, and a grade 2 sensory neuropathy that progressed to grade 4 was dose-limiting (PMID 22887890). In adults, a phase II study in sarcomas closed its osteosarcoma/Ewing/rhabdomyosarcoma arm for low accrual and reported minimal single-agent activity in soft-tissue sarcoma, with painful neuropathy prominent (PMID 15739208). The reading for a tumour board is therefore narrow: the in-vitro sensitivity of RMS at concentrations below the paediatric C_{max} is real, the paediatric dose and toxicity profile are established, and in the paediatric trials we found single-agent bortezomib produced no objective responses in solid tumours (ADVL0015: 15 enrolled, 11 assessable, no objective responses). If it is ever reached for, it would be as a component of a combination — which is what the sarcoma trial itself recommended — and the aneuploidy argument of this report is a reason to ask whether MVA-driven tumours are the exception, not evidence that they are.


There is a second, immediate implication. Spindle poisons — vincristine, taxanes — depend on a competent mitotic checkpoint. With reduced BubR1 that arrest is less robust; a case of rhabdomyosarcoma with PCS/MVA treated with **reduced-intensity chemotherapy** is on record (PMID 31184400 — a single case, not a guideline and not a safety estimate; PubMed carries no abstract for it, so we cite only what its title states and do not name the regimen). This is a mechanistic caution for the treating oncologist to weigh against probability of cure — the report supplies no clinical evidence that would make it a contraindication of any grade — and never a reason to withhold curative therapy. If a mitotic-poison-sparing regimen is ever needed, there is **no validated replacement agent**: a proteasome inhibitor would be a bench hypothesis to test first, not a substitution anyone could make at the bedside.

We state the boundary explicitly: this is a hypothesis for a tumour board, not a prescription.

The tumour-stage question, specified. *Research setting:* an active MVA-associated tumour. *Comparison:* tumour-derived material against constitutional normal-cell comparators under matched exposure conditions, with the tumour's standard regimen as the treatment comparator. *Question:* does adding bortezomib produce an incremental effect on the tumour material with a reproducible tumour/normal separation? *Exposure condition:* the tested schedule must be justified against clinical exposure and pharmacodynamic evidence (duration of proteasome inhibition), not against the 312 nM peak alone. *Decision (provisional planning thresholds — design choices, not findings — to be frozen before any material is obtained):* compare vehicle, bortezomib, the standard regimen, and standard regimen plus bortezomib, with viable-cell recovery at a fixed post-washout time as the primary endpoint; require at least 20 percentage points of incremental tumour killing from adding bortezomib, at least 10 points of tumour-normal separation, and no more than 10 points of incremental killing in either normal comparator, with Δ_T the incremental tumour killing (standard regimen plus bortezomib minus standard regimen, on vehicle-normalised viable-cell recovery), Δ_N the same quantity in each normal comparator and $D = \Delta_T - \Delta_N$: advance only if the simultaneous lower confidence bound of Δ_T exceeds 20 points, both lower bounds of D exceed 10 points and both upper bounds of Δ_N stay below 10 points, with culture-level resampling and replication in two independently initiated experiments; a normal-comparator lower bound above 10 points is *harm*, a tumour upper bound below 20 points is *insufficient benefit*, anything else is *inconclusive*. The operating characteristics of this rule have not been simulated; failure of the constitutional-fibroblast hypothesis (§6) does not decide this question. *Prioritisation:* there is no comparative basis today for ranking this work above trametinib or another unresolved candidate, and we say so rather than assert one.

5. Why any endpoint must be measured per cell

Any trial of the above needs a readout for **aneuploidy burden**. We asked whether the data this hackathon provides could supply one, and computed the answer instead of asserting it.



Detection limit of bulk WGS for mosaic aneuploidy

Figure 1 shows the **first-order** analysis only: mean $|BAF - 0.5|$ against a fixed 0.005 noise floor, crossing at $\approx 14\%$ of cells. **The later calibrated limit of 2.07% is not on this plot** — it comes from a different statistic (Ψ over 125-site windows against a Monte Carlo null at $\kappa = 2$) and cannot be read off this curve. Neither figure is a measurement of this child's genome; see “What we do not claim about this child” below.

We are not proposing a new assay. The correct endpoints exist and are standard: metaphase karyotyping with premature chromatid separation scoring is the diagnostic criterion for MVA; the micronucleus assay is an international standard (OECD TG 487); and single-cell DNA sequencing has been used to assay chromosome copy number genome-wide since Knouse et al. 2014 (PMID 25197050). We cite that work for the method, not for our conclusion — its own finding is that aneuploidy is *rarer* in liver and brain than previously reported.

What we contribute is the number.

The analysis used the **VCF only**, so no GC-LOESS normalisation was applied. **That was our choice of input, not a limit of the challenge:** the organisers do distribute the raw reads (84.7 GB of FASTQ), and the statistic that ends up carrying the argument does not need them. An earlier version of this paragraph blamed the absence of alignments for the absence of GC normalisation, and it was wrong on both halves. Four analyses were run, of which three are refinements of the same B-allele-frequency statistic and one (depth-based dosage) is orthogonal;

they are not four independent tests. In a trisomy, heterozygous sites split symmetrically, so the *mean* BAF does not move and the correct statistic is $\text{mean } |\text{BAF} - 0.5|$.

The measured standard deviation of $\text{mean } |\text{BAF} - 0.5|$ across the 22 autosomes at fixed depth (DP 40–48, 857,435 sites) is **0.0035**, so the 0.005 noise floor we use — arising from GC content, mappability and segmental duplications — is a ≈ 1.4 -SD criterion. Simulating at DP 44 against it gives a **first-order detection limit of a clonal whole-chromosome trisomy present in $\sim 14\%$ of cells**; a conventional 2-SD criterion moves it to 16% and 3-SD to 20%. Every stricter choice makes the assay look worse, so the conclusion does not rest on the threshold.

A later analysis improved that limit, and we report both because they answer slightly different questions. Replacing the fixed 0.005 floor with a calibrated null — Ψ over 125-site windows, the measured Beagle switch process, the empirical segment-size distribution, and 40,000 Monte Carlo replicates per cell — gives a limit of **2.07% of cells at an overdispersion $\kappa = 2$** , and 1.46% at $\kappa = 1$, 2.95% at $\kappa = 4$ (`mosaico/RESULTADOS.md` §3; every value in `mosaico/ventana.json` at the chosen window $W = 125$). Read that number with its three conditions attached: **it is a simulated limit of detection under a fitted null, not a measured sensitivity, and it is conditional on κ** , which we estimate from these data at ≈ 2.3 rather than knowing independently. Under it, the 1–2% per chromosome that variegation implies sits **at the edge of what is achievable rather than far below it** — a better limit than the 14% figure, but one that leaves the variegated signal at the edge of detection rather than comfortably inside it.

Read sampling is not what sets that limit. The simulation subtracts the binomial baseline at DP 44, so the 0.005 threshold a signal must clear is the *measured dispersion between chromosomes* — systematic library bias, not counting noise. The depth-based dosage test points the same way by a separate route: Poisson error when averaging a whole chromosome at $44\times$ is $\sim 0.03\%$, far below the observed deviations. These are two different statistics and we do not conflate them; both imply that sequencing deeper would barely move the limit.

As an **illustrative bound, not a measurement of this child**: if 30% of cells were aneuploid — the order of magnitude of the diagnostic criterion, not a figure from this patient — spread across 22 autosomes, each chromosome would be affected in **1–2% of cells**, and gains and losses of the same chromosome cancel in bulk. Against the first-order 14–16% floor that is **7–16× below** it; against the calibrated 2.07% limit it is **at the boundary**. The conclusion survives either way, but it does not survive *because of* the limit, and the argument that does not depend on any limit is this one:

A perfectly balanced variegated mixture is not identifiable from bulk allele fractions. Model a fraction g of cells gaining homologue A, an equal fraction gaining B, and matched losses of each. Mean copy number is exactly 2 and mean allele fraction exactly 0.5, **by symmetry**. Simulated under identical conditions, detection stays at **0.1–0.2% — the false-positive rate — from $f = 2\%$ up to $f = 40\%$** (`mosaico/RESULTADOS.md` §4). That is not low power. Two biologically distinct cell populations produce the same observable distribution, so no depth and no statistic separates them.

That proof is about the mixture we modelled, not about every MVA genome. Real variegation need not be perfectly balanced between homologues or between gains and losses, and a sufficiently asymmetric pattern would leave a signal. What we have shown is that the architecture the disease name describes contains a regime that bulk DNA cannot reach even in principle — enough to disqualify the endpoint, not enough to say “bulk WGS can never see MVA”.

And **bulk averaging conceals balanced variegation by construction**, at any depth, and cannot recover the distribution of chromosome states across cells, although asymmetric mixtures may remain detectable. The cytogenetic hallmark of MVA, premature chromatid separation, leaves no trace in DNA sequence at all.

Several candidate signals we did observe coincided with technical features — coverage, segmental duplications (22q11, centromeres), or partially GC content — partially, because our measured GC–coverage correlation was $r = 0.43$ from a sparse sample, and chr21, the largest deviation, is GC-poor. That is an insufficient measurement rather than an established explanation, and we report it as such.

What we do not claim about this child, and why

We ran the calibrated statistic on the proband and **we withdrew the result**. It is not reported here as a negative finding, and the improved limit must not be read as one either. Three reasons, all from our own files:

- **We calibrated one test and applied another.** The pre-registered Ψ threshold is exceeded on **22 of 22 autosomes** at every κ setting. That is not 22 mosaic events; it is a null model that does not describe these data. The between-chromosome robust z-score we actually reported was chosen after seeing the data and was never calibrated (`mosaico/RESULTADOS.md` §5).
- **The residual signal is uncontrolled.** What remains is a global overdispersion of $\approx 2.3\times$ the binomial expectation, shared by all 22 autosomes and correlated at the ~ 200 kb scale. From the VCF alone we cannot say whether it is technical or biological.
- **The external control was pre-registered, it ran, and it came back not conclusive.** Ten unrelated 1000 Genomes individuals on two fixed windows (`mosaico/PREREGISTRO.md` amendment 4, committed `6405d9e`; result `ce94b89`). The proband fell **below** the control envelope on both statistics and both chromosomes — an outcome the amendment had not listed, so the pre-registered verdict is *not conclusive*. The cause was declared in advance with its direction: the controls carry no GQ filter and were called jointly across 3,202 samples against a singleton call for the proband, which inflates their dispersion. A κ of 46 in a healthy individual is a property of the calling, not of biology. **The control did not fail to run; it failed to be a control** (`mosaico/CONTROL_RESULTADO.md`).

One narrow statement survives, and we make only it: against those ten genomes, window-scale overdispersion ranges from $\kappa \approx 6$ to 46, and the proband's is the **lowest in the comparison** — 2.74 on chr8 against a control range of 9.5–46.0, and 4.60 on chr17 against 6.1–15.0 (`mosaico/control_resultado.json`; these are the thinned window statistics the comparison used, not the $2.27\times$ chromosome-wide excess reported in `mosaico/RESULTADOS.md`) — so, with the caveat that the comparison itself was judged not conclusive, a κ of 2–5 is not large by the standards of these controls. **That is not a clinical negative.** We are not reporting that this child has no clonal lesion above 2% of cells; we

are reporting that our attempt to measure it did not produce an interpretable answer. Resolving it needs a control processed identically from reads, which we did not have.

Consequence: no proposal whose endpoint is “reduced aneuploidy burden” can be evaluated with these data. Any screen must measure per-cell mis-segregation.

6. The go/no-go experiment — an eight-week assay window

Scoped to be executable by one laboratory in one window, not as a grant application.

Question. Is there a reproducible difference in bortezomib response between the sampled patient and control cultures across nominal concentrations spanning the reported total-plasma peaks (without assuming equivalent clinical exposure), and can an increase in the micronucleus endpoint exceeding the provisional margin be excluded among the survivors?

What this experiment can and cannot decide, stated first. Dermal fibroblasts from this child are *constitutional* MVA cells, not tumour cells. A sensitivity difference between them and unaffected controls therefore measures a response difference between the sampled lines — attribution to a transferred proteotoxic dependency requires mechanistic follow-up — and it does **not** measure tumour selectivity, and on its own it is as compatible with systemic toxicity as with therapeutic promise. That is why the go criterion below is *not* “patient cells are more sensitive”. A negative therefore **fails to support transfer without closing the tumour-board question**, which would need aneuploidy-high tumour-derived material. We say this now rather than after the data.

The quantity that governs this design is the fraction of cells carrying *any* aneuploidy, not the fraction carrying a *given* chromosome. A cell is under proteotoxic load whichever chromosome it gained, so the ~30% of the diagnostic criterion is the right denominator here, while the 1–2% per chromosome is the right one for §5. A pre-registered power calculation ([potencia/PREREGISTRO.md](#), committed before the run) gives **power 0.897** at 30% of cells with an

assumed $4\times$ EC50 selectivity, $n = 3$ and $\alpha = 0.05$, against **power 0.063** under the per-chromosome reading (both from the original two-parameter model). An earlier version of this section used the per-chromosome figure and concluded a negative would be ambiguous by design; that hedge rested on the wrong denominator and is withdrawn as a statement about statistical power — although the four-parameter follow-up below restores a version of it for a different reason: at $n = 3$ the design is underpowered under the declared model, so the remedy is n , not the denominator.

Read 0.897 as an idealised scenario, not as this experiment's power. It holds *for an assumed $4\times$ selectivity, in a $\sim 30\%$ aneuploid culture, at 10% well-level CV, and none of those three has been measured in these cells.* The simulation is also weaker than the design we pre-registered: it fits a two-parameter Hill curve where the pre-registration promised four, and its noise model has no biological-replicate term. Both push the number upward, so 0.897 is optimistic **within its own model** — it is neither a validated power estimate for patient cells nor a mathematical ceiling on what a better-specified version of the experiment could achieve (`potencia/RESULTADOS.md` §5b).

We then ran the model we had promised, pre-registered as a follow-up (`potencia/PREREGISTRO_4P.md`, `RESULTADOS_4P.md`, 8 Sep), **and it retires the small- n branches.** Four free Hill parameters, a shared log-normal multiplier on each replicate culture's EC50 (CV_{bio} 0.20), a plate-level offset (CV 0.05), per-well noise (CV 0.10), the current 0–1,000 nM concentration grid, a Welch t on log₁₀ EC50, and fitting failures counted as non-rejections. At the central scenario ($f = 0.30$, $4\times$ selectivity, $n = 3$ per arm) **power is 0.167**. The two-parameter fit on the same simulated data gives 0.234, so the four-parameter fit itself costs 0.067 and the rest of the drop from 0.897 comes from what the re-run also changed — the biological and plate variability, together with the 0–1,000 nM grid and the Welch test; fitting failures were 0.0% in every scenario. Under that scenario the simulation first crosses 0.80 at **$n = 12$ independent cultures per arm at $f = 0.30$, $n = 23$ at $f = 0.20$** , and not within 24 at $f = 0.10$ — a conditional planning figure (600 simulated experiments per n , testing rejection of equal EC50s, not the complete advancement rule), not a validated sample size; at $n = 3$ it is 0.05 at $f = 0.10$ and 0.43 at $f = 0.50$, and CV_{bio} 0.10–0.30 moves it between 0.23

and 0.12. Two of the four pre-registered predictions were falsified — we expected fitting failures of at least 5% and expected the fit, not the biology, to be the larger cost — and both are reported. **What n counts:** independent biological replicate cultures per arm, from separate passages on separate plates, each yielding one fitted EC50; technical wells are not n . The simulation assumes independent culture-level errors; real cultures remain nested within one donor and share lineage, passage and batch effects, so separate passages and plates do not establish independence — passage and plate will be recorded and modelled, and inference is limited to the sampled lines. **Consequence, per the pre-registered mapping:** the $n = 3$ branch and the “ $n = 14$ at $f = 0.10$ ” branch of earlier versions are retired from this report; 0.897 remains only as the optimistic two-parameter reference; and a design that needs twelve independent cultures per arm from one biopsy is heavier than the eight-week assay window implies, which the timeline paragraph below now carries.

A second analysis, pre-registered and run on 1 September — the same day as the original power calculation (commit 45c4ed7) — gives us our own reason to doubt the 4×. Proteotoxic load is a dose, not a switch: for an illustrative low-burden scenario we model an MVA cell with **about one** altered chromosome, two arms (the treated culture’s distribution is unknown), against the eighteen-arm point at which we evaluated the dose response — close to the mean of the solid DepMap stratum (`potencia/dosis.json`, `arms_mean = 18.6`), which is its middle and not its top. We have not counted arms in the specific lines that gave the < 40 nM EC50; the eighteen is the high end of the range over which we measured the dose response, not a measurement of those lines. In DepMap, in the one lineage where the aneuploidy–proteasome dependency is significant at all, PSMB5 dependency moves **0.09 residual SD at two arms against 0.84 at eighteen** (`potencia/RESULTADOS_2_DOSIS.md`, pre-registered as 45c4ed7).

State precisely what that does and does not do, because we got this wrong once. The pre-registration fixed its magnitude rule on the **solid** stratum. There the slope’s 95% interval covers zero ($-0.0118 \dots +0.0006$, $p = 0.076$), so the falsification clause applies and the exercise is **uninformative** for the stratum the rule was written about. The 0.093 ratio is **haematological**, where no decision rule was pre-registered, and CRISPR gene effect does not convert into an EC50 ratio in either

stratum. So the 4× central value is **unsupported for these cells and equally unrefuted**. An earlier version of this section reported the rule as having fired; it did not, and the correction is dated 2026-09-06 in RESULTADOS_2_DOSIS.md §3. At an assumed 2× selectivity the original two-parameter model gave **0.387**; the declared four-parameter model gives **0.068** at $n = 3$. The honest summary is that **we do not know which regime these cells are in, and the experiment is what would tell us**.

It still buys a protocol change, not just a corrected sentence. Power falls steeply with the aneuploid fraction — under the declared model, at $n = 3$, 0.096 at $f = 0.20$ and 0.050 at $f = 0.10$ (0.589 and 0.196 under the original two-parameter model) — and **nobody has counted metaphases in these fibroblasts** — culture selection may change the fraction in either direction, so it is measured at the passage used for treatment. So the fraction becomes a **gating measurement** rather than an assumption, taken on the culture that will actually be treated, at the passage that will be treated. The endpoint is a **count of cells carrying at least one numerical chromosome abnormality**, recorded together with **how many abnormal chromosomes each abnormal cell carries**, because the second analysis shows the per-cell burden matters as much as the fraction.

PCS scoring is a companion, not this measurement. Premature chromatid separation is a centromere-cohesion phenotype and the diagnostic criterion for MVA; the aneuploid fraction is a count of abnormal chromosome *numbers*. The same metaphase slides yield both, at no extra cost, but they are different endpoints and an earlier version of this section conflated them.

Sample size, under the declared model: **about 12 independent cultures per arm if the measured fraction is ≈ 0.30 , 23 at ≈ 0.20 , and below ≈ 0.20 do not run the bulk viability assay** — go to a per-cell endpoint, because the required n grows past what a biopsy-derived culture can supply (RESULTADOS_4P.md; the earlier $n = 3$ / $n = 14$ branches came from the optimistic two-parameter model and are retired). Every branch is computed at the same assumed 4× selectivity and CV_{bio} 0.20, so **measuring the fraction does not validate the**

branch it selects: it removes one unmeasured input of four. Treat it as planning, not as a validated operating rule. Details and every sweep in `potencia/RESULTADOS.md` and `RESULTADOS_4P.md`.

Consent, assent and ethics approval come before week 1, and they are not ours to give. Taking a skin biopsy from a child, establishing a fibroblast line from it and retaining that line require informed parental consent, the child's assent at an age-appropriate level, and approval by the research ethics committee or institutional review board of whichever centre runs this. We are an analysis team with no access to the patient, no relationship with the family and no institutional sponsor for this work; nothing below is authorised by us, and neither the approval nor how long it takes is within our control. The same applies to the parental sampling in `clinico/PARENTAL_SEGREGATION_ORDER.md`, which sets out its own consent requirements.

Week	Step	Readout
1–3	Patient dermal fibroblasts from skin biopsy, plus two matched controls, established, with the aneuploid fraction f measured, and then expanded to the replication the measured f requires (about 12 cultures per arm at $f \approx 0.30$, 23 at $f \approx 0.20$) across recorded passages and plates (dependence modelled, not assumed away) — a pre-assay phase of 4–8 weeks that precedes week 1 and sits outside the eight-week window. Gate: on the culture to be treated, count metaphases scoring (a) the proportion of cells with ≥ 1 numerical chromosome abnormality, (b) the number of abnormal chromosomes per abnormal cell, and (c) PCS, reported separately. (a) is the input the power model needs and it is unmeasured; (c) is the MVA diagnostic criterion and is not a substitute for it (<code>potencia/RESULTADOS.md</code> §3). Patient cells are required: RPE1 with reversine-induced aneuploidy reproduces Ippolito's own system and cannot test whether the dependency transfers to BUB1B-deficient constitutional MVA — we list it as a positive control, not a substitute. Without patient cells the go/no-go question is not answerable	Growth; baseline karyotype
2–4	Allele fate: allele-specific RT-PCR $\pm$ NMD inhibitor; quantitative western blot	Allele-specific RNA and size-resolved immunoblots assess transcript decay and protein products; total N/C-

Week	Step	Readout
	with N- and C-terminal antibodies	terminal signal includes protein from the other allele and cannot by itself assign allele fate — the RT-PCR arm decides; the blot alone does not
4–6	Bortezomib dose–response, 0–1,000 nM with dense sampling below 100 nM, 72 h, patient versus control versus an aneuploidy-high positive control (reversine-treated RPE1)	EC50 and the ratio between them. The grid includes 312 nM, the provisional nominal-concentration ceiling used for this assay; this does not establish clinical exposure equivalence. Extended from 0–400 nM after the power calculation: at a selectivity of 10× the control's EC50 sits at the old ceiling and is estimated at the edge of the data, doubling the spread of the estimate (<code>potencia/RESULTADOS.md §4</code>)
4–6	Combination arm: bortezomib + everolimus	Tests the predicted antagonism
6–8	Micronucleus assay (OECD TG 487) on survivors of every condition, plus FISH for 3–5 chromosomes	Micronucleus frequency as a proxy for mis-segregation (filter 5); centromere/kinetochore labelling discriminates whole-chromosome loss from breakage (the discrimination OECD TG 487 calls for) and FISH adds chromosome-specific information for the chromosomes probed; proliferation is recorded because a change in division rate alone moves the count

Timeline, stated honestly in the heading it belongs in. Eight weeks is the *assay* programme, and it holds only once a patient fibroblast line exists — and twelve independent cultures per arm must be expanded from it first. Establishing $15\text{--}20 \times 10^6$ fibroblasts from a fresh biopsy typically takes 4–8 weeks on its own (`track2/DATOS_CANDIDATO.md §7`), so from a standing start the planning figure is **12–16 weeks**, not eight

— and the same source warns that growth can be slower in MVA1 fibroblasts, which would push it further. We use “eight-week experiment” as shorthand for the assay window and nowhere as a promise of an answer in eight.

Advancement criteria, fixed before the data. We state them in the direction that avoids rewarding toxicity:

1. **A reproducible response difference, requiring mechanistic follow-up before it is called a dependency:** an EC50 shift between patient and control fibroblasts whose 95% confidence interval excludes 1, **of at least the magnitude that a culture of the measured aneuploid fraction can produce.**

We had this criterion wrong and an adversarial review caught it. It previously demanded a shift “at least equal to the near-euploid-versus-highly-aneuploid contrast in Ippolito Fig. 6o”. That contrast is between two **homogeneous** populations. A patient fibroblast culture is a **mixture**, and under our own model a mixture cannot produce it: with 30% of cells aneuploid and a 4× selectivity *within* those cells, the **bulk** EC50 ratio is **1.50×**, not 4×. The old criterion was unreachable by construction — the experiment would have failed it even if the biology worked exactly as predicted.

The threshold must therefore be derived from the gating measurement, not imported from a homogeneous-population figure. At a 4× within-cell selectivity, the bulk ratio our model predicts is:

Measured aneuploid fraction f	Bulk EC50 ratio to expect	Max viability gap
0.10	1.14×	0.03
0.20	1.30×	0.07
0.30	1.50×	0.10
0.50	2.00×	0.16

Read the criterion off this table once f is measured, and treat a shift whose 95% interval lies below the corresponding row as a negative. A shift more than twice the row’s prediction is not a better result — it is a reason to suspect the mixture model or

systemic toxicity, which is what stop signal 4 below exists for. Every row assumes the same unmeasured $4\times$ within-cell selectivity, so this table is a pre-specified reading rule for the experiment, **not a validated threshold**. The primary threshold stays fixed; any table recalibrated from a selectivity estimate obtained in this experiment is exploratory and must be evaluated in a subsequent experiment.

2. **Effect at ≤ 312 nM.** That is the top of the *total plasma* C_{max} interval for the approved dose, used here as an order-of-magnitude ceiling for plausibility. It is not free drug, not intratumoral drug, and not a validated clinical threshold.
3. **Filter 5 read as a decision, not a slogan.** Micronucleus frequency among survivors (per 1,000 binucleated cells, OECD TG 487) in treated versus vehicle cultures of the same line. Let Δ be the treated-minus-vehicle difference and M the harm margin, set here at 50% of the vehicle rate (a design threshold, not a validated one; the vehicle rate and its uncertainty enter M through the same bootstrap). **Evidence below the margin requires the upper 95% confidence bound of $\Delta - M$ to lie below zero.** An interval crossing zero is *inconclusive* — read as not shown to pass filter 5, never as a pass; an interval wholly above zero is *harm exceeding the margin*. The proliferation index is recorded because a change in division rate alone moves the count, and micronuclei are classified with centromere/kinetochore labelling (the discrimination OECD TG 487 calls for), with FISH for 3–5 chromosomes adding chromosome-specific information for the chromosomes probed, so that an increase is attributed to mis-segregation only when the labelling says so. Micronuclei are scored only under prespecified proliferation, cytotoxicity and assay-validity criteria (OECD TG 487); unevaluable conditions are classified inconclusive, and the endpoint estimates frequency among evaluable dividing survivors — it cannot exclude injury in lost or non-dividing cells.
4. **A stop signal, not a go signal:** if patient fibroblasts are markedly more sensitive than both control donors, that raises concern about constitutional-cell toxicity in an MVA patient; the comparison cannot determine the mechanism (constitutional toxicity and selective killing of constitutional aneuploid cells are not mutually

exclusive, and RPE1 is not an isogenic comparator). RPE1 serves to confirm assay responsiveness; patient-cell toxicity is interpreted on its own, without requiring a particular ordering across cell types. That outcome stops the constitutional-cell programme too, and it is the outcome we would most easily have mistaken for success.

Two research questions, two decisions. Each result is classified as *advance*, *insufficient response* (a precise estimate below the required threshold), *demonstrated harm* or *inconclusive*; failing any of 1–3, or triggering 4, blocks advancement of the *constitutional-cell* programme (an inconclusive result permits no advancement claim and requires redesign, not rejection of the hypothesis): it measures drug response and potential normal-tissue vulnerability in the sampled lines. It cannot establish tumour benefit, and a negative result does not automatically stop tumour-specific investigation. Any tumour-directed proposal requires a separate comparison using relevant tumour material and normal-cell comparators, with its own decision rule.

Falsifying results, stated at the precision the design supports. If patient fibroblasts are no more sensitive than controls, the experiment **fails to support transfer of the dependency into these cells** — it does not falsify a dependency in aneuploidy-high tumour material, which is the setting §4 actually proposes and which this design does not sample. That distinction is the difference between “we did not see it here” and “it is not there”, and we hold ourselves to the first. If the micronucleus interval for $\Delta - M$ lies wholly above zero, that fails the constitutional-cell safety criterion and stops advancement on that evidence regardless of how well it kills; it does not by itself resolve tumour-selective benefit, and an interval crossing zero is inconclusive, not a pass.

The stage this design does not contain. Deciding tumour selectivity requires aneuploidy-high tumour-derived material with a near-euploid comparator, which no eight-week window starting from a skin biopsy can supply. The fibroblast arm decides the constitutional-cell hypothesis; it informs, but does not by itself decide, whether that second stage is built, and it is not a substitute for it.

This design can fail, which is the point of proposing it.

7. Actionable now, without any drug

The clearest clinical action supported by published guidance today is not a molecule.

There is a published consensus for MVA specifically, and it should be the anchor. The 2024 update from the AACR Childhood Cancer Predisposition Workshop covers mosaic variegated aneuploidy by name and states: *“This group recommended renal ultrasound surveillance every 3 months from birth until age 7, for all MVA conditions including those with an unknown genetic cause”,* together with *“regular clinical assessment including review of systems to identify signs of rhabdomyosarcoma and other malignancies”,* avoidance of radiation exposure, and HPV vaccination (PMID 39264246). The same source records that cancers in MVA1 are *“predominantly Wilms tumor and rhabdomyosarcoma, but MDS, AML and ALL have also been reported”*.

Provenance note, 2026-09-06. These three passages — the 3-monthly-to-age-7 interval, the tumour spectrum, and the statement in §9 about the 15 published MVA2 individuals — come from the **full text** of PMID 39264246, not its abstract. On 2026-09-06 we re-verified all three against the open-access full text (PMC11705613, archived verbatim in `evidencia/consenso_fulltext_2026-09-06.xml`): the interval is attributed there to the SIOP-Europe Host Genome Working Group and SIOP Renal Tumor Study Group, with the AACR group stating that it *supports* those recommendations; the MVA2 sentence reads *“at the time of writing, none of the 15 individuals with MVA2 described in literature developed cancer”,* which describes the published literature, not a followed cohort, and is not a statement of zero risk. Nothing in our recommendation is original to us here; a clinician should read the consensus itself before acting.

We had reached the 3-monthly-to-age-7 schedule independently, by extension from UK Wilms-risk recommendations that stop at age 5 (PMID 16857697). We report that convergence rather than our derivation: **the schedule below is consensus guidance, not our inference**, and a family should be given it as such.

Where the consensus leaves a gap, and it is the gap that matters for this child. The window closes at seven. A 2026 case report describes an MVA3 girl whose Wilms tumour appeared at three and whose orbital embryonal rhabdomyosarcoma appeared at **twelve** (PMID 42595739) — five years past the end of renal-ultrasound surveillance, at a site renal ultrasound does not image. Renal ultrasound is a kidney-directed test; rhabdomyosarcoma is not a kidney tumour, which is why the consensus pairs it with clinical assessment rather than with imaging. For a child who has *already had* an embryonal rhabdomyosarcoma, we therefore suggest an examination frequency and a review point at age 10; the continued clinical assessment itself is already supported by the consensus, which sets no end date for it — the frequency and the review point are ours, not guidance.

Surveillance windows must be anchored to the child's current age, which is not in the dataset and must be confirmed before use. If the age is unknown, that is the first thing to establish — nothing below is actionable without it. This is a discussion framework for an expert cancer-predisposition clinic, not a substitute for one:

Period	Surveillance	Frequency
Birth to age 7	Renal ultrasound (we suggest extending the field to liver, retroperitoneum and pelvis)	Every 3 months — consensus recommendation for all MVA conditions (PMID 39264246). The wider abdominal field is our addition
To age 10, then reviewed	Full paediatric and oncological examination: abdomen, nodes, head and neck, orbit , skin, genitalia, soft tissues	3-monthly to age 7, then 6-monthly. The consensus asks for clinical assessment for rhabdomyosarcoma without setting an end date; the frequency and the age-10 review point are our suggestion , prompted by the age-12 orbital ERMS in PMID 42595739. Examines sites outside the renal-ultrasound field; detection performance and benefit of this schedule are unquantified
Post-ERMS	Relapse follow-up per tumour protocol	Coordinate to avoid duplicate imaging
Through childhood	Full blood count with differential	Leukaemia is reported in MVA1, but routine blood-count surveillance has no demonstrated benefit here; whether to do it, and how often, is an individual clinic decision
Lifelong	Family education; annual predisposition clinic	Urgent review for mass, haematuria, abdominal distension, persistent pain, neurological deficit, bleeding or weight loss

If the child is already past 7, missed scans should not be “caught up”: a baseline assessment and an individualised decision by the predisposition team is the right course.

We do **not** propose serial AFP, repeat CT, PET-CT or routine whole-body MRI while asymptomatic — no demonstrated benefit in MVA, weighed against sedation, false positives, cost and radiation. The consensus is explicit that **radiation exposure should be avoided** in these disorders

(PMID 39264246), which is an argument against imaging-heavy surveillance and, separately, a consideration for any future treatment plan.

One implication for the parents, which the phenotype document already hints at. The clinical document lists recurrent spontaneous abortion (HP:0200067) — an obstetric history, not a feature of the child. If the two variants are in *trans* — which is the interpretation this report adopts but does not demonstrate (§10) — each parent would be expected to carry one of them; segregation testing is what would establish that, and it should precede any carrier-specific interpretation. A 2026 study of two unrelated families with unexplained recurrent pregnancy loss identified novel heterozygous *BUB1B* variants and found **significantly elevated premature chromatid separation rates in the carriers' lymphocytes** by G-banding and centromere FISH, with a trend toward reduced BUBR1 (PMID 42434306). Those variants were classified VUS and the cohort is two families, so this is a hypothesis, not a finding about this family. But it is cheap to act on and it is squarely in the family's interest: a PCS assay on parental lymphocytes and formal genetic counselling would test whether the parents show the cellular phenotype those carriers showed — it would not establish that it caused the pregnancy losses — and would inform any future pregnancy. It is proposed for discussion with clinical genetics; local cost and turnaround are unknown.

We have written the request out so that it does not have to be reconstructed. `clinico/PARENTAL_SEGREGATION_ORDER.md` is a one-page laboratory request a clinician can adapt and sign: both loci, two Sanger amplicons with primer sequences, and an interpretation table written **before** the result, covering all five outcomes — including the *cis* result, which would refute our own Track 1 interpretation outright. The primers were designed against GRCh38 with common 1000 Genomes SNVs and RepeatMasker intervals excluded from the primer regions, and each of the four was then counted across the entire assembly: all occur exactly once. **The RepeatMasker exclusion is what caught the bad pair** — the first design the software returned sat inside an *AluY* element 272 bp from the *c.2210* site, with roughly a hundred exact copies in the genome. The genome-wide count is **confirmation, not the filter**, and we state it that way rather than claiming we sifted a bad pair out at the counting step (`clinico/`

PARENTAL_SEGREGATION_ORDER.md). Either way the failure it avoids is real: a primer like that does not fail in review, it fails in the laboratory, and in a segregation test allele drop-out reads as “*this parent does not carry it*”.

What segregation can and cannot deliver. A *trans* result supports the compound-heterozygous model and contributes PM3-type evidence. **It does not by itself reclassify the missense as pathogenic**, and we do not treat it as doing so. A *cis* result is the informative one in the other direction: it would refute the Track 1 interpretation outright.

This is the cheapest experiment in the project that speaks to phase, to a future pregnancy and to HP:0200067 at once, and it needs blood rather than a new laboratory.

8. What we do not propose

- **MPS1/TTK inhibitors and reversine** — not approved, and they weaken a checkpoint that is already insufficient. Reversine is a research tool for *inducing* aneuploidy, the opposite of the goal.
 - **Apcin, proTAME** — APC/C tool compounds, not marketed.
 - **DCZ0415** — experimental TRIP13 inhibitor: wrong gene, not approved.
 - **AICAR and 17-AAG** — genuine hits in the Tang screen, neither marketed. Useful as in vitro positive controls only.
 - **Metformin** — rejected at filter 3 as unevaluable on the mechanism proposed for it (§3.1).
 - **Senolytics** — the BubR1 senescence evidence rests on *genetic* clearance of p16^{Ink4a} cells via the INK-ATTAC transgene (PMID 22048312), which is not a transferable drug; no agent is approved with a senolytic indication; and senescence is an anti-tumour barrier worth preserving in a cancer-predisposition syndrome.
 - **mTOR inhibitors combined with proteasome inhibition** — predicted antagonism (PMID 31530568).
 - **Bortezomib for this child today** — fails filter 4, as set out in §4.
-

9. Scalability

- The **five-filter framework**, particularly filters 3 and 5, applies as *questions* to any repurposing exercise in a rare disease: filter 3 is an exposure plausibility screen, not a universal criterion, and filter 5's mis-segregation form is the MVA instance of the general question *does treatment worsen the disease-relevant defect or select for it?* It is the transferable product of this work.
- **The MVA genes are not interchangeable.** MVA2 arises from biallelic *CEP57* (PMID 21552266) and behaves differently from MVA1 and MVA3 on exactly the axis that matters here — see the *CEP57* caveat below.
- The **proteotoxic vulnerability**, *if* §6 shows it transfers, would extend most plausibly to **TRIP13**-related MVA, which shares both the tumour risk and the severity of checkpoint failure: “individuals with biallelic TRIP13 or BUB1B mutations have a high risk of embryonal tumors ... their cells display severe SAC impairment” (PMID 28553959) — and the 2026 case in §7 is a TRIP13 patient with two embryonal tumours (PMID 42595739). **It should not be assumed for CEP57.** Yost reports that CEP57-related MVA “is not associated with embryonal tumors” with cells showing “minimal SAC deficiency”, and the 2024 consensus is blunter still: **“At the time of writing, none of the 15 individuals with MVA2 described in literature developed cancer”** (PMID 39264246). the cited checkpoint findings do not establish a proteotoxic burden in *CEP57* cells; that transfer is unmeasured. No proteasome data exist for either, and §10 states that transfer to MVA1 itself is unproven.
- **The target to watch is UBE2H.** The 2026 paired CRISPR screens nominate it as a top aneuploid-selective dependency linked to mitochondrial proteostasis (PMID 42094535, preprint). It has no approved inhibitor, so it fails filter 1 today and is not a candidate — but it is a research target nominated by that preprint, and a repurposing exercise repeated in two years should check it before starting again at metformin.
- **A reusable worksheet** is included as the Appendix: the five filters as blank rows, with the metformin and bortezomib verdicts as the worked example, so a team working on a different rare disease can apply the framework without reading this report.

- **The framework was transferred prospectively to the other two MVA genes, and it gives different answers.** `report/TRANSFER_MVA2_MVA3.md` holds the worksheet filled for MVA2 (*CEP57*) and MVA3 (*TRIP13*) with the same candidate and the filters held fixed. For **MVA2 (*CEP57*)**, under the same standards as MVA1: filter 2 passes on plausibility (constitutional mosaic aneuploidy is the defining phenotype, PMID 21552266 — a verdict first set to “no”, then “not evaluable”, and now aligned with the plausibility standard applied to every candidate; `report/TRANSFER_MVA2_MVA3.md`, amendment of 5 October), filter 3 is the same generic cancer-model comparison with disease-specific transfer unestablished for all three genes, and **filter 4 is not evaluable without a patient** — the literature supplies no tumour-treatment setting: an active malignancy **has not been reported** in the 15 published MVA2 individuals (PMID 39264246). That differs from MVA1 only at filter 4’s setting, where the clinical setting requiring reassessment demonstrably exists. It stays **conditional for MVA3**, which shares both the embryonal-tumour risk and the severe checkpoint failure. A valid framework can return the same verdict for three diseases; what the different verdicts for MVA2 and MVA3 show is how gene-specific evidence changes the worksheet entries — a demonstration of the framework, not of its predictive performance, and not evidence that the drug behaves differently in them.
- **The retrieval benchmark is packaged as a tool, not described as one.** `replay/mva_replay.py` takes a candidate gene, two alleles, an HPO set and a background VCF, and returns the planted score, the rank, the ClinVar-whitelist-off counterfactual and an unrelated-HPO control. It is a single file with no dependencies beyond what an Exomiser installation already needs, and a team working on a different rare disease can spike their own gene into a public genome in an afternoon. Its built-in check that a control HPO term is not annotated to the planted gene is what caught the contamination described in `hpo/RESULTADOS.md` — in our own control set.
- **And we tested that sentence rather than leave it as a claim** (`replay/TRANSFER_TRIP13.md`, pre-registered and pushed before the run). From a fresh clone of this repository, the tool planted two public Likely_pathogenic *TRIP13* alleles — chosen by the frozen

pool rule, not by us — into three healthy GIAB genomes and queried them with six HPO terms transcribed from the abstract of the 2026 TRIP13 case (PMID 42595739). TRIP13 came back **rank 1 in 3 of 3 backgrounds (HG001, HG002, HG005) at 0.8332**, absent unspiked in HG001 and HG005 and scoring 0.0 at rank 212 in HG002, and rank 1 / 2 / 3 at 0.4240 under an unrelated phenotype — the same pattern as BUB1B. Each background took under 3½ minutes and about 2.4 GiB. The operational half was the useful one: the clone needed **four provisioning steps, three of which the tool's requirements list did not include**, and the regression tests the README said ran from a bare checkout did not. Both are fixed; the run that found them is in the repository: the provisioning script, and the log of every tool invocation with its exit code.

- **One command regenerates the Track 2 package.** `analysis/make_track2.sh` re-runs the DepMap follow-up (`depmap/05_seguimiento.py`: figure 5 and its tables) when the public downloads are present, renders the five-filter worksheet in both parts of this document from `report/candidates.tsv`, checks a declared set of reported numbers against the results files (`analysis/verificar_afirmaciones.py`: 27 regex anchors — 10 Track 1, 15 Track 2 report, 2 methods form — covering 31 recomputed facts; not every number in the document) and builds the PDFs. Without the public downloads it rebuilds the document from committed results; with them it reruns the specified analyses, and the log says which mode ran. Run from a fresh checkout of the repository it finished in 8 s with a peak of 179 MB, naming the one missing input — the ~700 MB DepMap 24Q4 and PRISM 24Q2 downloads, which are not bundled — and building both PDFs from the committed copies of the figure and tables (`evidencia/make_track2_fresh_checkout_2026-09-08.log`).
- The **gene panel** (BUB1B, CEP57, TRIP13, CENATAC, MAD1L1, MAD2L1BP, CEP192, BUB1, SMC5, TRIM37, CENPE) and the artefact controls used in the mosaicism analysis transfer to any MVA or PCS workup.
- The **micronucleus → scDNA-seq endpoint** applies to mitotic CIN disorders. It does **not** transfer to Fanconi anaemia or Bloom syndrome, whose diagnostic assays are DEB/MMC chromosome breakage and sister-chromatid exchange respectively — a different lesion, not substitutable.

10. Limitations

- Phase is not established: the called heterozygous markers do not support ordinary short-fragment bridging, and the tested reference-panel route was uninformative.** There are no parental samples; the challenge distributes a VCF and raw reads but no alignments, we did not align the reads, and ordinary short-read fragments would not be expected to bridge these two variants — the only heterozygous site between them lies 6.8 kb from one and 4.1 kb from the other, beyond any short-read fragment, so read-backed phasing would need long or linked reads; the two variants sit **10,911 bp** apart and the caller emitted no phase tag for either. Reference-panel phasing fails for a separate and quantifiable reason: inside that interval the child carries exactly three heterozygous sites — the two candidate variants and `chr15:40216470` — and **all three are absent from the 1000 Genomes 30× phased panel** (3,202 individuals); the nearest panel-present heterozygous markers are 16.8 kb upstream and 809 bp downstream, with nothing between them to link the alleles. A Beagle 5.5 run against that panel keeps 1,391 of 1,452 window variants and drops all three that matter. The absence is what the frequencies predict: expected copy counts in 6,404 haplotypes are 0.64 for the nonsense and **0.006** for the missense. Two heterozygous variants in a singleton could be cis or trans, the evaluated VCF markers and reference-panel analysis did not weigh one over the other, and rarity does not phase them. What supports *trans* is that the child is affected, that MVA1 is recessive, that truncating-plus-missense is the recurrent pattern, and that the phenotype matches. **It remains a diagnosis pending segregation, and parental genotyping at the two loci is the cheapest experiment in this report.**
- The retrieval pipeline is itself insensitive to phase.** In a controlled test on public data, Exomiser 15.1.0 reads phase into its ACMG evidence (trans → PM3, cis → BP2) but not into the compound-heterozygous model: two pathogenic nonsense alleles declared on the *same* chromosome are still called AR_COMP_HET, still classified PATHOGENIC and still ranked first, at an absolute

combined-score difference of 0.0034. Nothing in the Track 1 retrieval would have flagged a *cis* configuration.

- **p.Asn1002Lys has no functional assay**, and its apparent ClinVar entry is for a different nucleotide.
- **The bortezomib EC50 is a conservative upper bound** read from a figure panel rather than a tabulated value, and the Cmax is a label interval rather than a point estimate; the ratio of 5.8–7.8 should be read as an order-of-magnitude statement computed on total, not free, drug.
- **Ippolito et al. studied cancer aneuploidy**, not constitutional variegated mosaicism. Whether the proteotoxic dependency transfers is precisely what the proposed experiment tests; we have not assumed it.
- **No trial specific to MVA1** was found in ClinicalTrials.gov, Open Targets, ChEMBL or CMap/LINCS. Absence of an interface result is not proof of absence.
- **The power estimate rests on unmeasured inputs.** Under the declared four-parameter model with an assumed CV_bio of 0.20 the design needs about 12 independent cultures per arm; CV_bio, the aneuploid fraction and the 4× selectivity are all unmeasured in these cells, and the original 0.897 is an optimistic two-parameter reference.
- **The child's current age is unknown to us**, and the surveillance schedule depends on it.

A closing note

The family of this child published their son's genome so that strangers might help. They will read what comes back.

We have tried to write something useful rather than something reassuring. The honest answer to “what drug should he take today” is none — and for the obvious candidate we identify a concentration mismatch for the proposed complex-I route and no usable aneuploidy-selective potency estimate, rather than proposing it softly. The useful answer is a mechanism that matches his disease, an approved drug whose plasma levels are not excluded by the concentrations at which cancer cells respond, an explicit statement of the condition under

which it would become a question for a tumour board — a changed comparator, not an established benefit — and an experiment — an eight-week assay window, 12–16 weeks from a fresh biopsy, more if twelve cultures per arm must be expanded — that could fail to support it in this child’s own cells. That experiment would not close the tumour-board question on its own, and §6 says so.

That is not a cure. It is a prepared answer to a question this family may unfortunately have to ask.

Appendix — The five-filter worksheet (reusable)

*The transferable product of this work. Rendered from `report/candidates.tsv` by `report/render_candidates.py`, so the table and the machine-readable file are the same object; every candidate this report evaluated is a row, with a blank row for a new disease. These five questions organise this candidate assessment; satisfying them would not establish clinical benefit or authorise treatment. The assessment makes exposure plausibility and disease-worsening risk explicit. Statuses describe evidence, not permission to investigate: *pass** — the stated criterion is supported in the named setting; *fail* — evidence opposes it there; *untested* — no direct test was performed; *not evaluable* — a required input is missing; *not reached* — not assessed because an earlier filter stopped the candidate; *not excluded* — only the specified total-plasma comparison; *exposure unresolved* — available pharmacokinetic information does not settle exposure plausibility. Unresolved evidence can motivate research (bortezomib’s status) but cannot establish clinical suitability. Filter 2 is applied as mechanistic plausibility; demonstrated disease-specific activity is a separate evidence field, unestablished for every candidate here.**

Candidate	F1 approved	F2 mechanism	F3 exposure	F4 safe in this patient	F5 margin excluded	Verdict
metformin	pass	pass	not evaluable — none: no aneuploidy-selective effective concentration exists; complex I route needs mM vs uM plasma	not reached	not reached	rejected at filter 3 unevaluable — Marketed stand-in for the AICAR/AMPK screen hit; the complex I mechanism cannot operate at plasma concentrations and an AMPK-mediated aneuploidy-selective effect is untested
bortezomib IV	pass	pass	not excluded — EC50 <40 nM read from Fig. 6o of PMID 39247952 (5 near-euploid vs 5 highly aneuploid cancer lines, 72 h); Cmax 89-120 ng/mL = 231-312 nM at 1.3 mg/m ² IV (FDA label NDA 021602); ratio 5.8-7.8 on total plasma drug; 83% protein-bound so the free peak is ~39-53 nM (~99 nM on the label's 223 ng/mL cohort), about 1.0-1.3x (2.5x) a 40 nM reference; comparable free-medium potency unknown	fail	untested	fails filter 4 for this child; conditional for active tumour — Proteasome dependency of aneuploid cancer cells; paediatric neuropathy 18% (motor 8%) in a trial with combination chemotherapy, on pre-existing muscle atrophy; own DepMap test attenuated in solid lineages; RMS essentially absent from the drug screen
bortezomib SC	pass	pass	exposure unresolved — Cmax 20.4 ng/mL = 53.1 nM SC (free ≈ 9 nM); nominal ratio 1.33 on total drug at the peak, AUC equivalent to IV; a free-fraction correction cannot be applied to one side only	fail	untested	exposure unresolved; fails filter 4 — Subcutaneous route withdrawn at filter 3
carfilzomib	pass	pass	not evaluable — no IC50 in an aneuploidy-stratified model; Cmax 2.89 uM at 56 mg/m ² known, no valid denominator	not reached	not reached	not evaluable — Same mechanism class; no aneuploidy-stratified effective concentration published
ixazomib	pass	pass	not evaluable — beta5 enzymatic IC50 3.4 nM is not a cellular aneuploidy-selective value; Cmax 171 nM at 4 mg oral	not reached	not reached	not evaluable — No valid denominator; own DepMap correlation ran the wrong direction (rho +0.028)
hydroxychloroquine	pass	untested	not evaluable — Tang's hit was chloroquine, which did not differentially inhibit the human CIN lines in the same paper	not reached	not reached	not evaluable mechanistically weakened — Mechanism not supported in human CIN lines
everolimus sirolimus	pass	untested	not evaluable — predicted to antagonise proteasome inhibition by relieving translational burden; whole-blood, not plasma, Cmax reported	not reached	not reached	not evaluable possibly counterproductive — Predicted antagonism with bortezomib; muscle-directed mTOR rationale in monoallelic mice points the other way and is not support for the combination
trametinib	pass	pass	not evaluable — only relative IC50s recoverable from PMID 39251587; 0.45 nM sensitising dose is sub-lethal, not an IC50; Cmax 22.2 ng/mL = 36.1 nM at 2 mg/day (SmPC §5.2)	not reached	not reached	not evaluable gap most worth closing — Aneuploid RPE1 clones are MEK-inhibitor-sensitive; approved from age 1 with dabrafenib for BRAF V600E low-grade glioma; first candidate to re-run filter 3 once an absolute aneuploid IC50 is published
AICAR	fail	pass	not reached — screen hit in isogenic trisomic MEFs and human trisomy-7 colonic cells	not reached	not reached	fails filter 1 — Not marketed; in-vitro positive control only
17-AAG HSP90i	fail	pass	not reached — 200 nM active dose in the Tang screen, not an IC50; no FDA/EMA-approved HSP90 inhibitor	not reached	not reached	fails filter 1 — Not marketed; class limited by organ toxicity
reversine MPS1 TTK inhibitors	fail	fail	not reached — research tools that induce aneuploidy	not reached	untested	fails filter 1 and 2 — Weaken a checkpoint that is already insufficient
apcin proTAME	fail	untested	not reached — APC/C tool compounds	not reached	not reached	fails filter 1 — Not marketed
DCZ0415	fail	fail	not reached — experimental TRIP13 inhibitor	not reached	not reached	fails filter 1 and 2 — Wrong gene for MVA1; not approved

Candidate	F1 approved	F2 mechanism	F3 exposure	F4 safe in <i>this</i> patient	F5 margin excluded	Verdict
senolytics	not evaluable	untested	not reached — BubR1 senescence evidence rests on genetic INK-ATTAC clearance, not a drug	not reached	untested	not proposed — No specific compound is nominated by the cited evidence, so the marketing criterion cannot be applied to the class; senescence is an anti-tumour barrier in a predisposition syndrome
UBE2H inhibitor	fail	pass	not reached — top aneuploid-selective dependency in paired CRISPR screens (preprint)	not reached	not reached	target to watch — No approved inhibitor exists
<i>(your candidate)</i>						

How to fill it in. Filter 3 needs a real denominator: an effective concentration measured in a model stratified for the disease mechanism. If none exists, write *not evaluable* — do not substitute an IC50 from an unrelated model, because that manufactures a ratio with no meaning. Filter 4 is patient-specific by construction, so a “fail” here is a statement about one person and one moment, not about the drug. Filter 5 asks the question repurposing exercises forget: among the cells that survive the drug, is the underlying defect made worse beyond a prespecified margin? A machine-readable version of this worksheet, with the full status vocabulary (pass / fail / untested / not evaluable / not reached / not excluded / exposure unresolved), the model context behind each filter-3 entry, a source locator and a reason, is `report/candidates.tsv`.

Appendix B — Methods description (Track 2 form)

Q2 — Describe your approach in detail (variant/mechanism → candidate medication)

We went from variant to candidate in five steps; exposure was most often unevaluable, and the direction-of-effect question remains untested for the one candidate not excluded.

Step 1 — define the working hypothesis. Two heterozygous BUB1B variants, presumed in *trans* (phase not established; Q9): `c.2210T>G` p.Leu737Ter, a nonsense allele that, if translated, predicts a 736-residue product lacking residues 737–1050 (314 residues), including the kinase domain (UniProt O60566 annotates it at 766–1050), paired with `c.3006T>G` p.Asn1002Lys inside that domain. Working hypothesis: loss of function, matching the truncating-plus-missense pattern of viable MVA1 — a truncating allele paired in *trans* with a functionally impaired missense allele, phase and residual function unestablished. Reading the pair as null-plus-hypomorph is our reading of the human truncating-plus-missense pattern (PMID 20516114), which does not state that complete biallelic loss is lethal, and neither do we; the

hypothesis that p.Asn1002Lys retains residual function is motivated by the Sieben mouse series (PMID 31738183), not by 20516114, and its rank relative to L1012P is not established.

Step 2 — follow the mechanism downstream, not at the lesion. No approved drug restores BubR1. We therefore asked what the lesion *produces* that is druggable. The chain is: weakened spindle assembly checkpoint → chromosome mis-segregation → constitutional aneuploidy → stoichiometric imbalance of protein complexes → proteotoxic stress. The last link is the one with approved drugs against it.

Step 3 — five filters instead of the usual two. Beyond the two questions every exercise asks — is it approved, is the mechanism plausible — we added three:

#	Filter	Question
1	Regulatory	Approved for marketing?
2	Mechanistic	Does it address the actual lesion or its consequence?
3	Pharmacokinetic	Is the achievable exposure excluded by the effective in-vitro concentration? (“not excluded” is the best outcome this screen can return)
4	Patient-specific safety	Compatible with <i>this</i> child’s comorbidities?
5	Direction of effect	Can an increase in mis-segregation among survivors beyond a prespecified margin be excluded?

Step 4 — apply them, and report what dies. Metformin — the candidate the literature points to, descending from the AICAR hit of Tang et al. (PMID 21315436) — is not evaluable at filter 3: no published study gives metformin an aneuploidy-selective effective concentration to compare against, and the complex-I mechanism proposed for it carries a separate objection — therapeutic plasma levels are micromolar while complex I inhibition requires millimolar (PMID

37343530). MPS1/TTK inhibitors and reversine are opposed at filter 2 by construction — they weaken a checkpoint that is already insufficient — and the quantitative filter-5 margin is untested for them.

Step 5 — the candidate not excluded, with its stopping condition.

Bortezomib passes filters 1–2, is not excluded by filter 3 on a total-plasma peak comparison (free exposure and duration at the target unresolved), and fails filter 4 *for this child*, which is a patient-specific verdict rather than a property of the drug. We state the condition under which it would become a question for a tumour board — an MVA patient with an active tumour, where the comparator is cytotoxic chemotherapy rather than nothing; a changed comparator, not an established benefit — and we propose an experiment — an eight-week assay window, 12–16 weeks from a fresh biopsy because establishing the fibroblast line takes 4–8 weeks on its own, more if twelve cultures per arm must be expanded — that could fail to support the dependency in this child’s own cells. It would not, on its own, close the tumour-board question, and we say so where we describe it.

Q3 — Generative AI declaration

Commercially-available generative AI was used. Provider, plan and relevant setting for each:

- **Anthropic**, Claude Opus and Claude Fable 5.1 via Claude Code, **Max** subscription plan, model-improvement data sharing disabled in account privacy settings.
- **OpenAI**, Codex (gpt-5.6), **Plus** plan, “Improve the model for everyone” disabled.
- **xAI Grok 4.6 via Cursor**, **Pro** plan, **Privacy Mode enabled** (zero data retention at Cursor and at the underlying model providers).
- **Qwen3-TTS 1.7B VoiceDesign** — open-weights, run locally on CPU with no cloud service and no API key. Used only to narrate the pitch video; not a commercial service, listed for completeness. Each clip was transcribed by a local recogniser and diffed against the written script before assembly, because this class of model can alter a word silently.

No block of the VCF and no genotype table was ever sent to a language model. Genomic work ran on local tools (bcftools, Exomiser) and on annotation APIs — Ensembl VEP and ClinVar — which return public annotation for submitted coordinates and acquire no rights over the input. The policy does not name these services; we read them as satisfying the two conditions it does set, namely no training on inputs or outputs and limited retention.

Q4 — Automated or manual candidate identification?

Primarily manual literature review and expert curation, with LLM-assisted retrieval. No drug-target database query produced the candidate. Interface queries against ClinicalTrials.gov, Open Targets, ChEMBL and CMap/LINCS returned no MVA1-specific drug pairing, which we report as a negative result rather than omit.

Q5 — Describe the manual review and curation

Three activities, all documented in the repository:

1. **Citation verification.** Every PMID was checked against PubMed or Europe PMC before use. This was not ceremonial: **three citation problems were caught, of three different kinds** — one wrong paper (a study of acromegaly cited as Wilms surveillance, PMID 17652220, replaced by 16857697); one merged attribution (two KARD papers under a single citation, separated into PMID 23789096 and 23345399); and one over-extension (PMID 28553959 was being used to carry the proteotoxic rationale to CEP57, which that same source contradicts — we withdrew that use and kept the paper only for the BUB1B/TRIP13 embryonal-tumour sentence its abstract actually states).
2. **Adversarial review by independent models.** Two different language models were asked to attack our own drafts. They found, among others, a claim of ours that contradicted **our own data file**: the report stated that Tang et al. used only stable trisomies, when our extraction record shows Figure 6 of that paper tests *Bub1b*^{H/H} hypomorphic MEFs. Corrected. All review rounds are in [evidencia/](#).
3. **Pharmacokinetic arithmetic, redone by hand.** Each C_{max} was converted from ng/mL to molar using the molecular weight,

against the FDA label section cited. This surfaced that our own margin used the ceiling of a label interval (89–120 ng/mL) rather than the interval; corrected to 5.8–7.8×.

Q6 / Q7 — Data sources

Publicly available drug, target and literature sources only; no proprietary data. The one controlled-access input is the challenge dataset itself (the proband's VCF and clinical document), used under the data transfer agreement described under *Data availability* and deleted within 30 days of the close.

- **Literature:** PubMed / Europe PMC, cited by PMID throughout.
- **Regulatory labels:** FDA labels and reviews — VELCADE (NDA 021602 s040; BPCA Clinical Review 2015), NDA 206927 s002, Kyprolis (NDA 202714 s010), Ninlaro (NDA 208462), Rapamune, Afinitor, Mekinist.
- **Variant and gene resources:** ClinVar, gnomAD v4, Ensembl REST and VEP, UniProt (O60566), OMIM and Orphanet identifiers, Human Phenotype Ontology.
- **Drug/target interfaces queried (all negative for MVA1):** ClinicalTrials.gov, Open Targets, ChEMBL, CMap/LINCS.
- **Pathway and model resources:** Reactome (R-HSA-69618), KEGG (hsa04110), MGI (MGI:1333889).
- **Assay standards:** OECD Test Guideline 487 (micronucleus assay).

Q8 — Proprietary data

None.

Q9 — How did you characterize the variant's mechanism?

Working hypothesis: impaired BUB1B function, with a candidate drug target downstream; phase and the missense allele's functional effect remain unresolved.

The nonsense allele removes the entire C-terminal kinase domain; the missense allele sits inside it. Cell lines derived from MVA patients with biallelic mutations show “an impaired mitotic checkpoint, chromosome alignment defects, and low overall BUBR1 abundance”

(PMID 20516114), and allele-specific effects — not merely total BubR1 quantity — drive phenotypic heterogeneity (Sieben 2020, PMID 31738183).

We were explicit about what we could not determine. Phase is not established, and the routes available to us were uninformative: there are no parental samples; the challenge distributes a VCF and raw reads but no alignments, we did not align them, and short reads could not phase two variants whose only intervening heterozygous site lies 6.8 and 4.1 kb away; and the two variants lie 10,911 bp apart with no GATK phase tag on either — a limit visible in the data itself, since GATK did phase three variants 5 bp apart elsewhere in the same VCF. Reference-panel phasing fails separately: all three heterozygous sites inside the interval are absent from the 1000 Genomes 3,202-sample phased panel, and Beagle 5.5 drops all three. Our retrieval pipeline is insensitive to phase in any case: two pathogenic nonsense alleles declared in cis are still called compound heterozygous, still PATHOGENIC and still ranked first, at an absolute score difference of 0.0034. Nonsense-mediated decay of the truncated transcript is unmeasured and we do not assert it either way without patient RNA. p.Asn1002Lys has no published functional assay, and its in-silico predictors disagree (REVEL 0.472, AlphaMissense 0.923), which is why we let it stand as a VUS.

We also asked whether aneuploidy burden could serve as a trial endpoint in these data, and computed that it cannot. A first-order calculation — mean $|\text{BAF} - 0.5|$ at DP 44 against a measured between-chromosome SD of 0.0035 — puts the detection limit at ~14% of cells for a clonal trisomy (16% under a stricter 2-SD criterion). A later calibrated simulation (Ψ over 125-site windows, Monte Carlo null) improves that to **2.07% at an overdispersion $\kappa = 2$** ; that is a *simulated* limit conditional on κ , not a measured sensitivity, and it places variegation's 1–2% per chromosome at the edge of what is achievable rather than far below it. The argument that does not depend on any limit is **non-identifiability**: for a perfectly balanced variegated mixture, mean copy number is exactly 2 and mean allele fraction exactly 0.5 by symmetry, and simulated detection stays at the 0.1–0.2% false-positive rate even with 40% of cells aneuploid. That proof covers

the balanced mixture we modelled, not every possible MVA genome; it disqualifies the endpoint for that case, and for this child no clinical negative is reported.

We ran the calibrated statistic on the proband and **withdrew the result**: the pre-registered threshold fires on all 22 autosomes, which indicates a null model that does not fit these data, and the pre-registered external control came back **not conclusive** because the control genomes were called differently from the proband. We report no clinical negative for this child.

Q10 — Time and effort

Approximately 30 hours of analyst time over the first two days (28–29 August 2026), of which the large majority went to literature verification and adversarial review rather than to computation; the later pre-registered rounds (30 August to 5 October) added several working days of analysis, verification and adversarial review, not logged to the hour. Compute was negligible: the genome-wide prioritisation ran in 53 seconds on a commodity workstation with no GPU, and the robustness controls in about ten minutes. Later rounds added the pre-registered analyses, on public inputs and, for the mosaicism work, the controlled-access proband VCF: eight in total — DepMap and its follow-up, MVA-Replay, mosaicism, power and its four-parameter follow-up, dose-response, and the TRIP13 transfer (VERIFY.md lists each pre-registration with its result file and amendments). The phase pilot is **not** one of them — it ran before any pre-registration and we do not count it. All of this is CPU-only except the mosaicism null calibration, which used one consumer GPU.

Q11 — Method abstract (≤ 500 words)

No approved drug restores BubR1 function. We followed the lesion downstream — weakened spindle assembly checkpoint, mis-segregation, constitutional aneuploidy, stoichiometric protein imbalance, proteotoxic stress — and asked which approved agent acts on the last link.

We applied five explicit filters. Beyond “is it approved” and “is the mechanism plausible”, we added a pharmacokinetic filter (is achievable exposure excluded by the effective in-vitro concentration?),

a patient-specific safety filter, and a direction-of-effect filter (can an increase beyond a prespecified micronucleus harm margin among survivors be excluded?).

Metformin, the candidate the literature points to, is not evaluable on filter three: no aneuploidy-selective effective concentration exists, and its complex-I mechanism needs millimolar levels against micromolar plasma (PMID 37343530). Bortezomib passes filters one and two (mechanism plausible in cancer aneuploidy; our own pre-registered test on public solid-tumour lines did not establish it there) and is not excluded by three. Aneuploidy level was associated with multiple myeloma patients' response to proteasome inhibitors (Ippolito 2024, PMID 39247952), on cohorts small enough (8 complete responders vs 50 progressive) that we treat it as motivation, not prediction. Label Cmax 231–312 nM (1.3 mg/m² IV) over a 40 nM reference bound read from a figure gives 5.8–7.8× on **total plasma** drug — Cmax/40, not Cmax/EC50; free peak ≈ 39–53 nM — not a therapeutic margin.

It then fails filter four for this child: peripheral neuropathy in 18% (motor 8%) of patients in a paediatric trial with combination chemotherapy, on top of existing skeletal muscle atrophy. But filter four is patient-specific. For an MVA patient with an active malignancy — where the comparator is cytotoxic chemotherapy — the comparator changes, though benefit remains unestablished; biallelic BUB1B disease carries a high risk of embryonal tumours (phase not established here). Our deliverable is a conditional tumour-stage research comparison, with unresolved exposure and selectivity requirements, plus a separate constitutional-cell vulnerability assay.

Strengths. Every candidate carries a stated status against a stated criterion, including our own; filter five stays untested; §8 lists what we excluded. We report a predicted antagonism: reducing translation protects CIN cells from proteasome inhibition (PMID 31530568), so mTOR inhibitors would antagonise, not synergise. The questions apply to any rare disease; re-run on MVA2 and MVA3 — a transfer inside the MVA group, not a general demonstration — they return different verdicts.

Limitations. The proteasome dependency is established in cancer aneuploidy, not constitutional mosaicism; §6 measures response and constitutional-cell vulnerability, not mechanism. Our own pre-registered DepMap analysis found the association attenuated in solid lineages, the lineage of this child's tumour. Under the pre-registered four-parameter model ($f = 0.30$, $4\times$ selectivity, CV 0.20) the equal-EC50 test has simulated power 0.17 at $n = 3$ and first exceeds 0.80 at $n = 12$; the full advancement rule is not powered. Phase is not established; parental genotyping is the cheapest informative experiment. No disease-modifying drug can be recommended from these data today; the clearest action supported by guidance is surveillance, not a molecule.

Acknowledgement

This work was made possible through the Hackathon, organized by Sage Bionetworks in partnership with the MVA Society, Hugging Face, and BEACON (The Benchmarking, Evaluation, and Assessment Consortium for Science), with prize sponsorship from AWS and Anthropic. We are deeply grateful to the child and their family who generously contributed their data and their story to advance research into this rare disease. We acknowledge their trust in making this Hackathon possible.

Data availability

The individual-level data analysed here is the Hackathon dataset (SageBio/mva-hackathon-2026-data), accessed under the challenge data transfer agreement accepted by Hugging Face account `cpu-16` on 2026-08-28, and cited as directed on the Hackathon Synapse page. **No patient sequence data or genotype-scale derivative is included in this document or in our public repository**, and every copy we hold will be deleted by 2026-11-23 with written confirmation to the organisers. The two diagnostic variants are named, as findings, which the challenge rules permit. All other inputs are public: Exomiser 15.1.0 with data release 2602, ClinVar, gnomAD v4, the Human Phenotype Ontology, Ensembl and VEP, the 1000 Genomes and GIAB genomes,

DepMap 24Q4, AlphaFold and the literature cited by PMID. Analysis code, pre-registrations and results: <https://github.com/cpu-16/mva-hackathon-2026>